

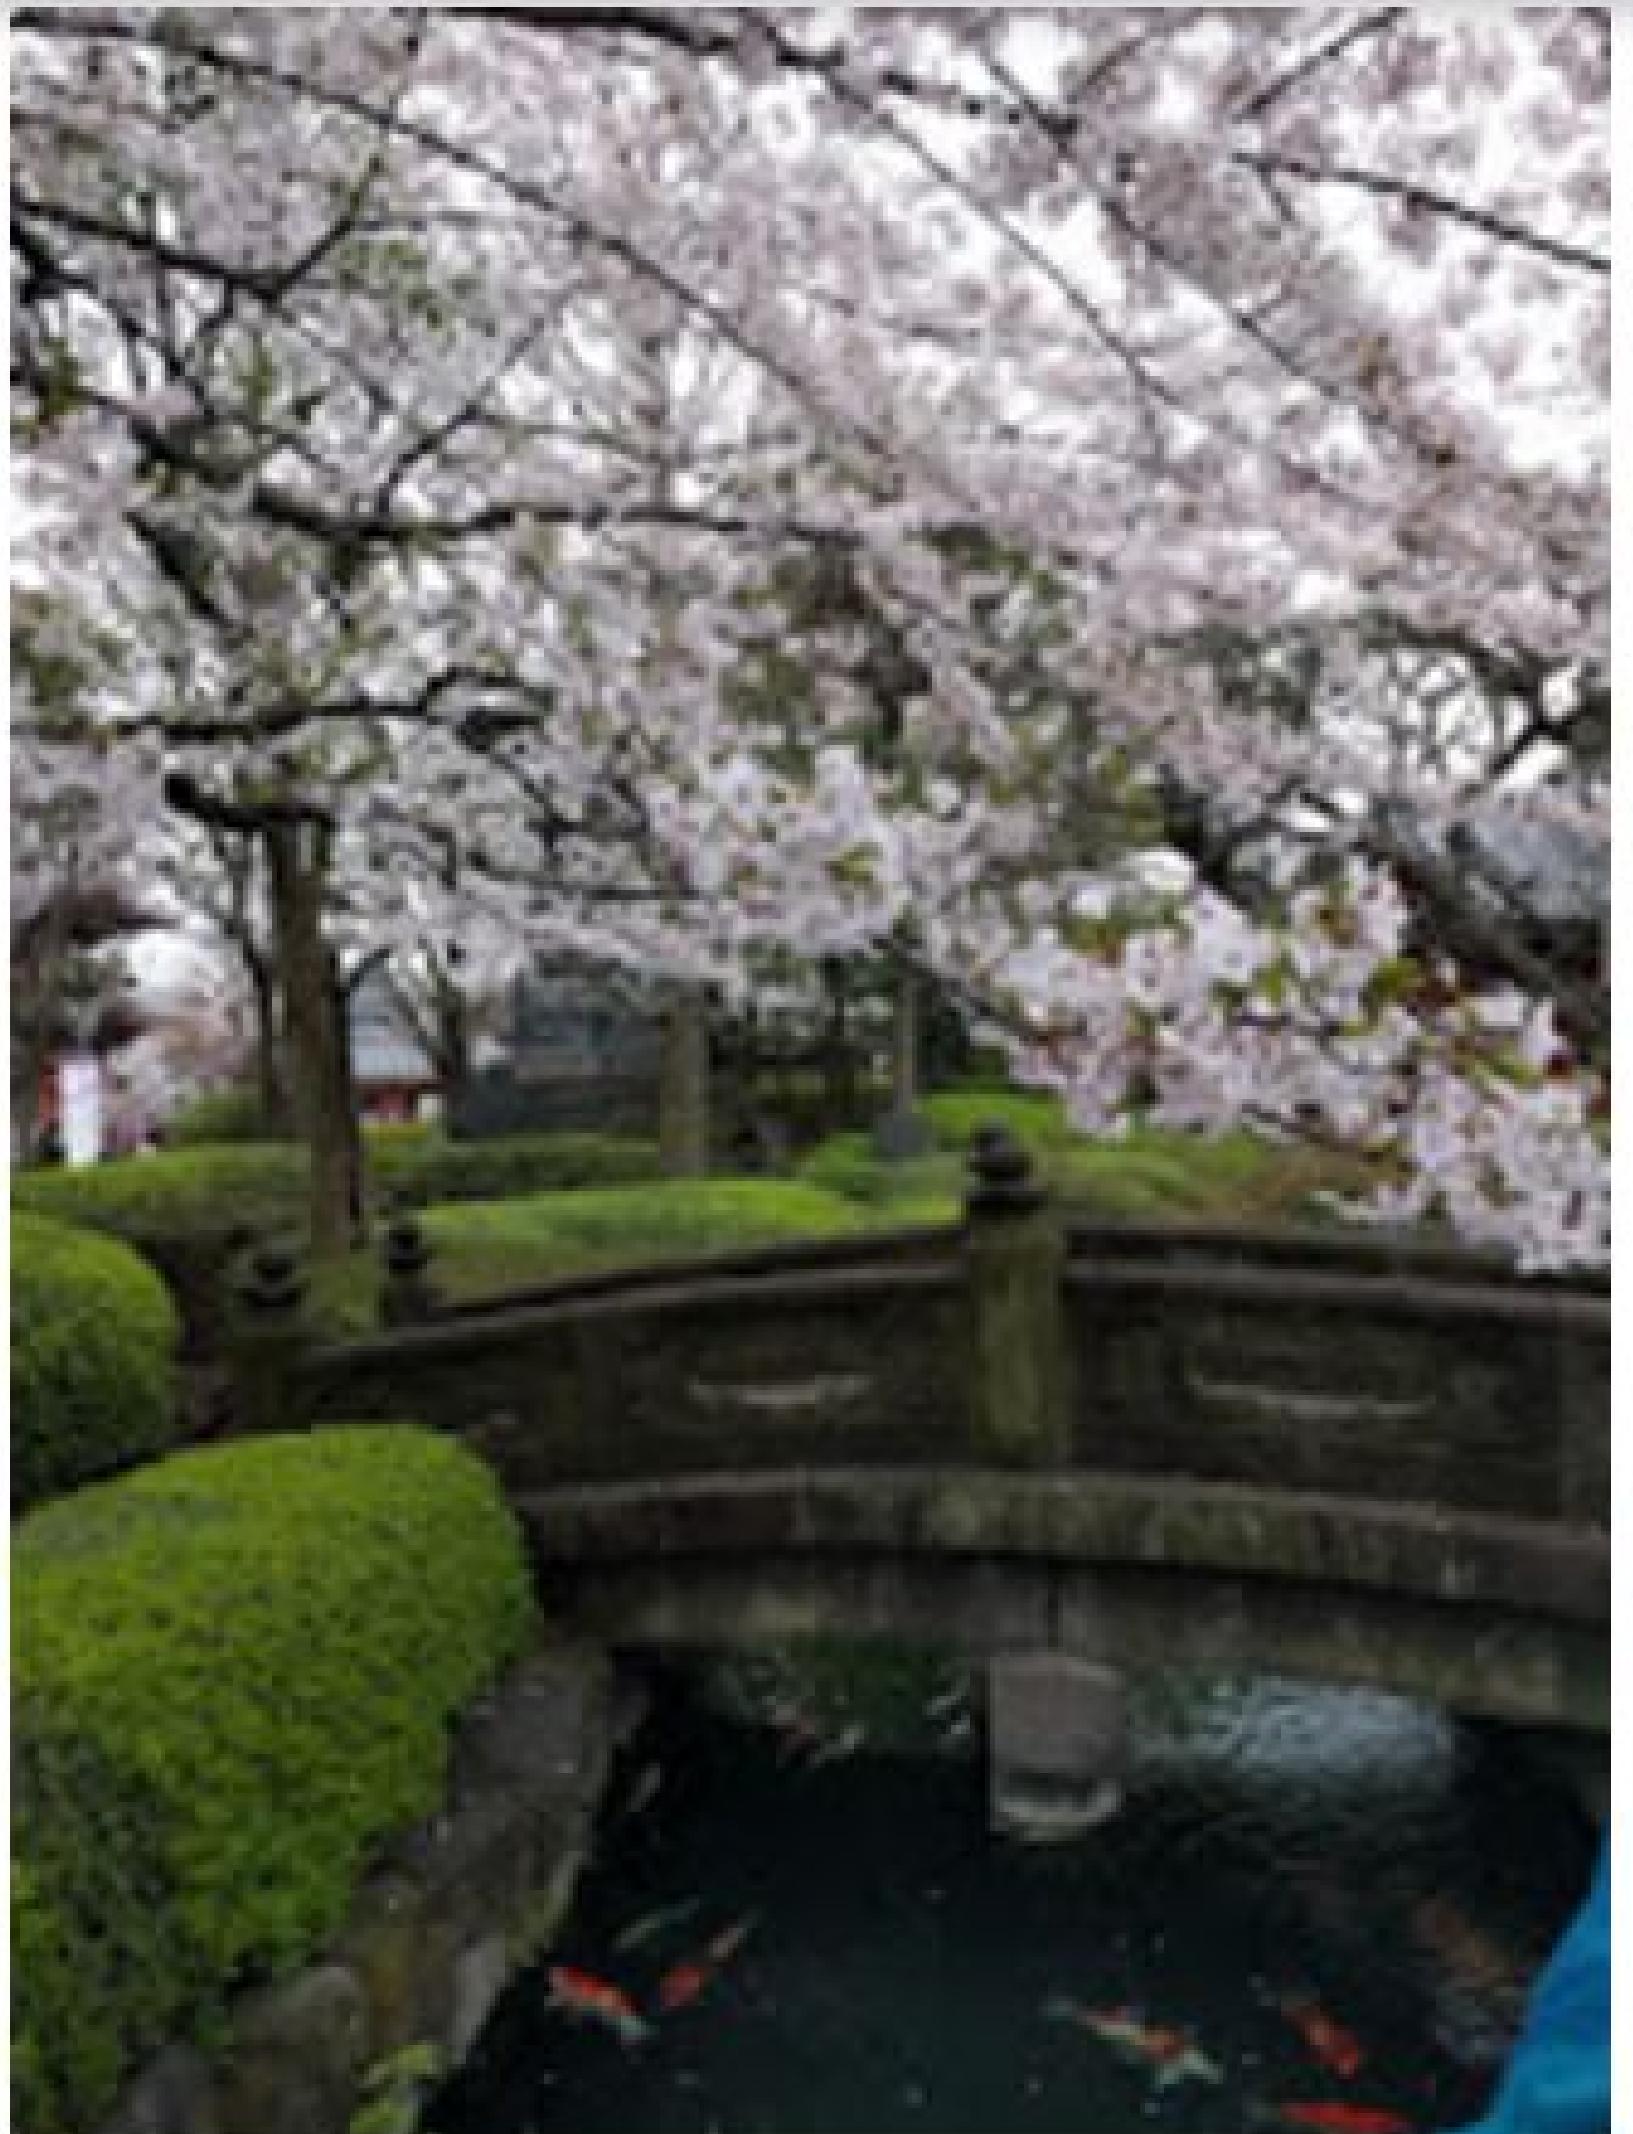
Effective TensorFlow for Non-Experts



Baohua Liao
liaobaohua@gmail.com

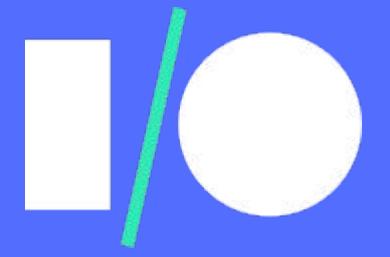


cherry blossom





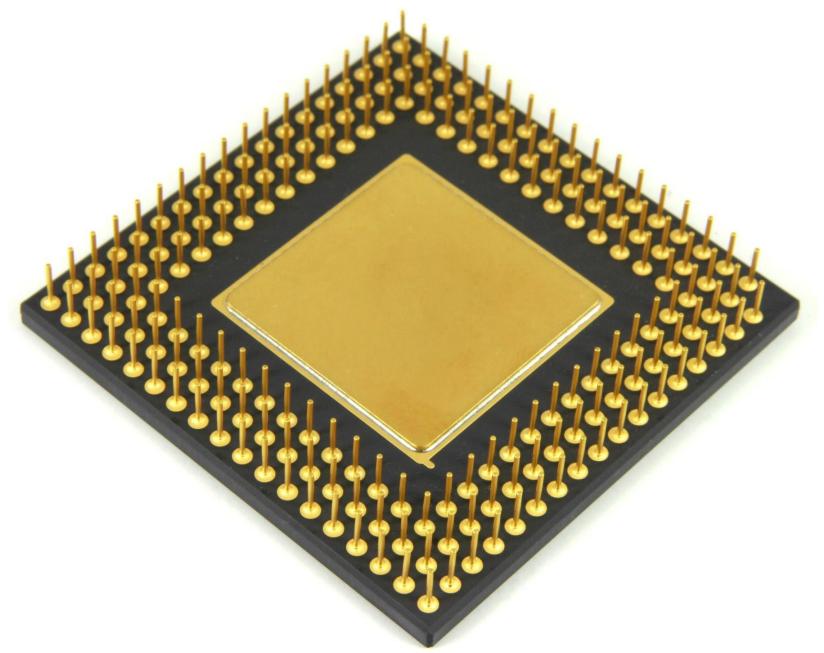
EXIT →



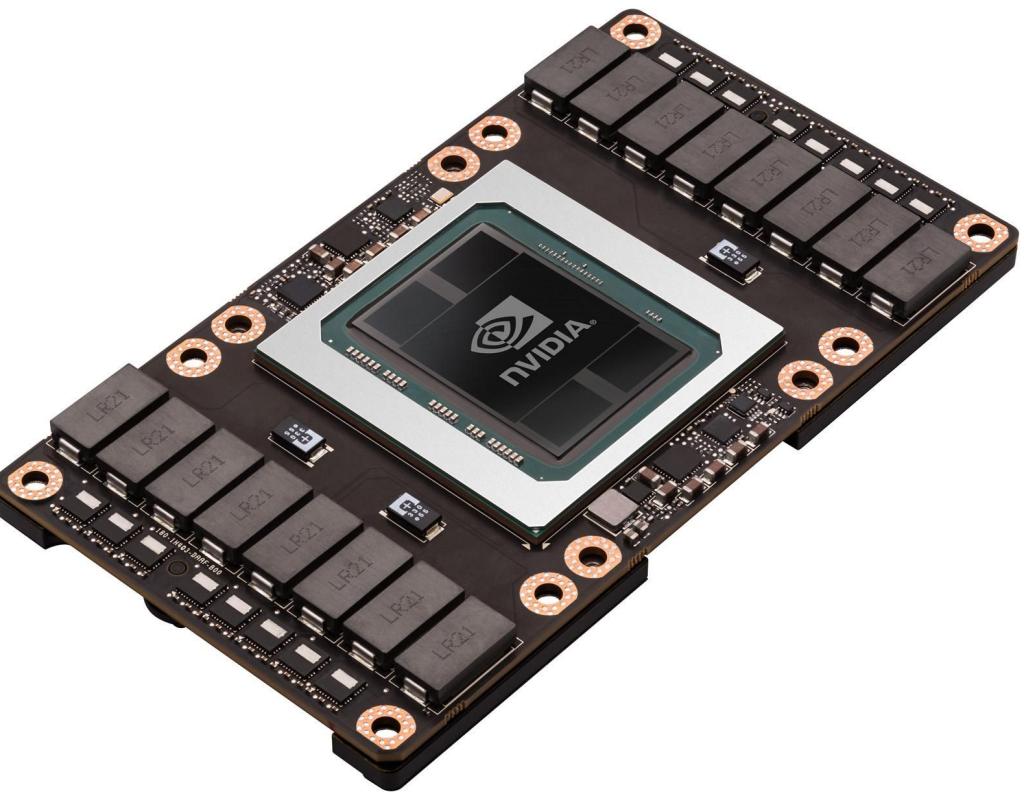
Complexity



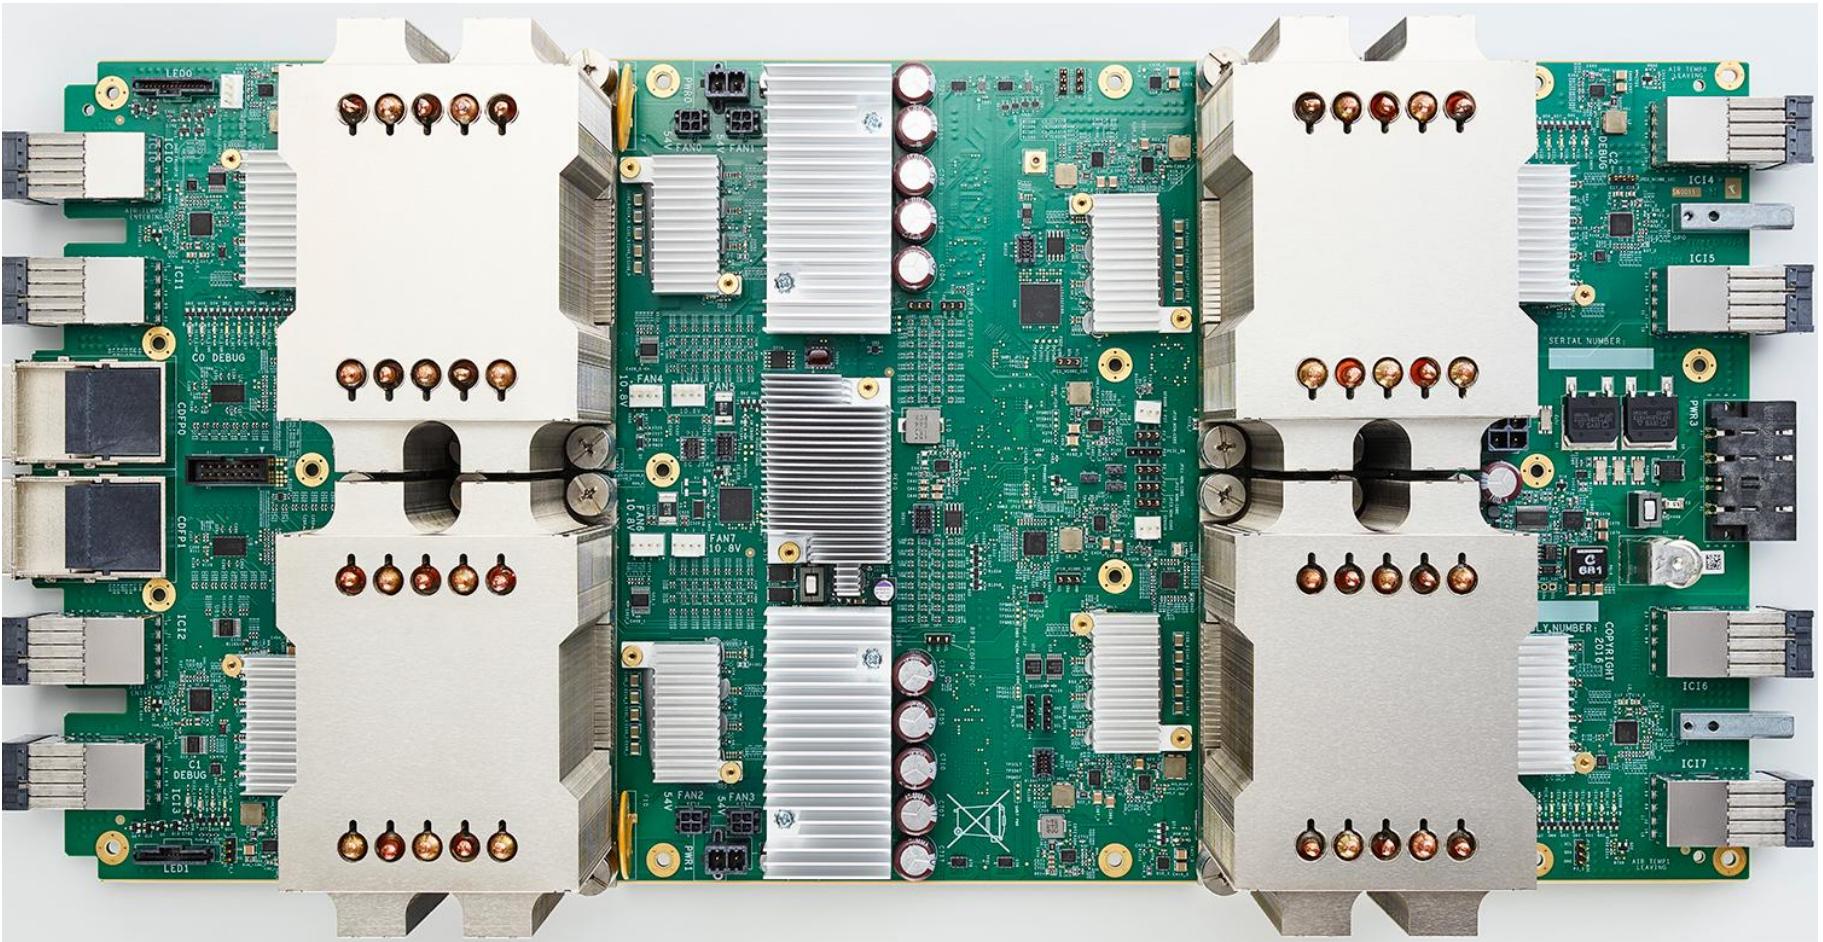




CPU



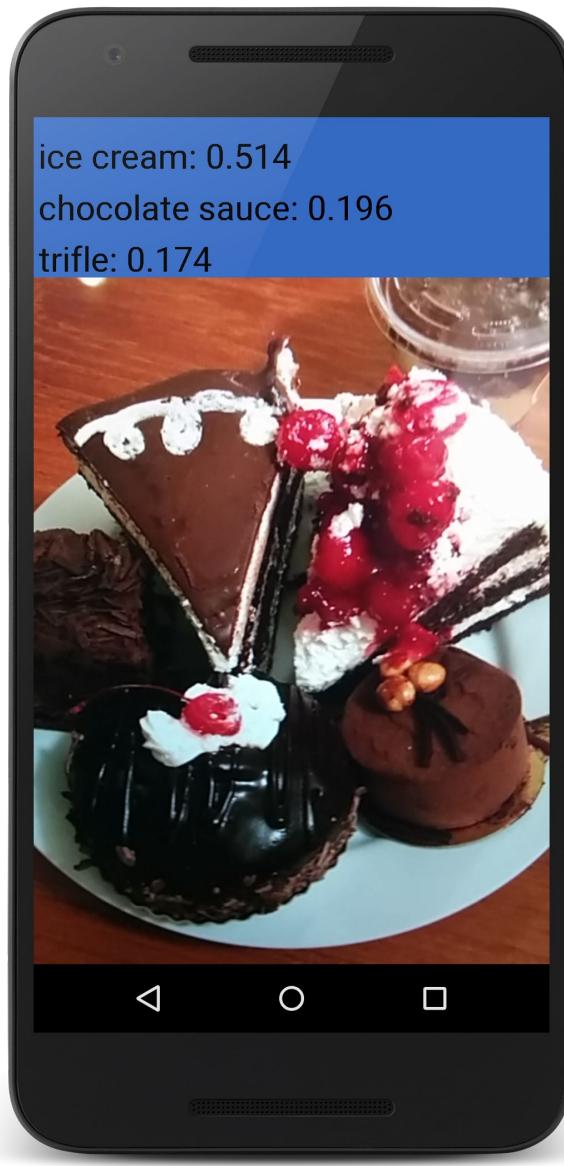
GPU



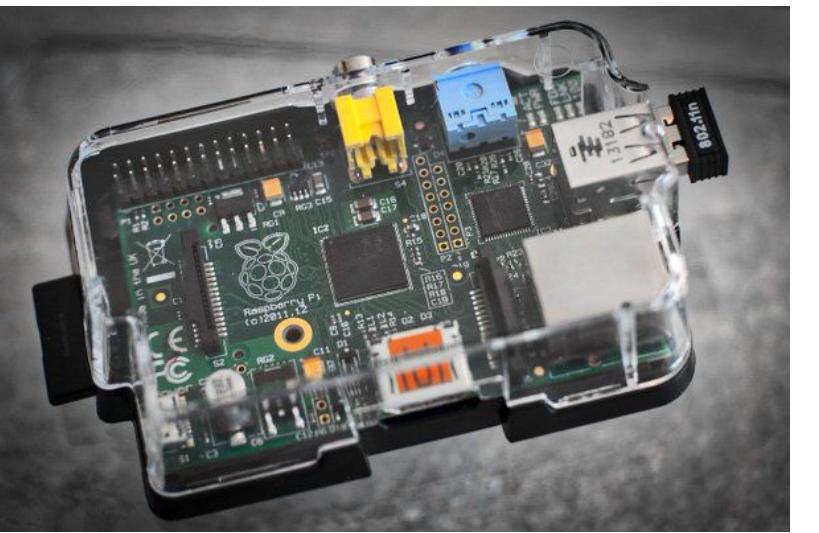
Cloud TPU



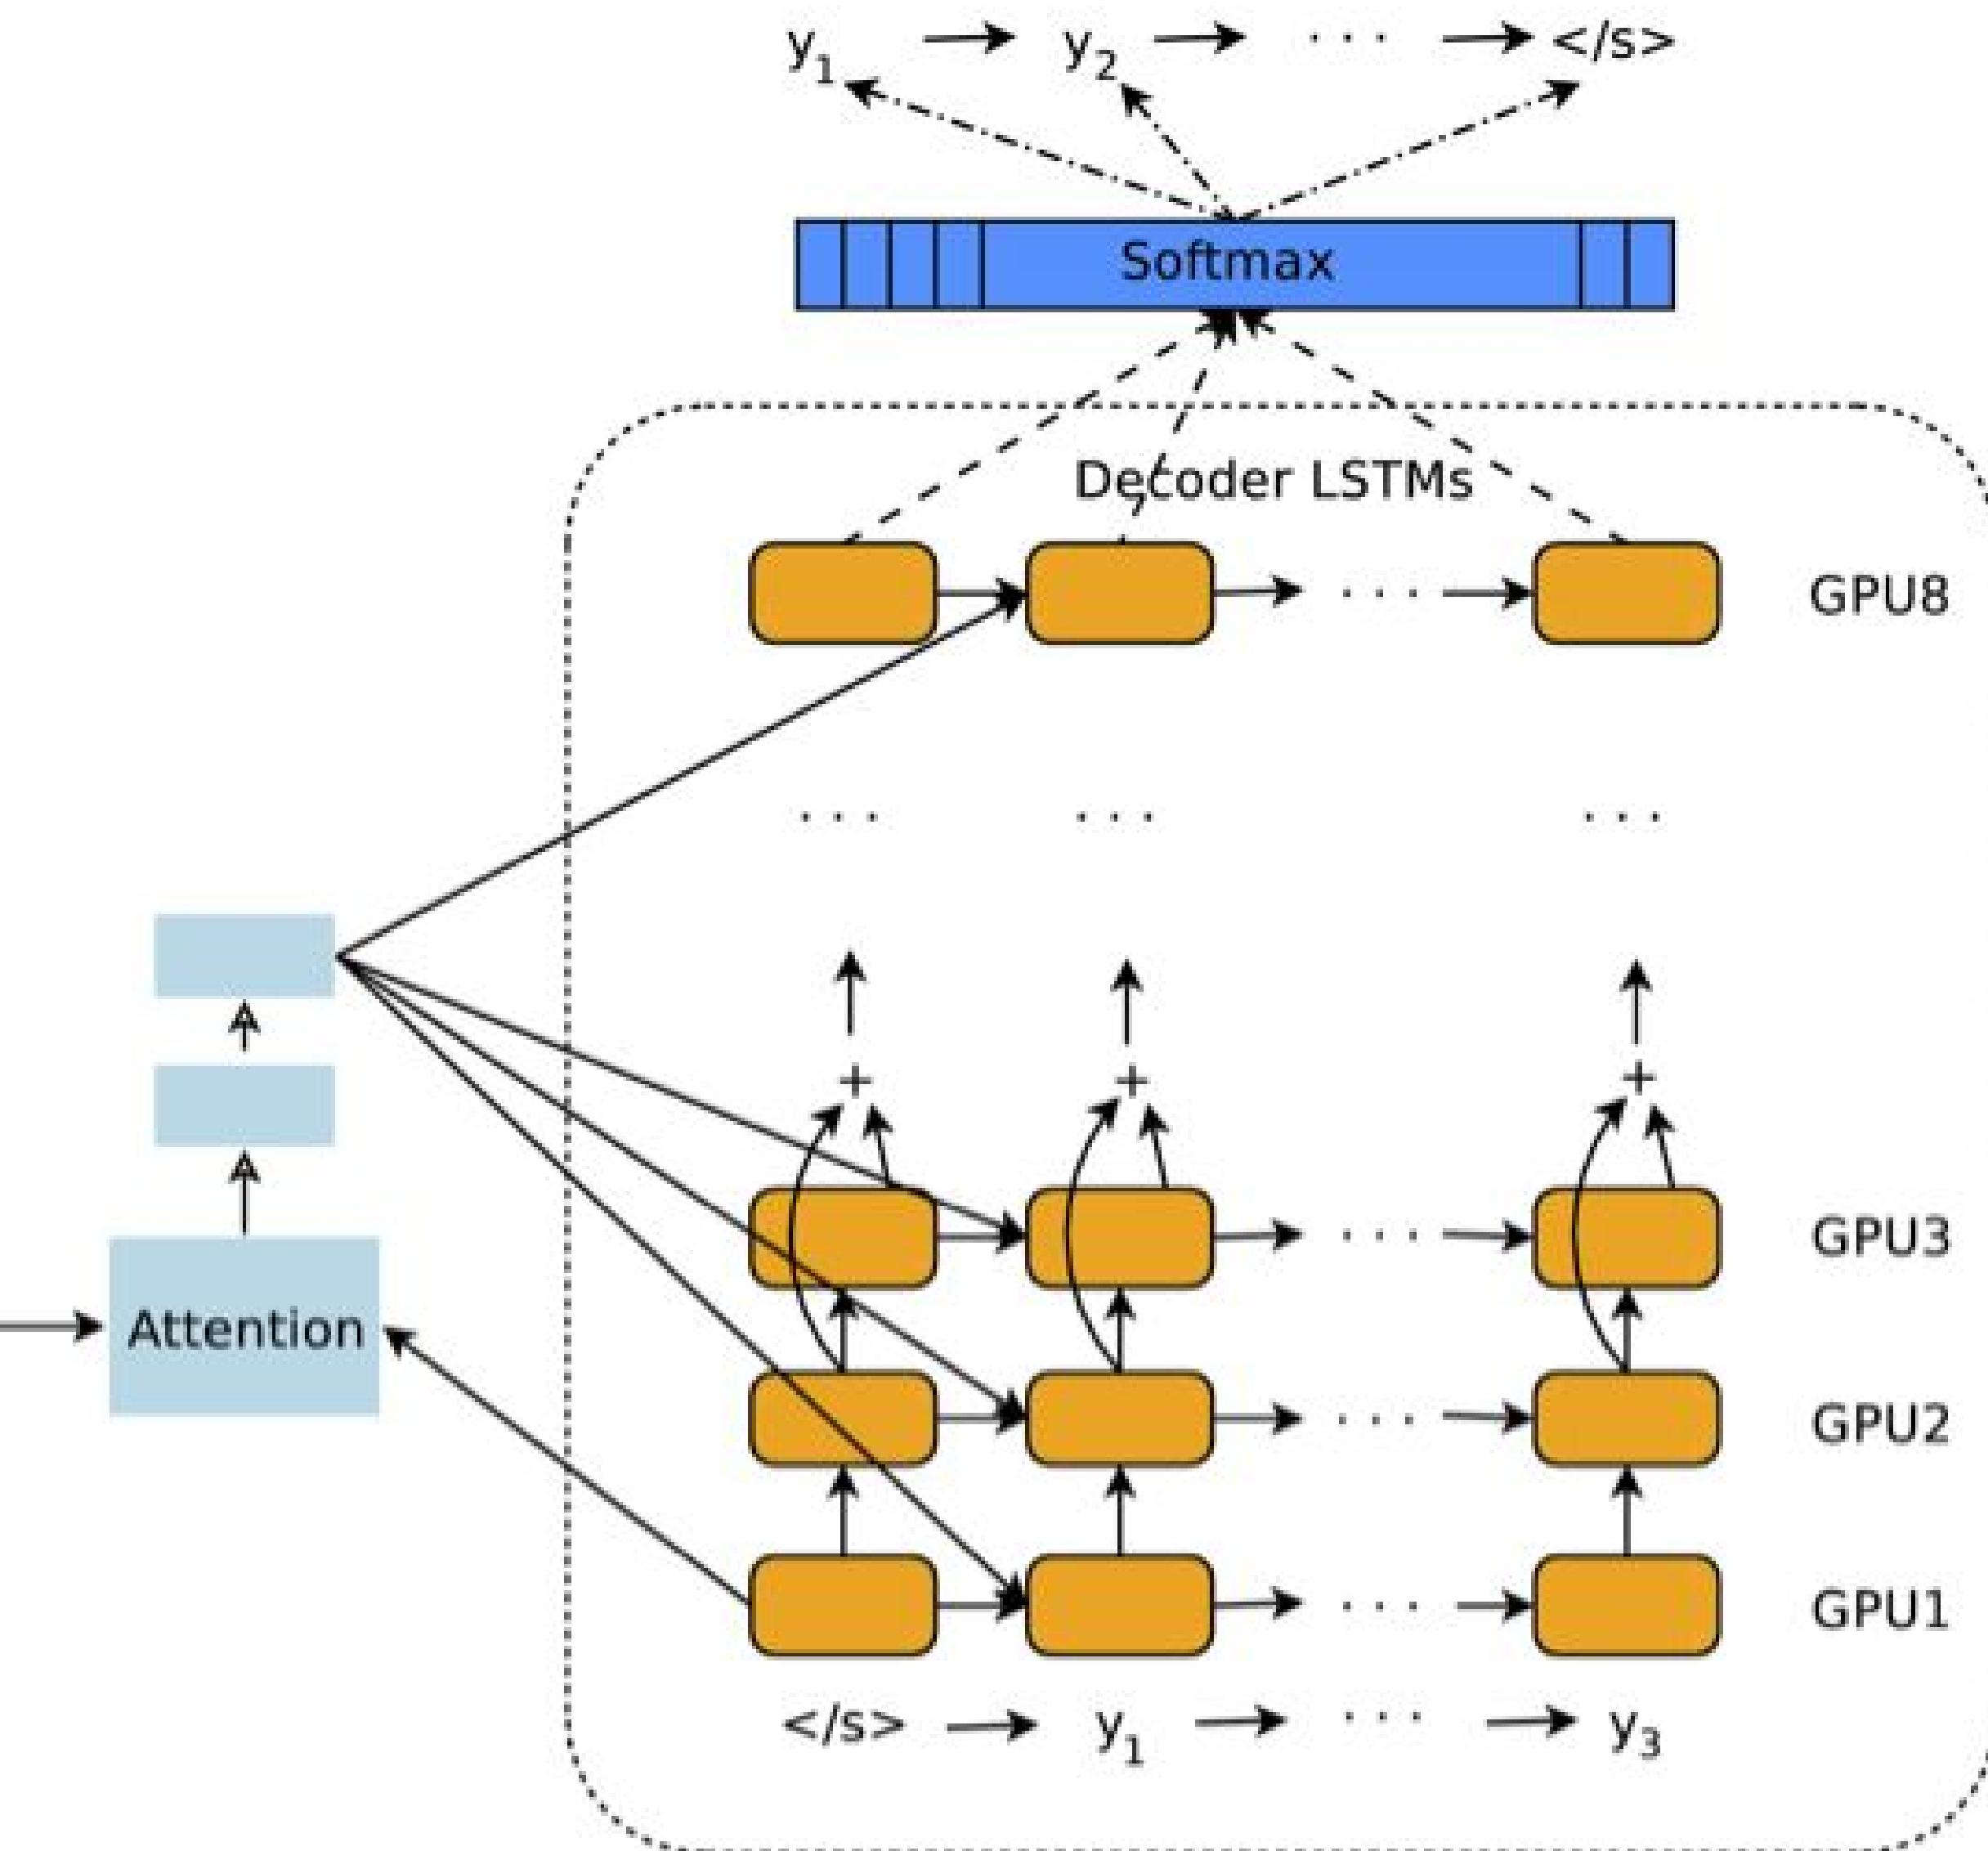
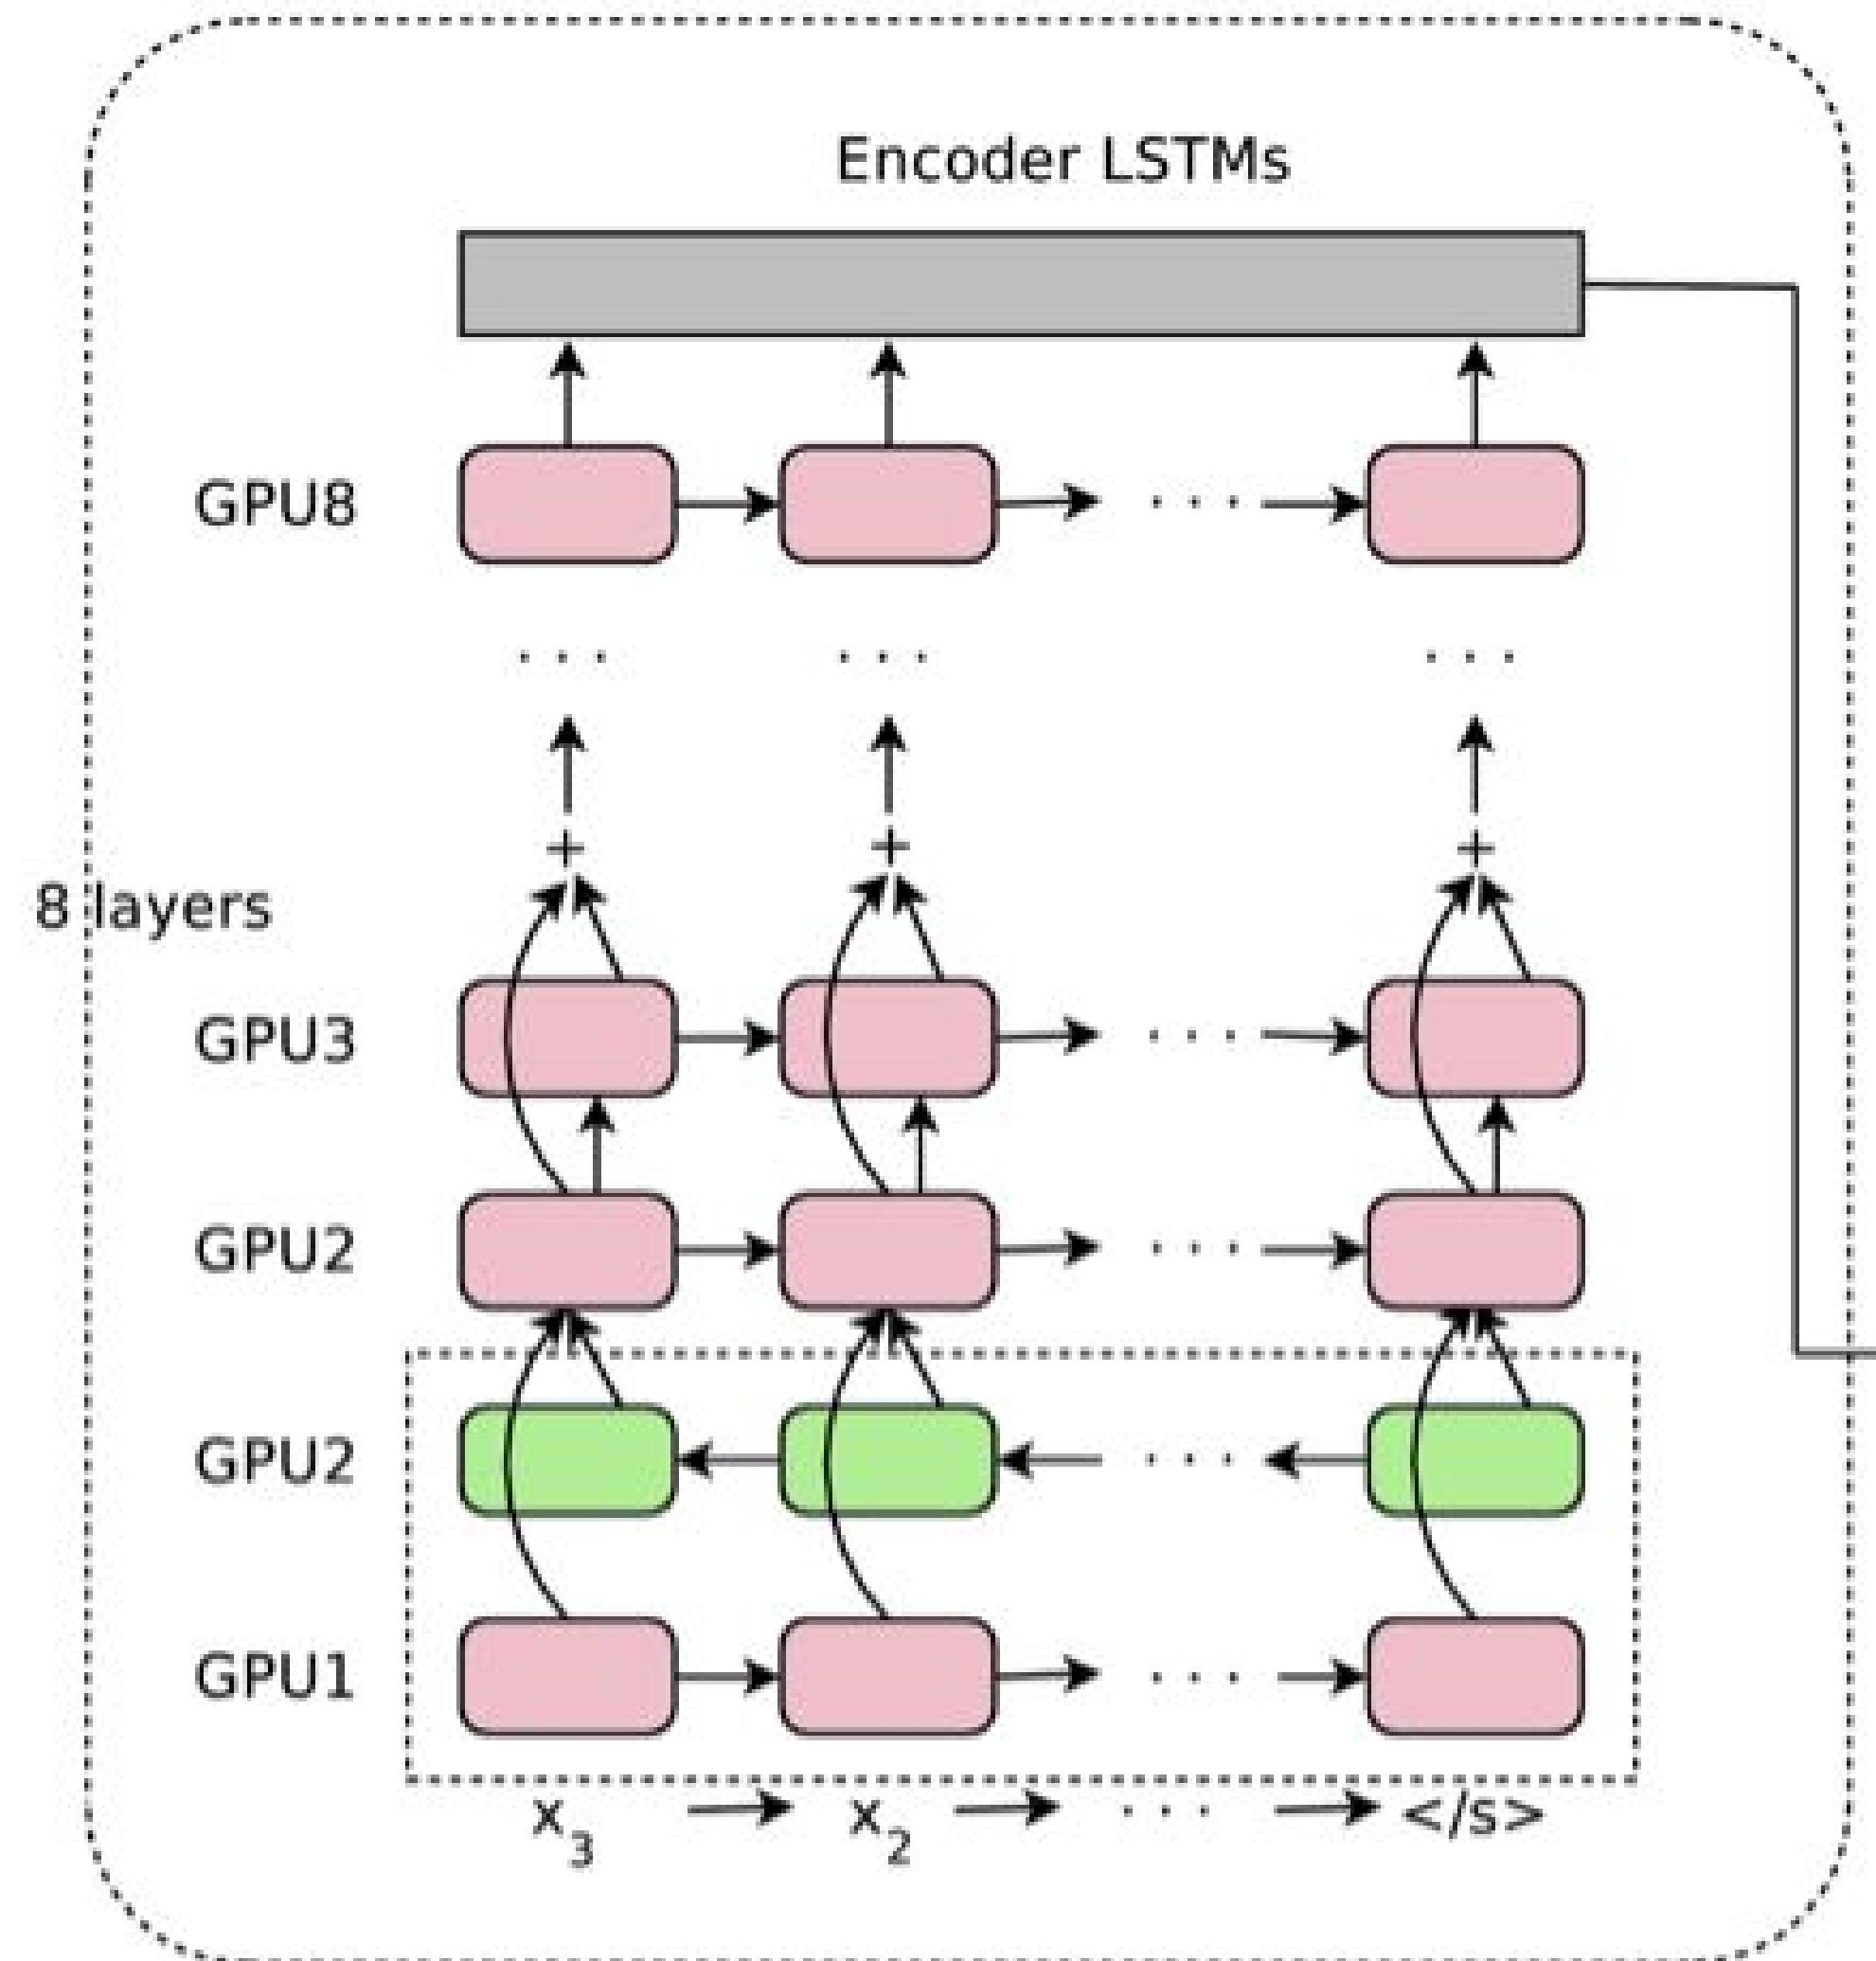
iOS



Android



Raspberry
Pi







Distributed out of the box



Distributed out of the box



Runs on CPU/GPU/TPU/mobile



Distributed out of the box



Runs on CPU/GPU/TPU/mobile



Fast and Flexible

Python Frontend

C++ Frontend

...

TensorFlow Distributed Execution Engine

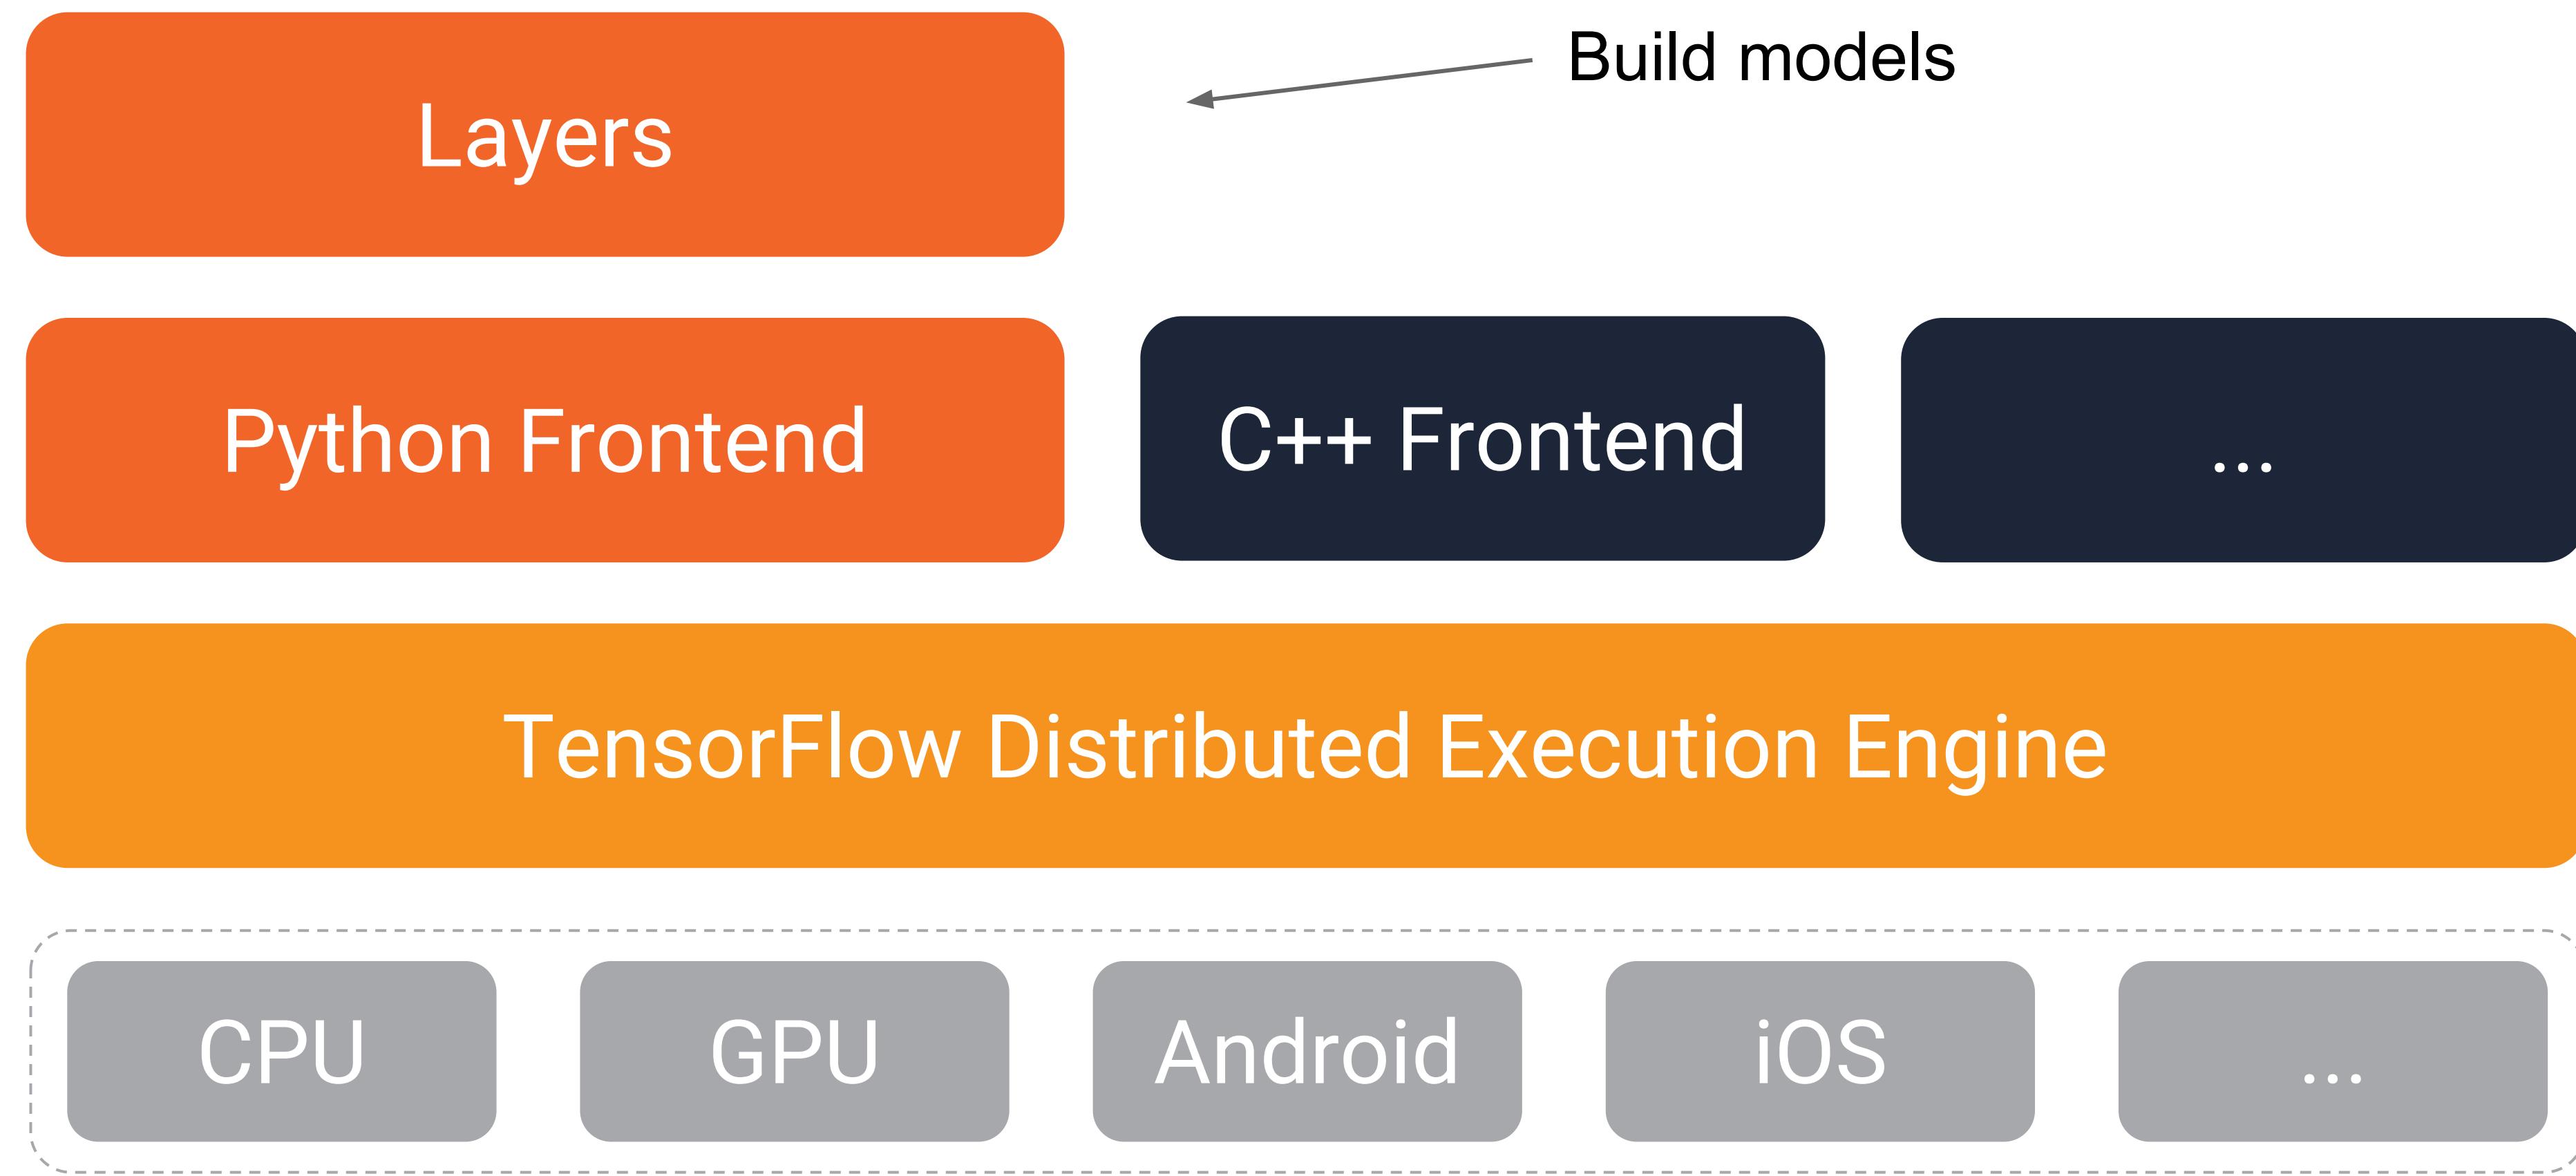
CPU

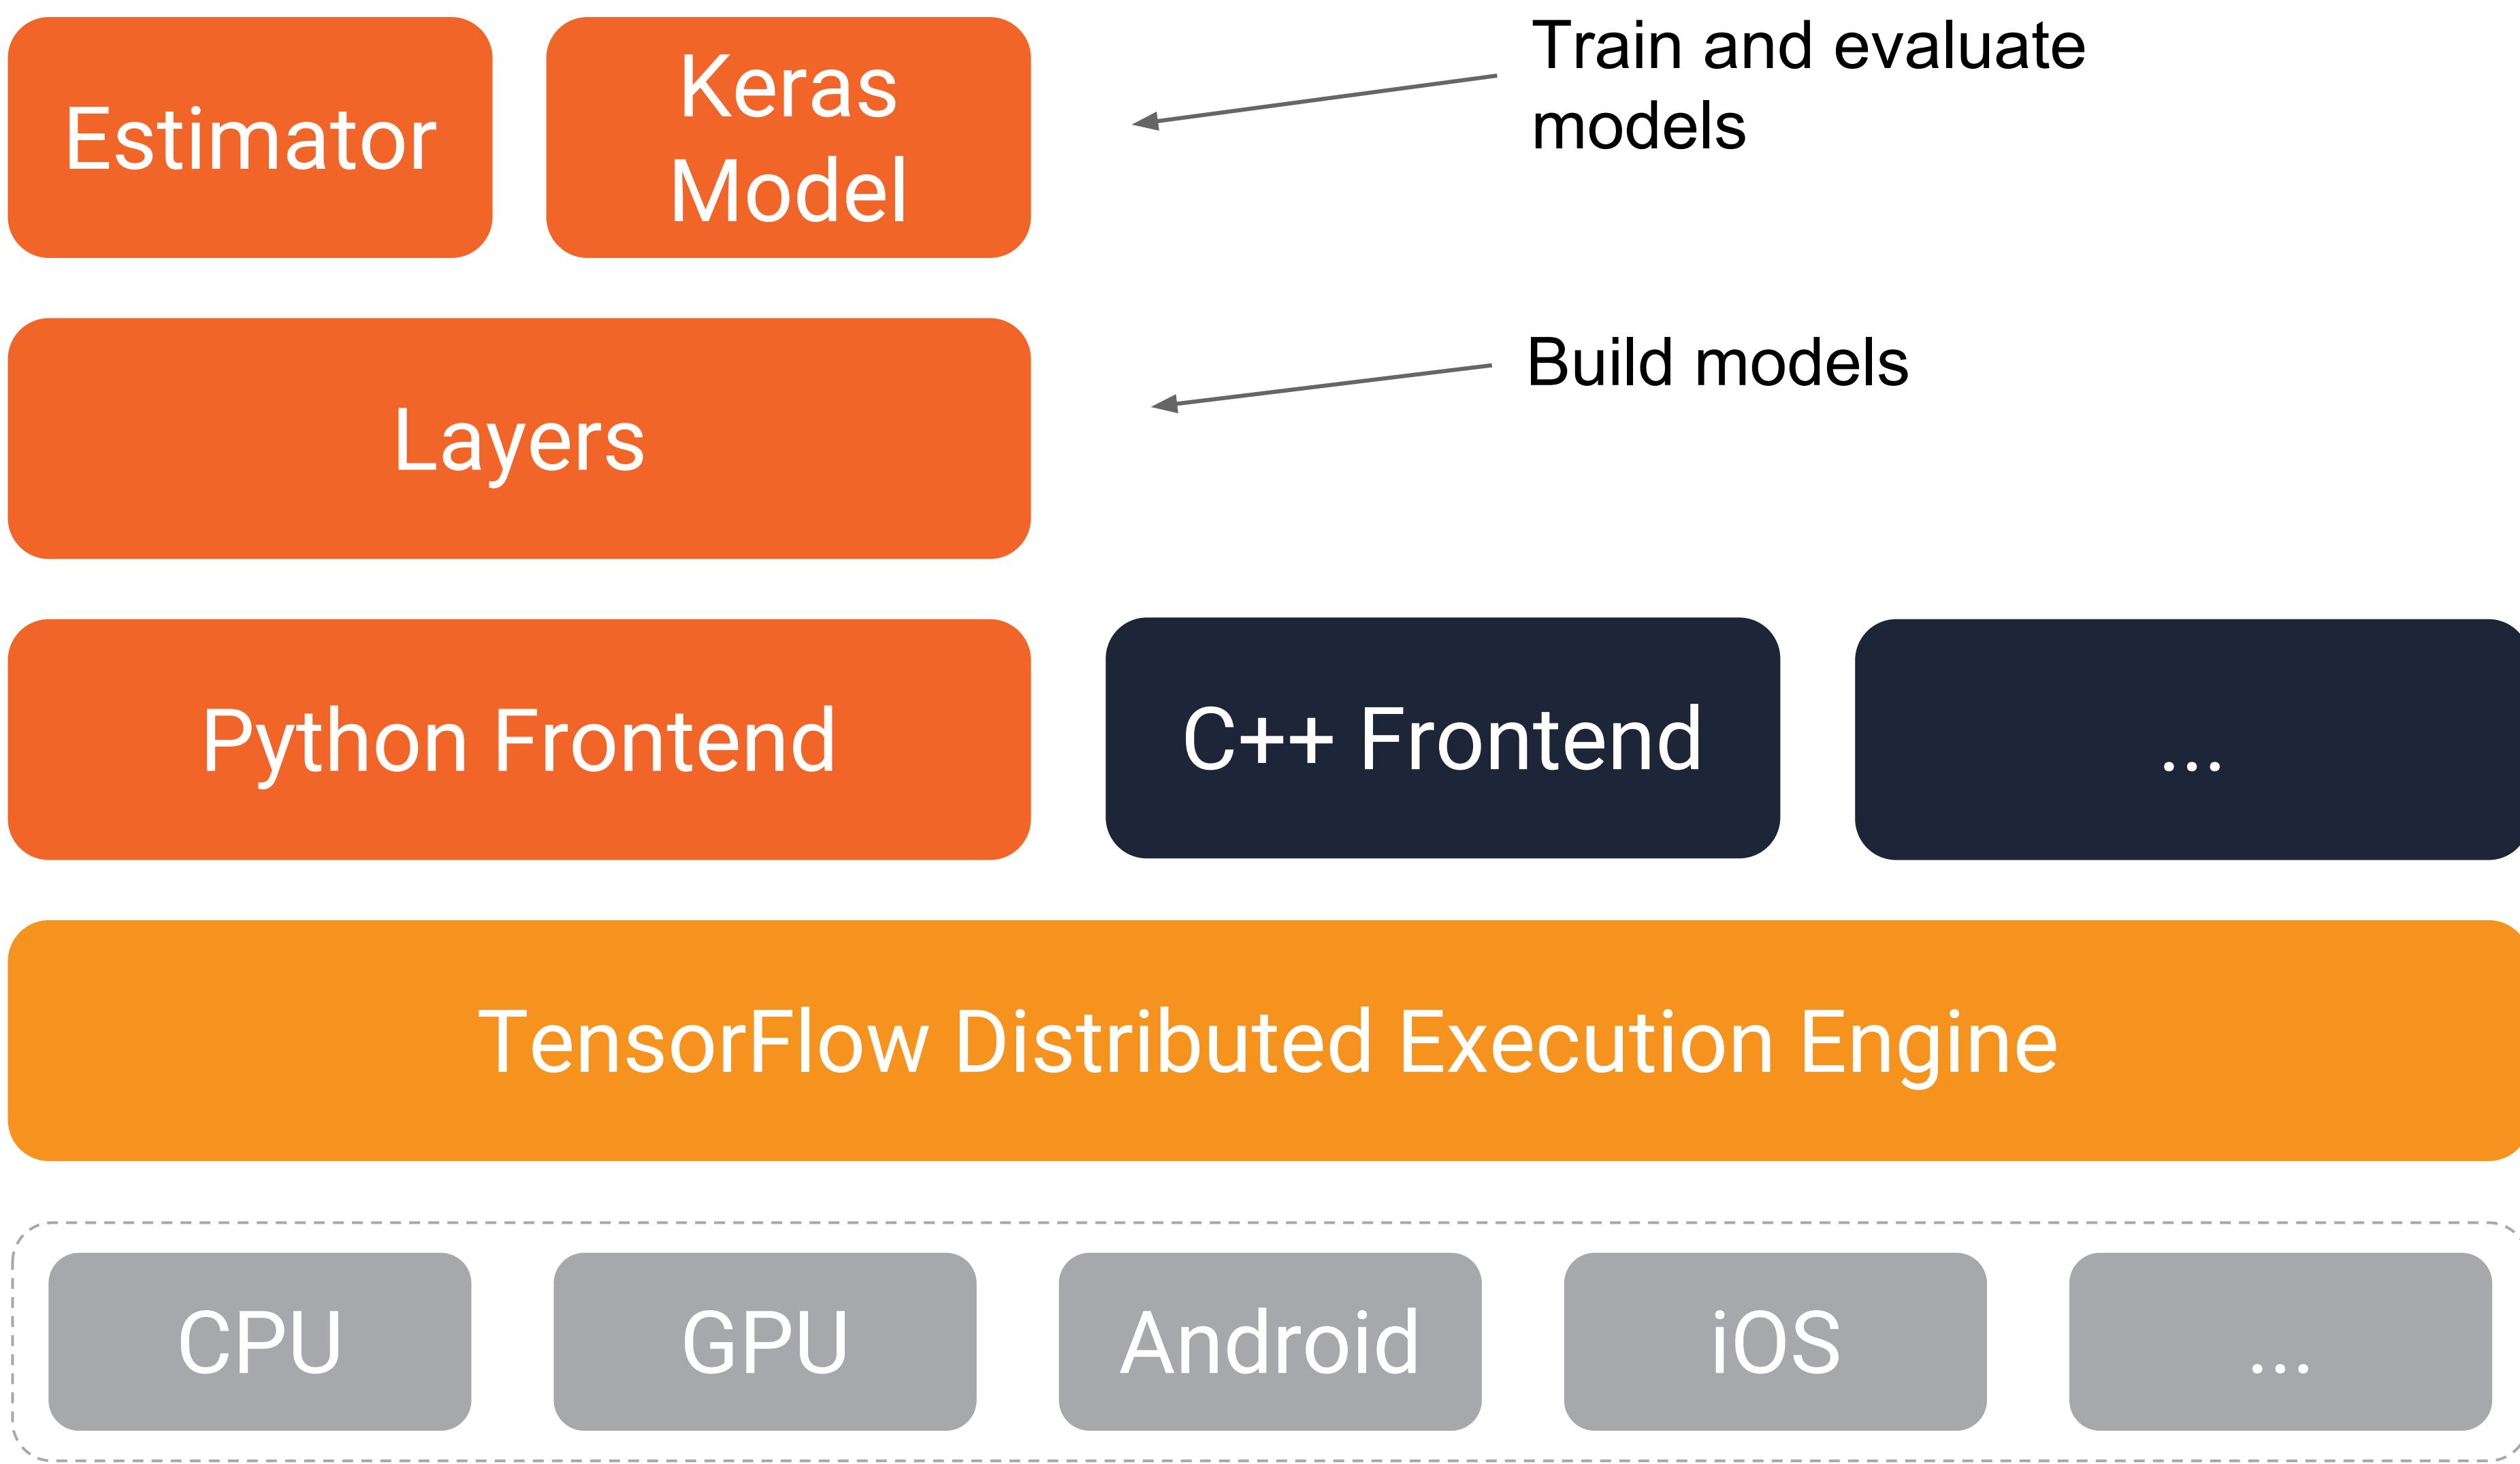
GPU

Android

iOS

...





Canned Estimators

Estimator

Keras Model

Layers

Python Frontend

C++ Frontend

...

TensorFlow Distributed Execution Engine

CPU

GPU

Android

iOS

...

Models in a box

Train and evaluate
models

Build models



My first Model
goo.gl/00gXiL


```
...  
regressor = LinearRegressor(feature_columns=[make, horsepower,  
                                              cylinders, ...])  
  
train_input_fn = pandas_input_fn(x=input_data, y=input_label,  
                                 batch_size=64, shuffle=True,  
                                 num_epochs=None)
```

```
...  
regressor = LinearRegressor(feature_columns=[make, horsepower,  
                                              cylinders, ...])  
  
train_input_fn = pandas_input_fn(x=input_data, y=input_label,  
                                 batch_size=64, shuffle=True,  
                                 num_epochs=None)  
  
regressor.train(train_input_fn, steps=10000)
```

INFO:tensorflow:global_step/sec: 113.22
INFO:tensorflow:loss = 6.07866e+07, step = 102 (0.884 sec)
INFO:tensorflow:global_step/sec: 131.905
INFO:tensorflow:loss = 7.26533e+07, step = 202 (0.758 sec)
INFO:tensorflow:global_step/sec: 125.879
INFO:tensorflow:loss = 3.64895e+07, step = 302 (0.794 sec)
INFO:tensorflow:global_step/sec: 131.558
INFO:tensorflow:loss = 4.84076e+07, step = 402 (0.760 sec)
INFO:tensorflow:global_step/sec: 124.394
INFO:tensorflow:loss = 4.316e+07, step = 502 (0.804 sec)
INFO:tensorflow:global_step/sec: 126.16
INFO:tensorflow:loss = 2.32498e+07, step = 602 (0.793 sec)
INFO:tensorflow:global_step/sec: 122.751
INFO:tensorflow:loss = 4.53009e+07, step = 702 (0.817 sec)
INFO:tensorflow:global_step/sec: 118.297
INFO:tensorflow:loss = 4.56962e+07, step = 802 (0.843 sec)
INFO:tensorflow:global_step/sec: 126.641
INFO:tensorflow:loss = 3.3779e+07, step = 902 (0.790 sec)
INFO:tensorflow:global_step/sec: 109.258
INFO:tensorflow:loss = 5.32131e+07, step = 1002 (0.915 sec)
INFO:tensorflow:global_step/sec: 135.997
INFO:tensorflow:loss = 3.2445e+07, step = 1102 (0.735 sec)
INFO:tensorflow:global_step/sec: 123.69
INFO:tensorflow:loss = 2.43225e+07, step = 1202 (0.811 sec)
INFO:tensorflow:global_step/sec: 136.363
INFO:tensorflow:loss = 2.86487e+07, step = 1302 (0.730 sec)
INFO:tensorflow:global_step/sec: 129.674

Write a regex to create a tag group X

Show data download links

Ignore outliers in chart scaling

Tooltip sorting method: default

Smoothing



Horizontal Axis

STEP RELATIVE WALL

Runs

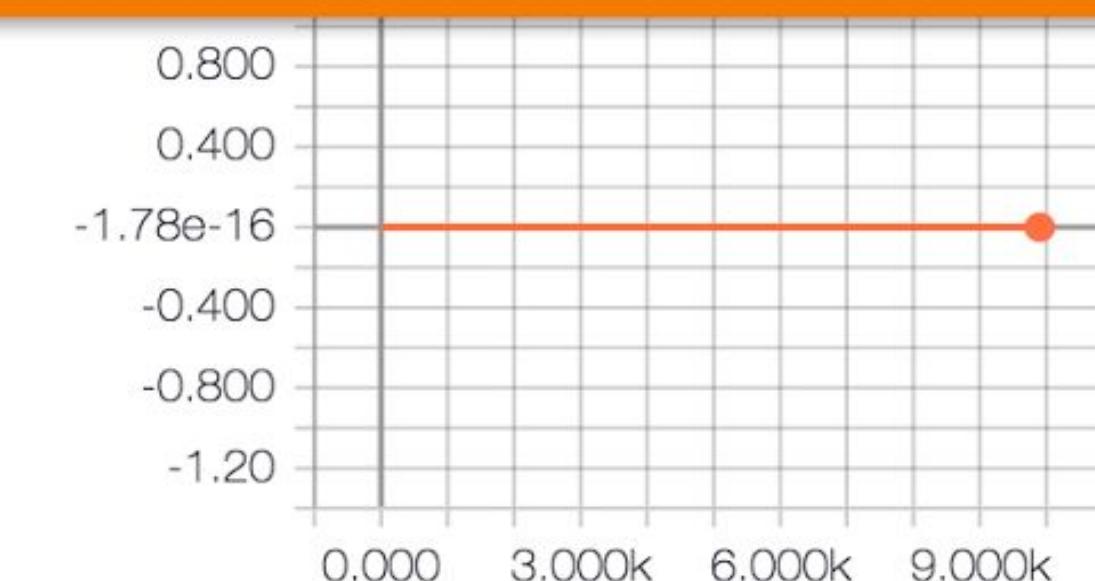
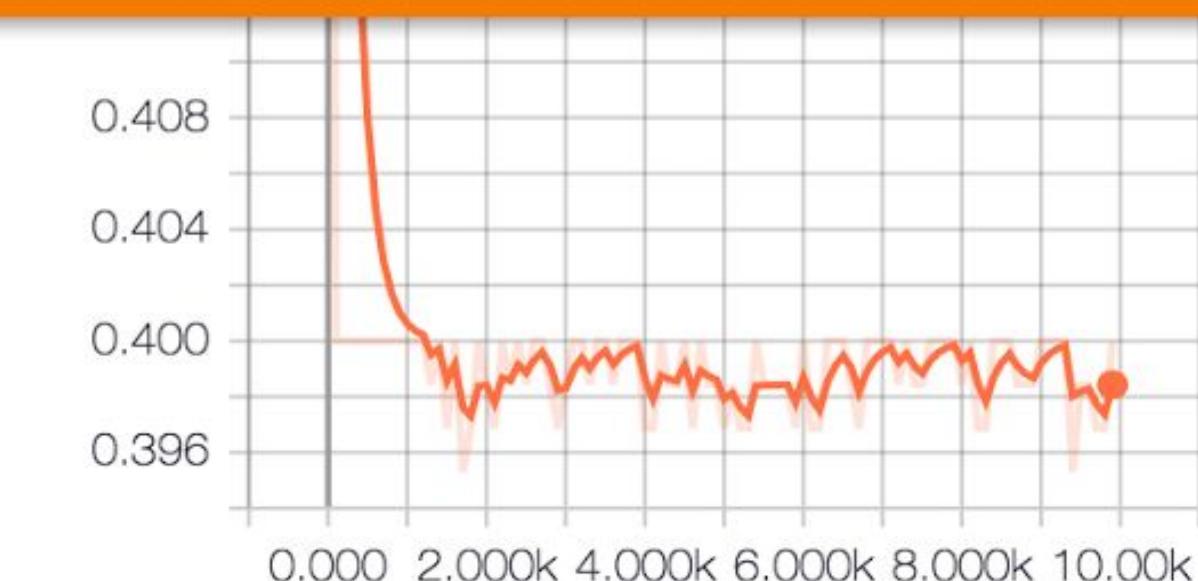
Write a regex to filter runs

○

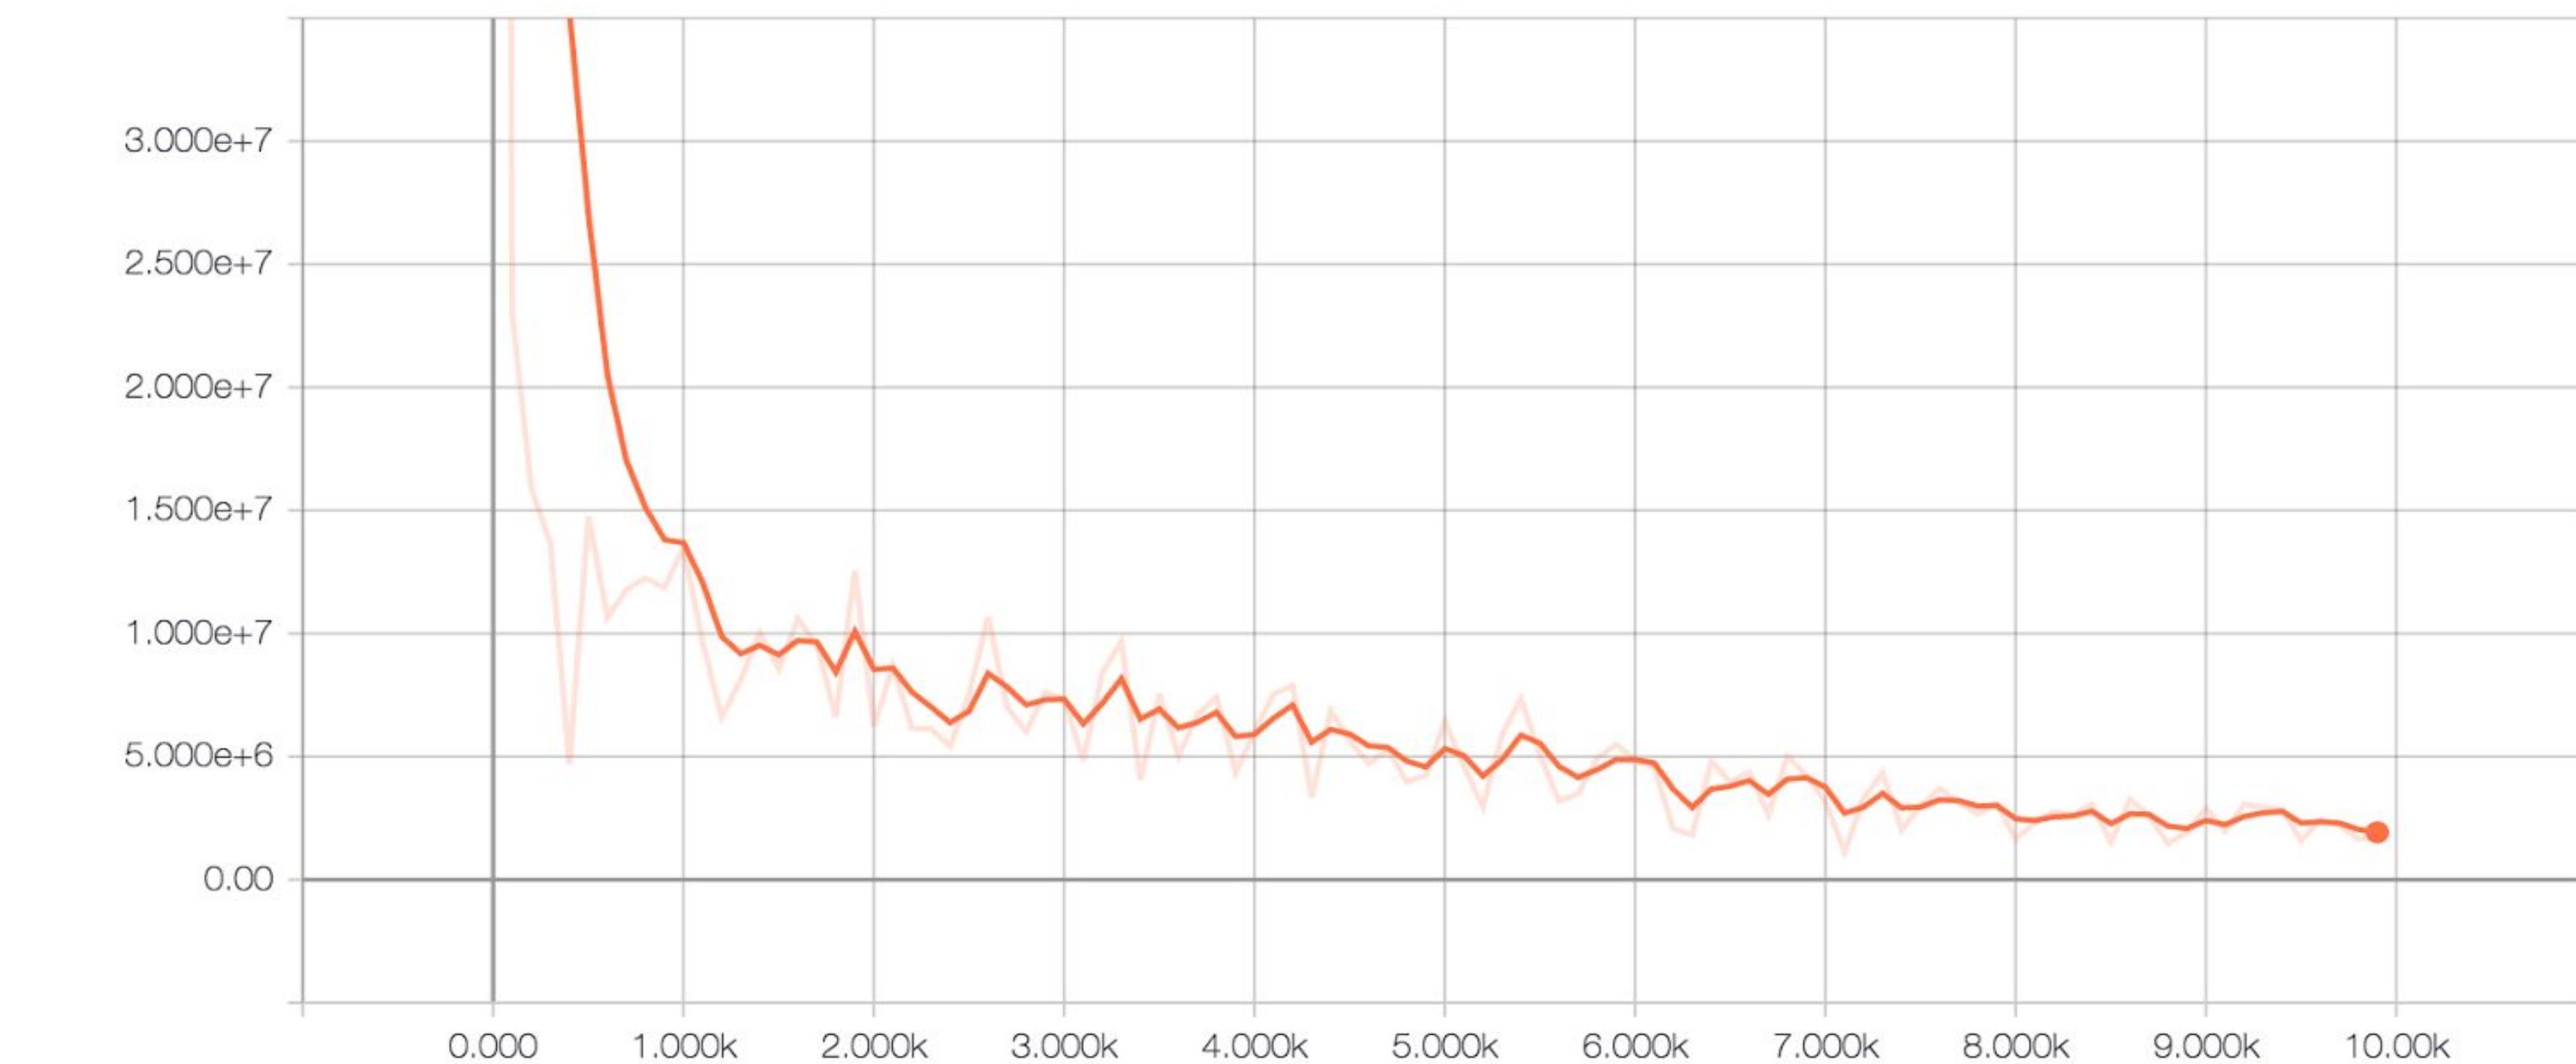
○ eval

TOGGLE ALL RUNS

/var/folders/qr/4w8_972d0w17wlh9xnly3y800_glr/T/tmpvSZZhV



dnn/regression_head/loss



enqueue_input

1

Fit to screen
 Download PNG
Run (2)

Session runs (0)

Upload

Trace inputs

Color Structure

Device

XLA Cluster

Compute time

Memory

colors same substructure

unique substructure

Graph (* = expandable)

Namespace*

OpNode

Unconnected series*

Connected series*

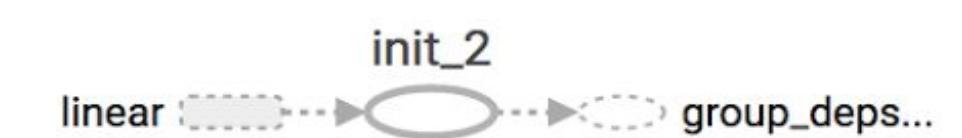
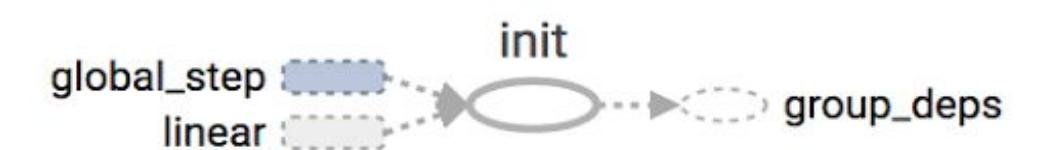
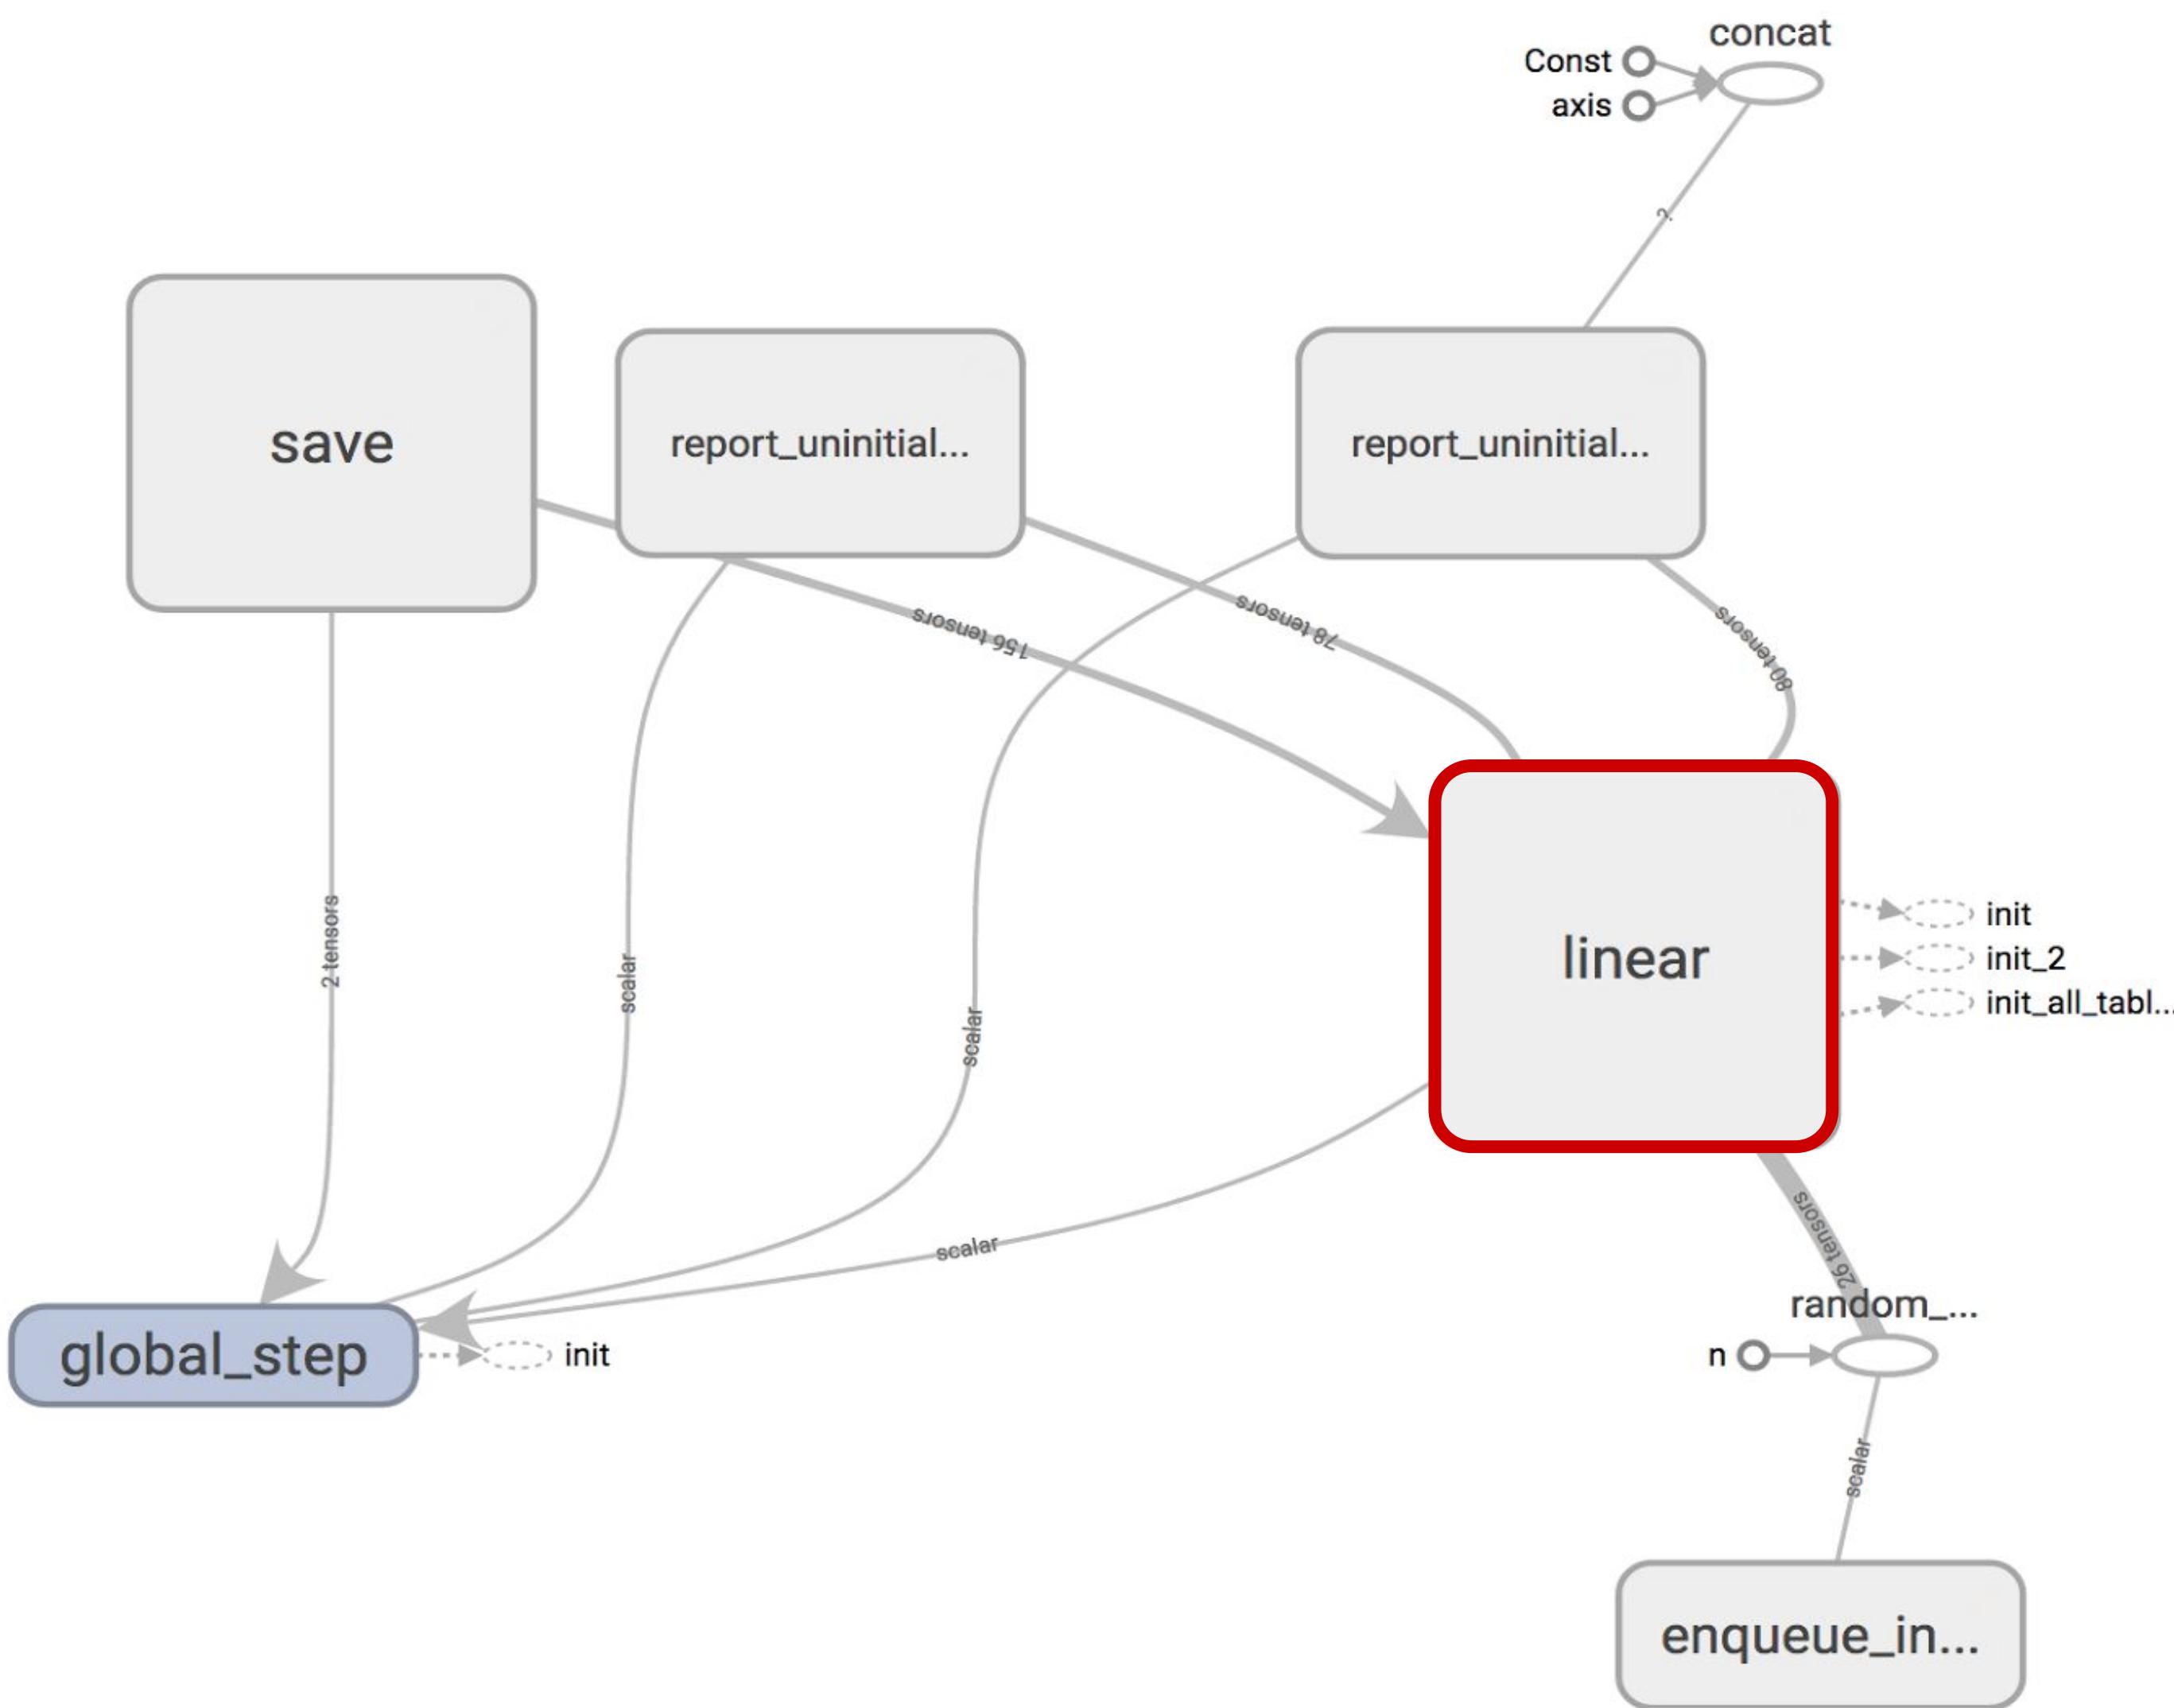
Constant

Summary

Dataflow edge

Control dependency edge

Reference edge





Fit to screen



Download PNG

Run

(2)

Session

runs (0)

Upload

Choose File

Trace inputs

Color Structure Device XLA Cluster Compute time Memory

colors same substructure

unique substructure

Graph (* = expandable)



Namespace*



OpNode



Unconnected series*



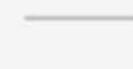
Connected series*



Constant



Summary



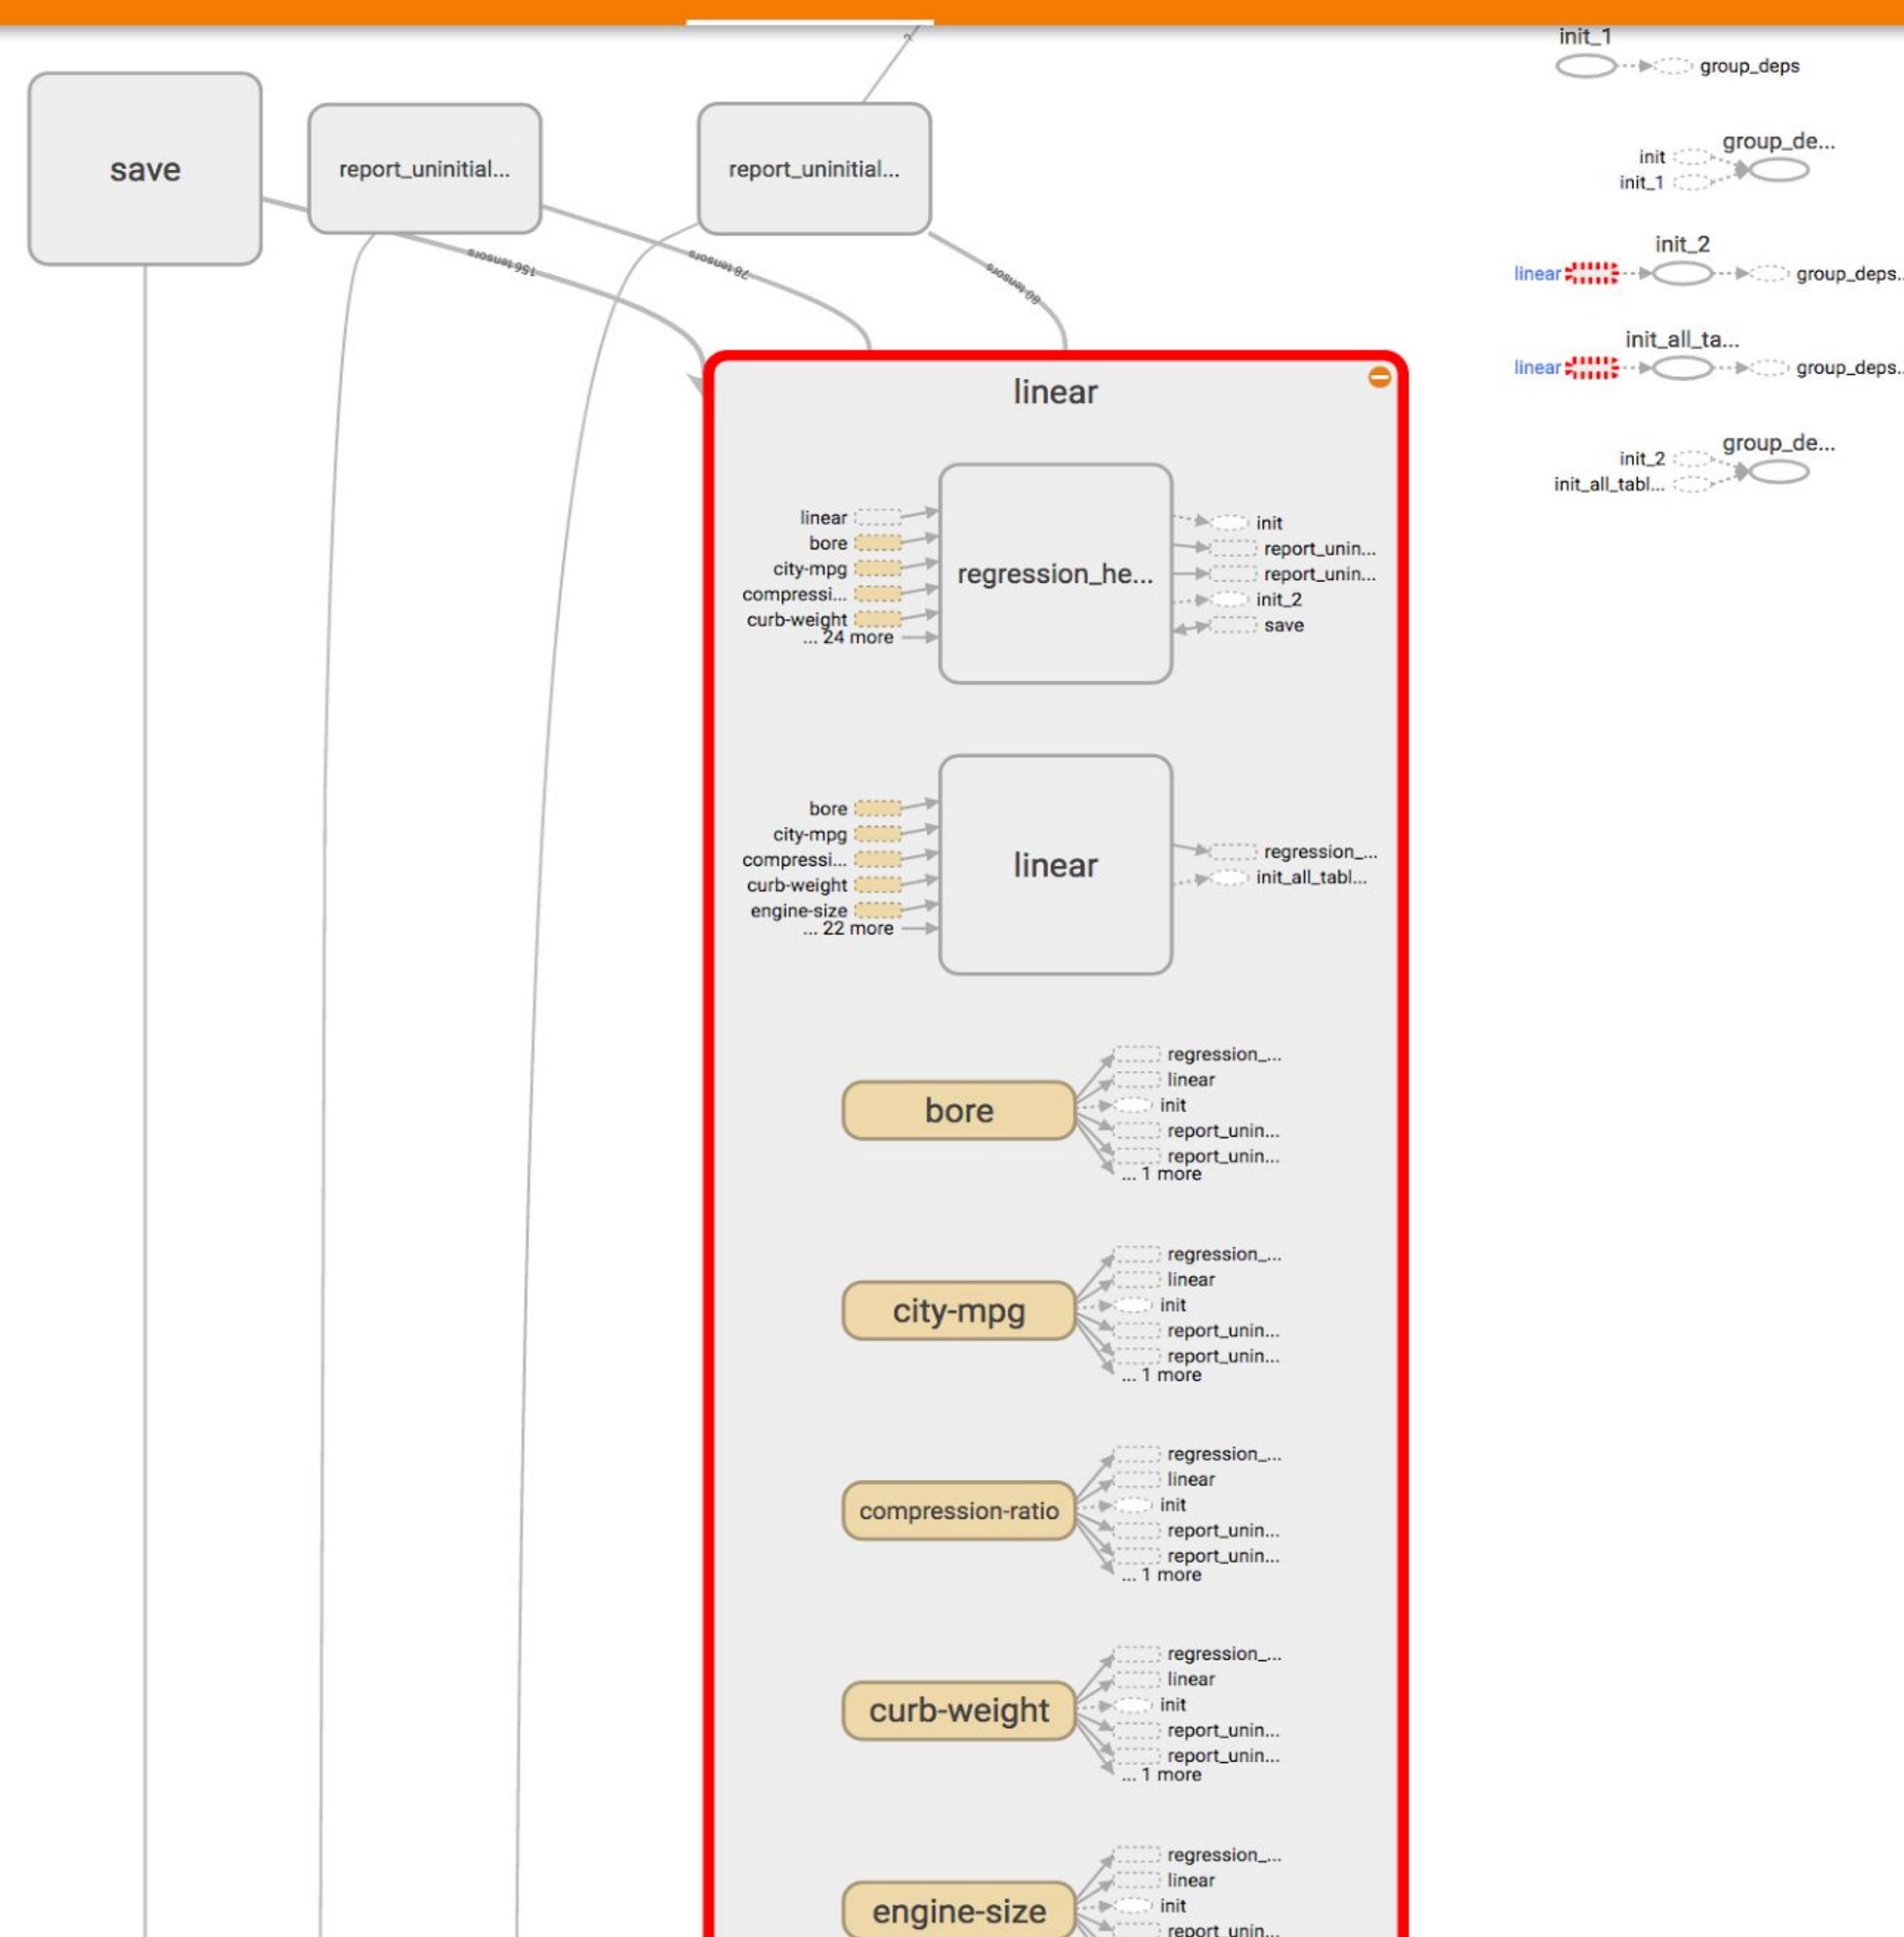
Dataflow edge



Control dependency edge

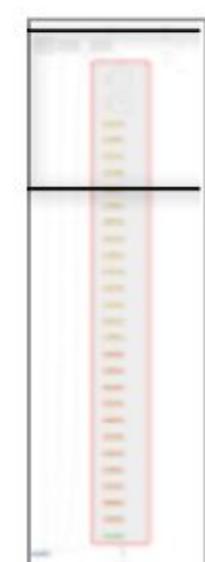


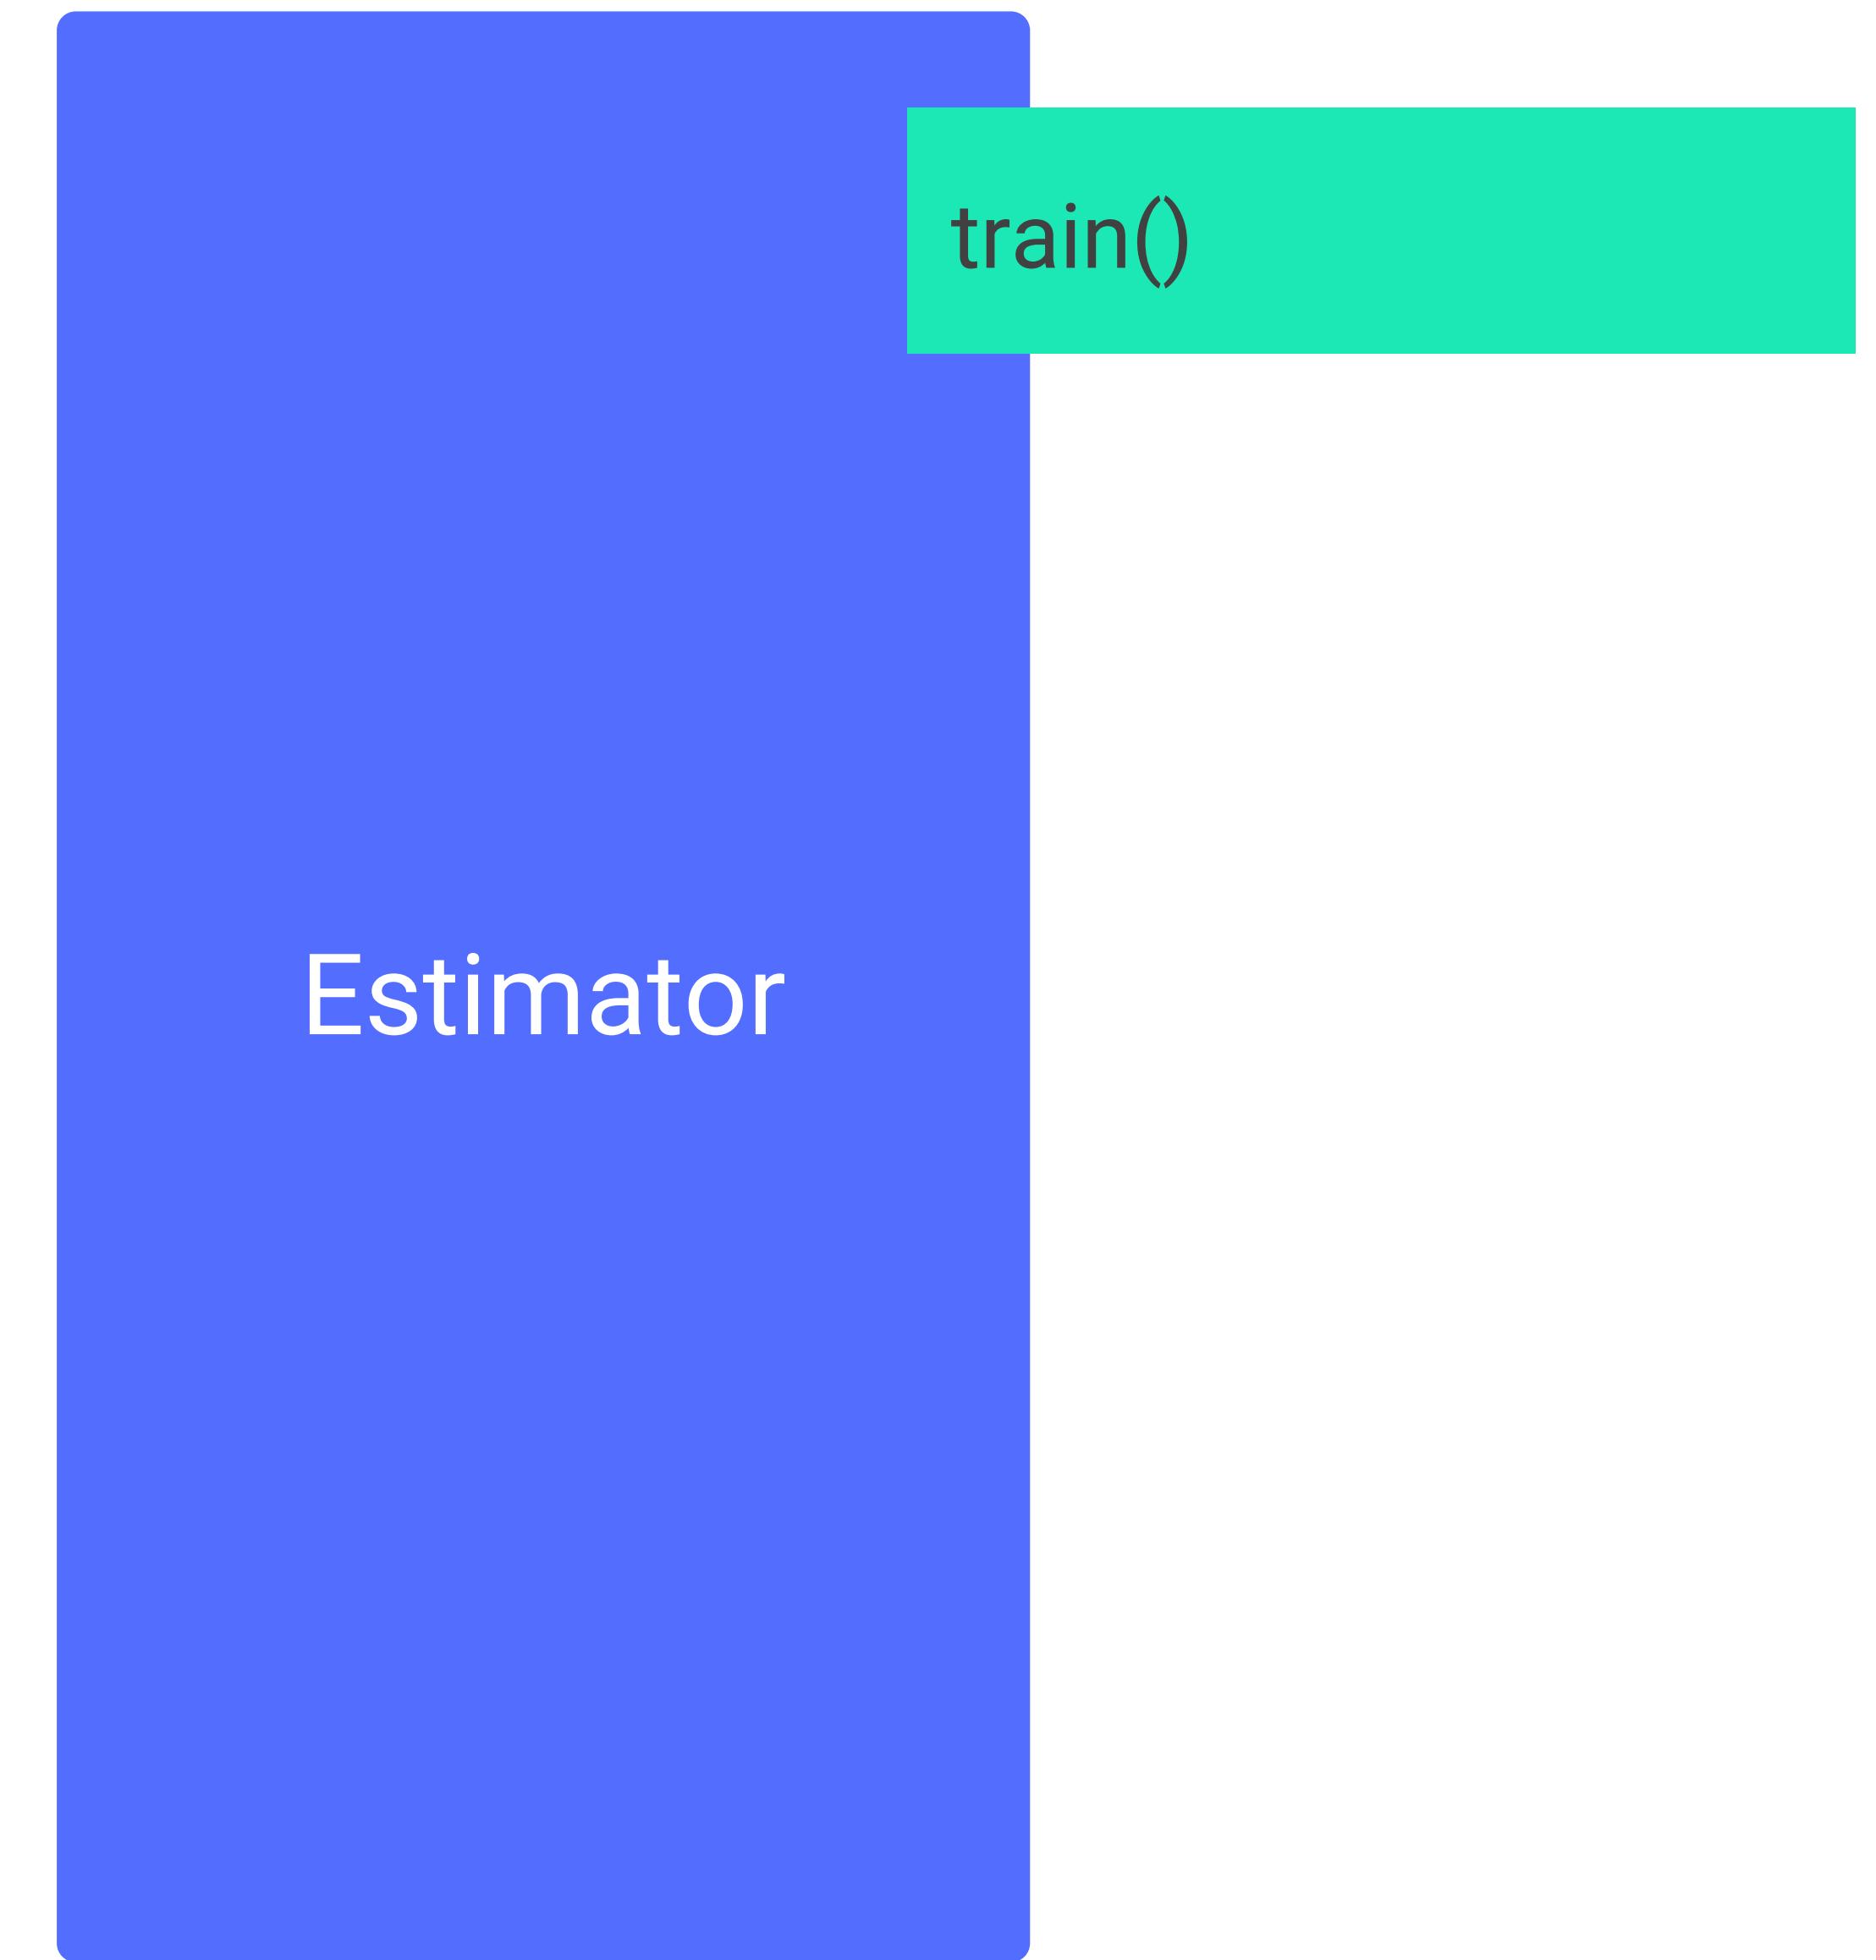
Reference edge

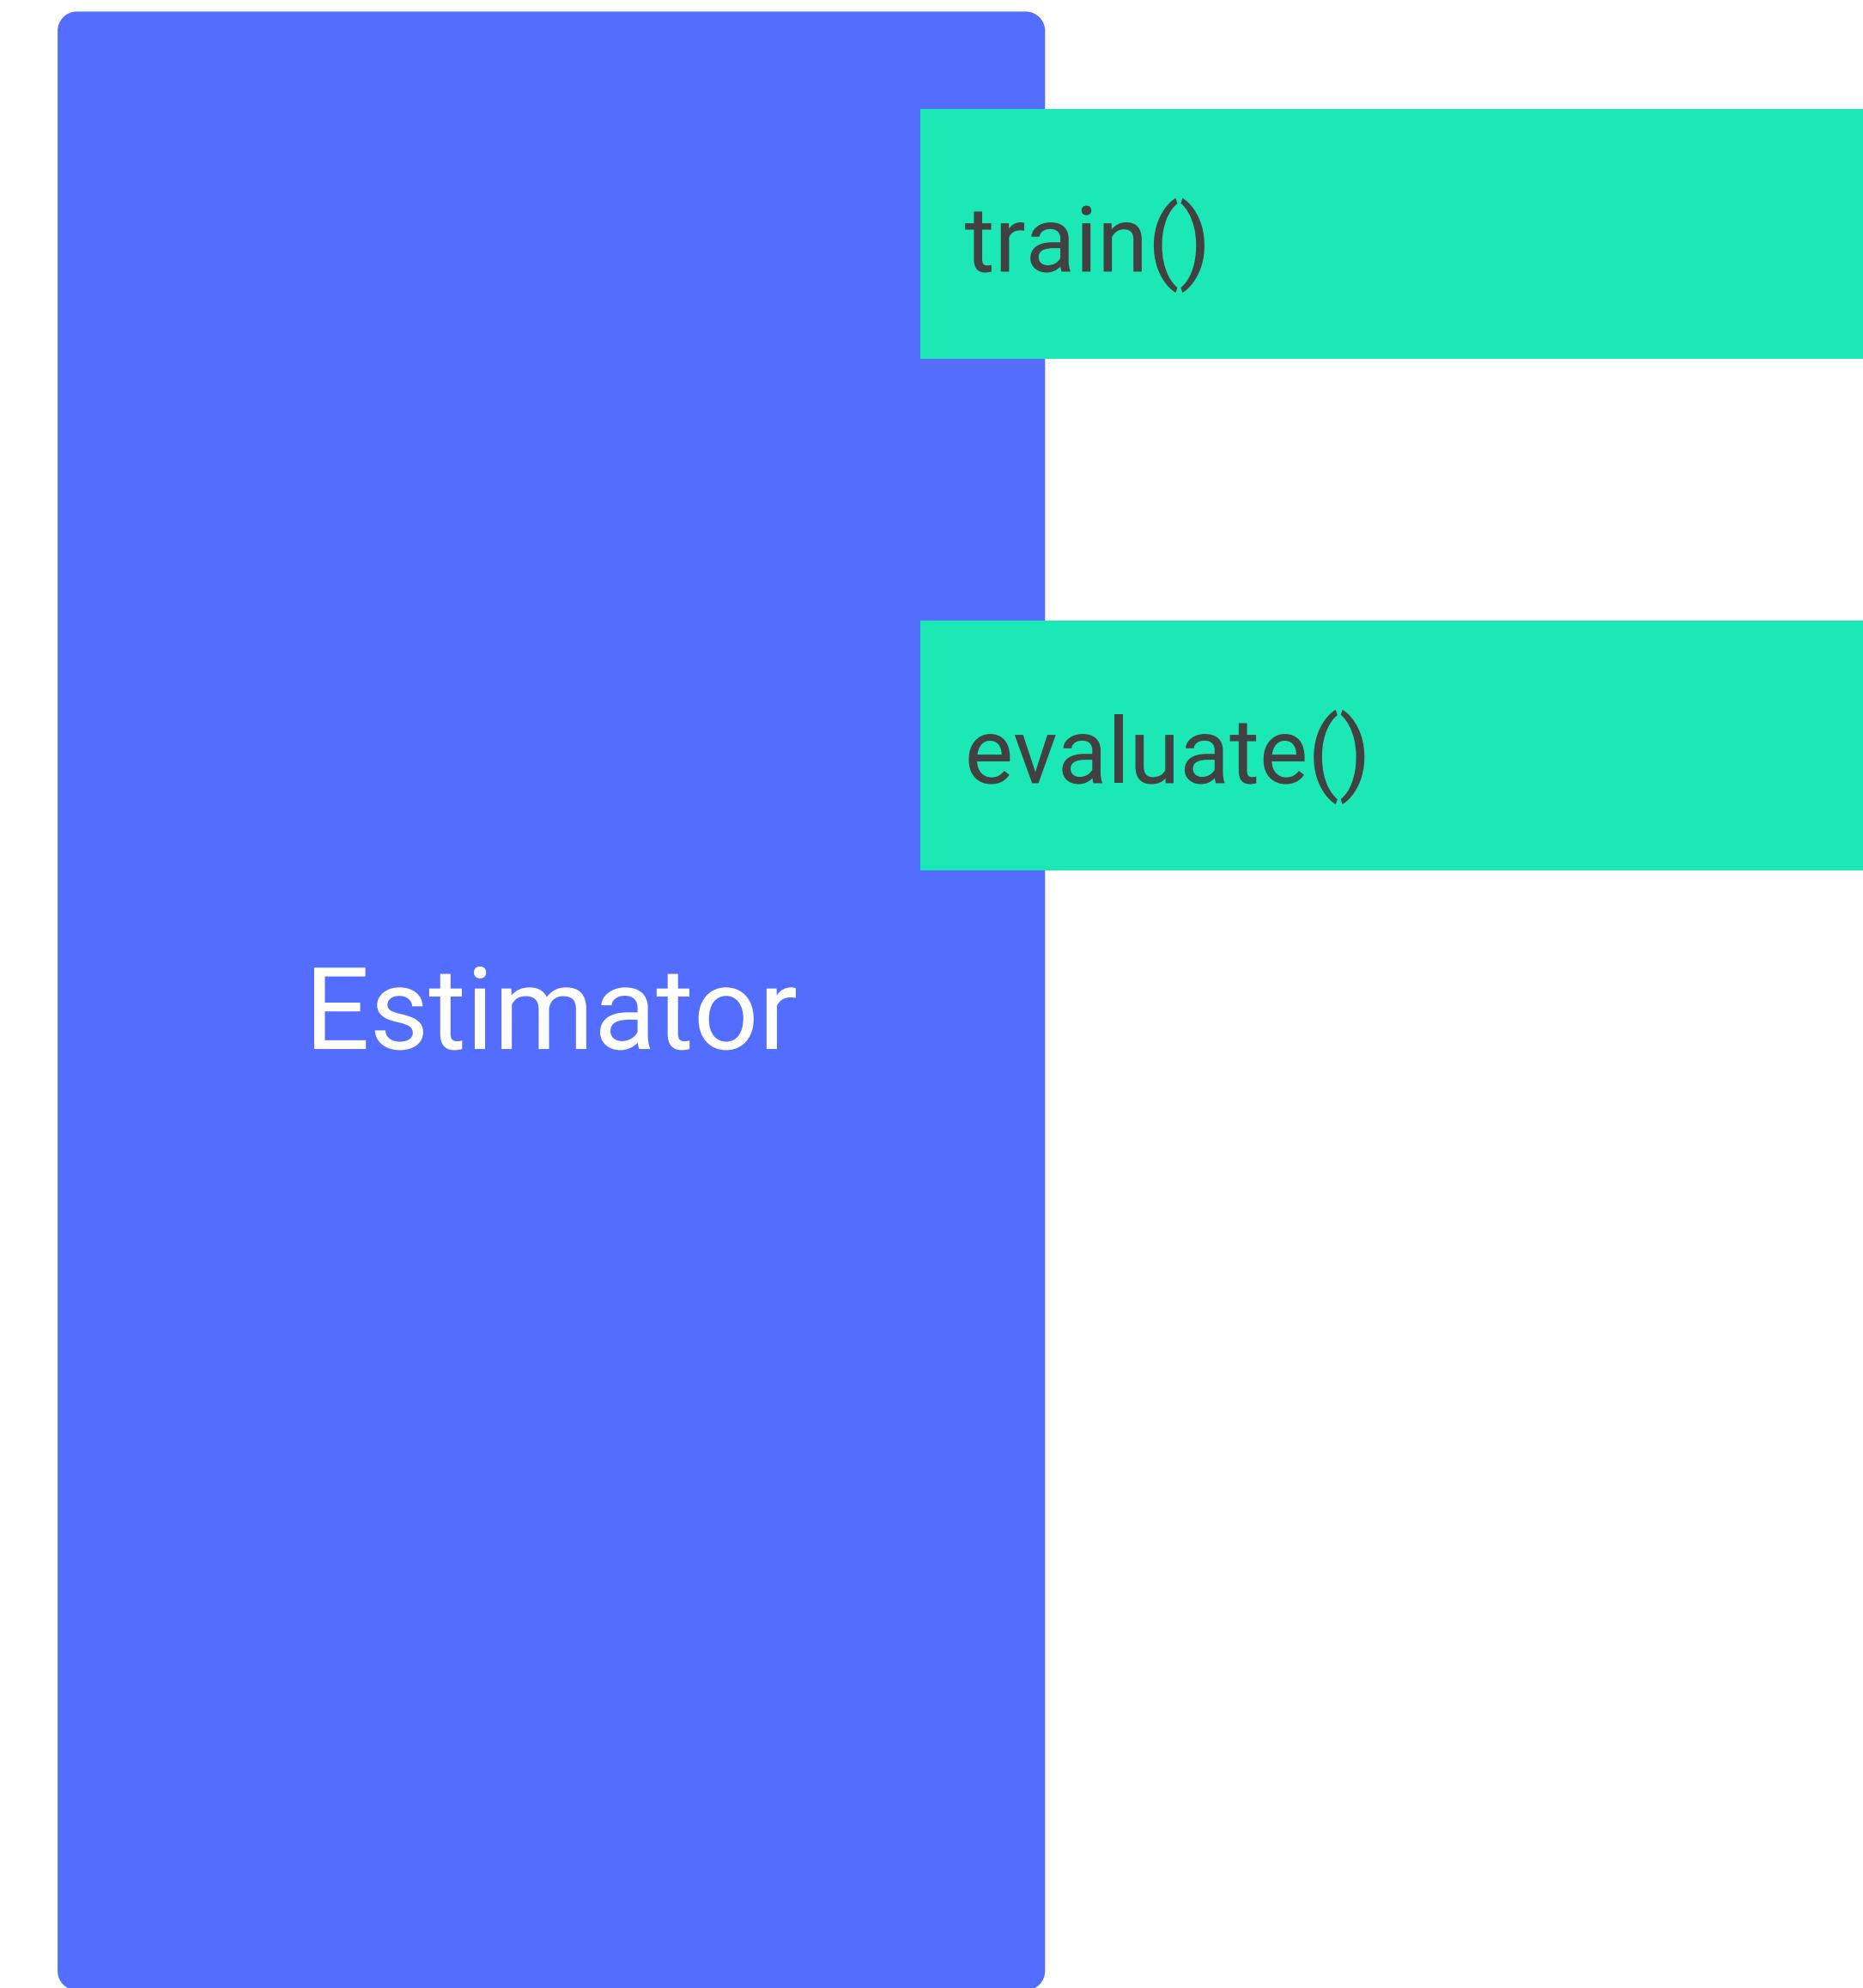


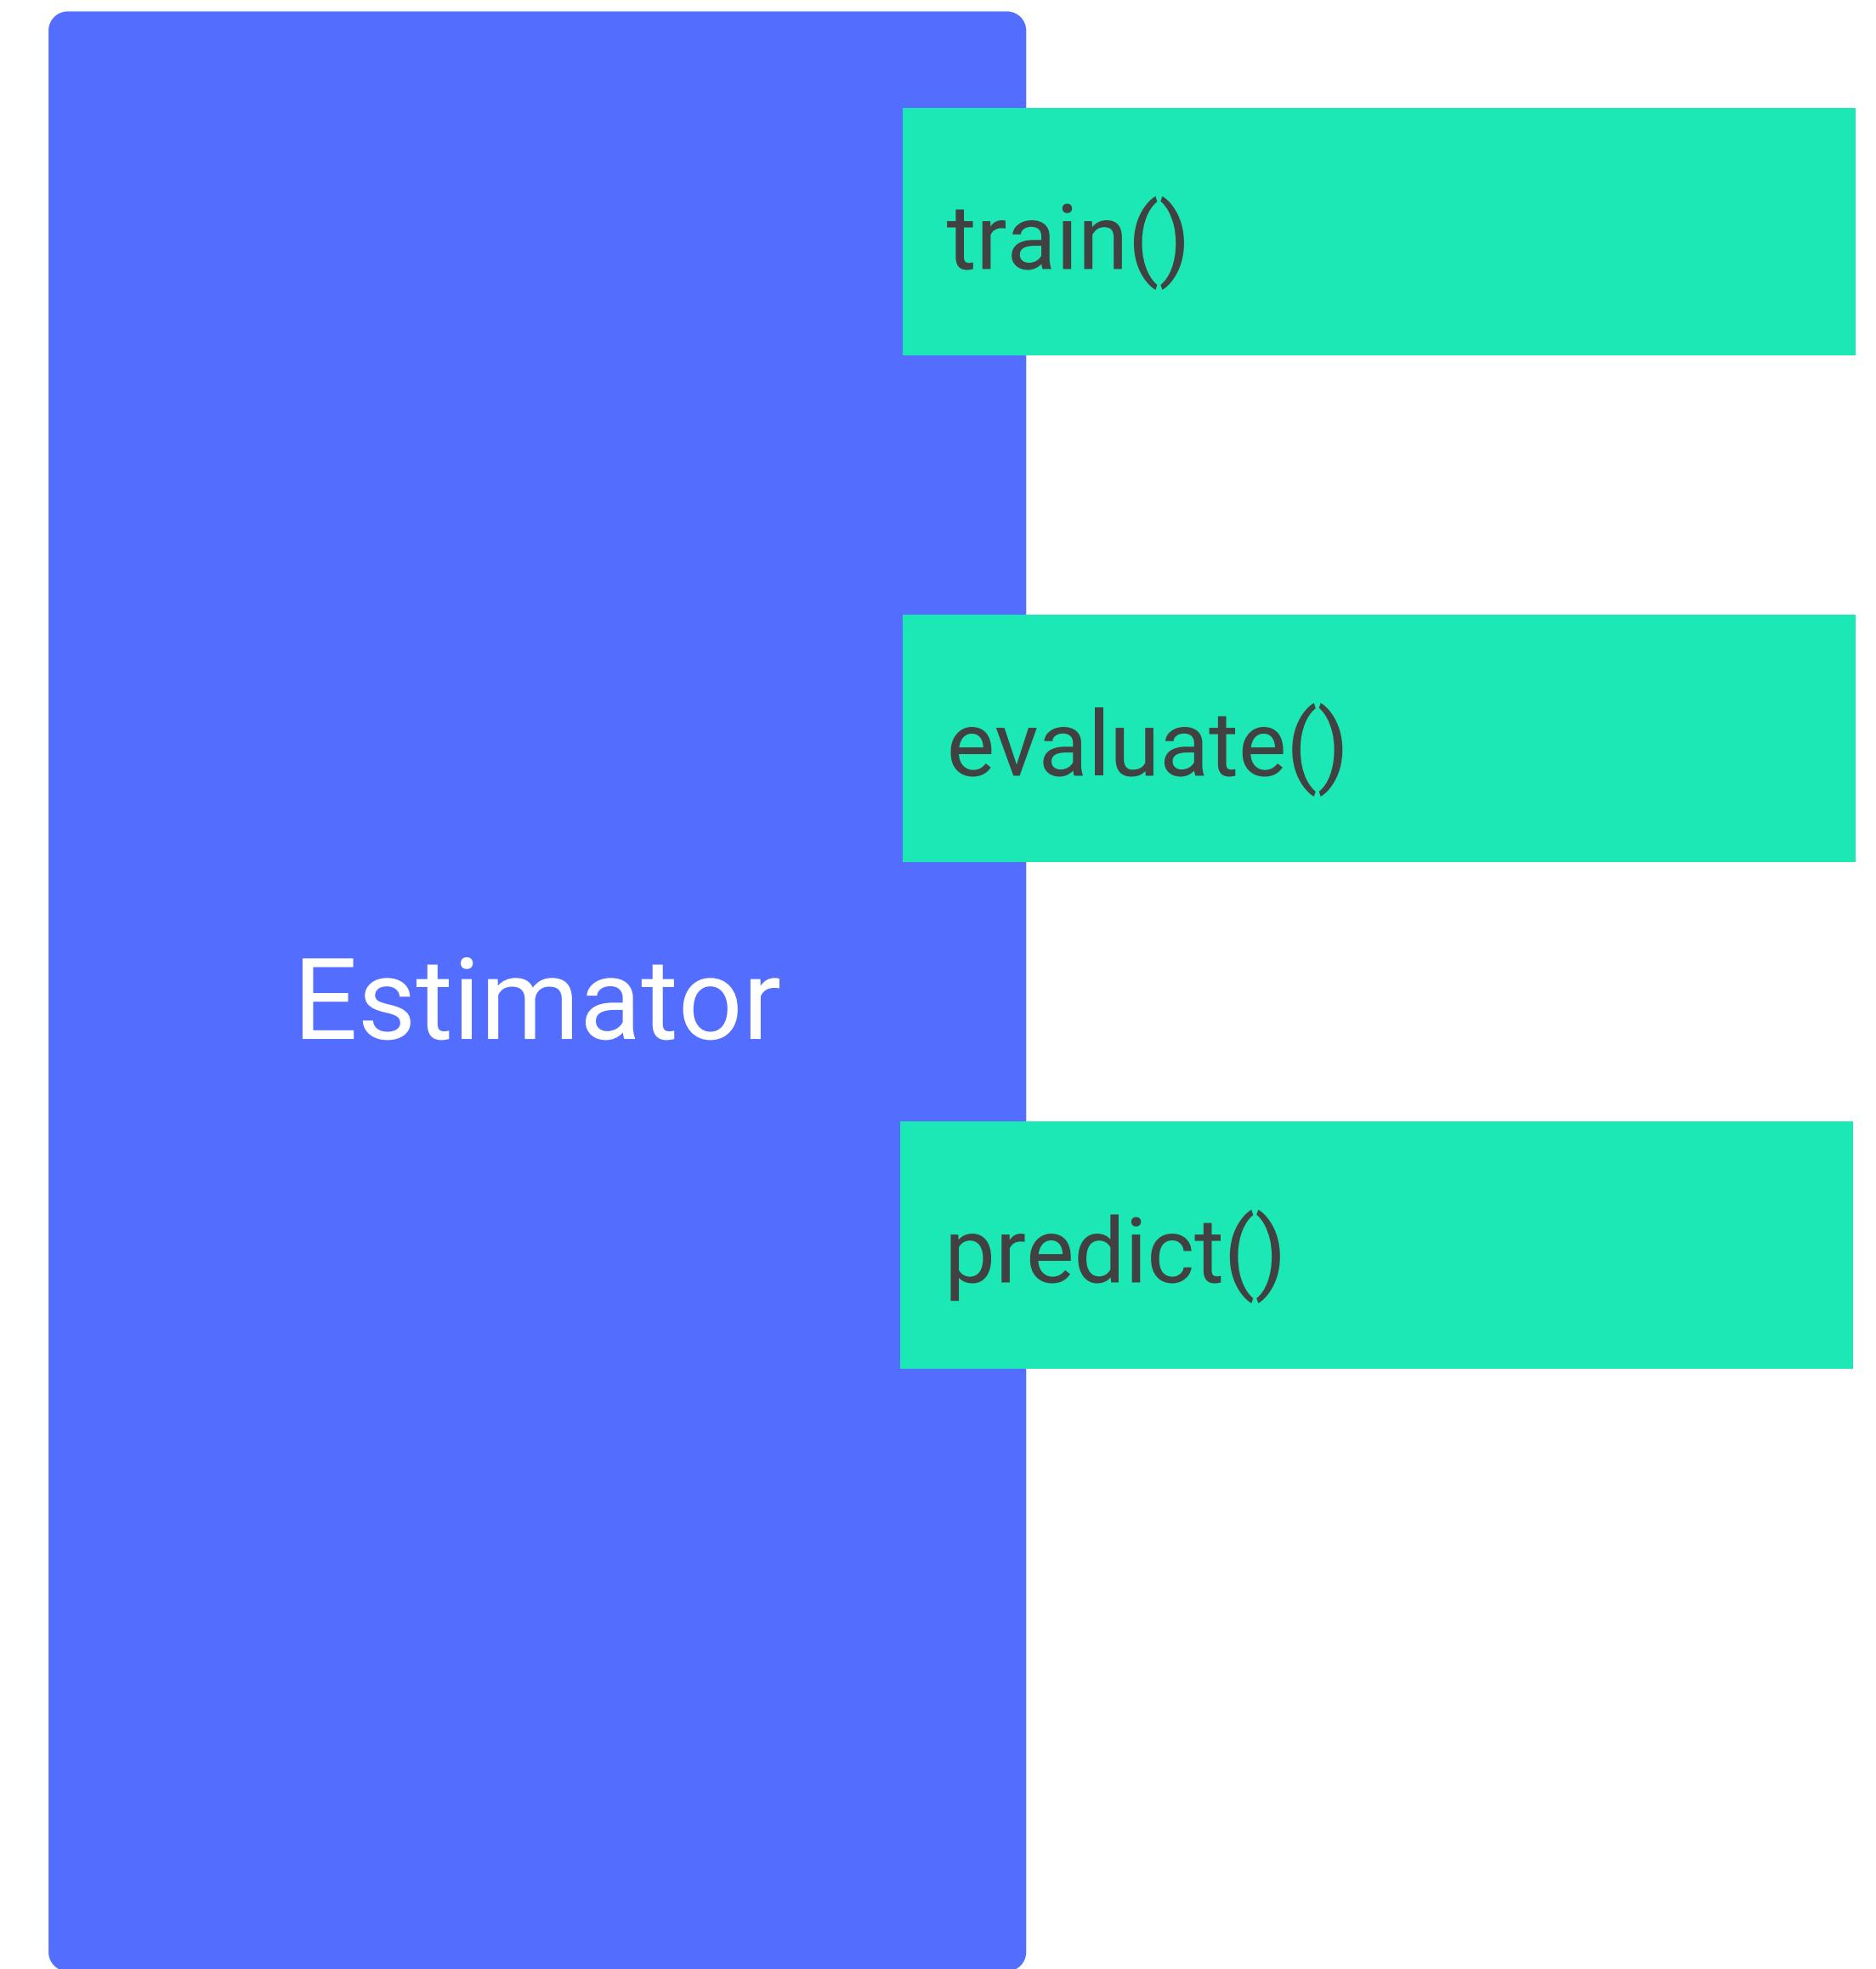
linear

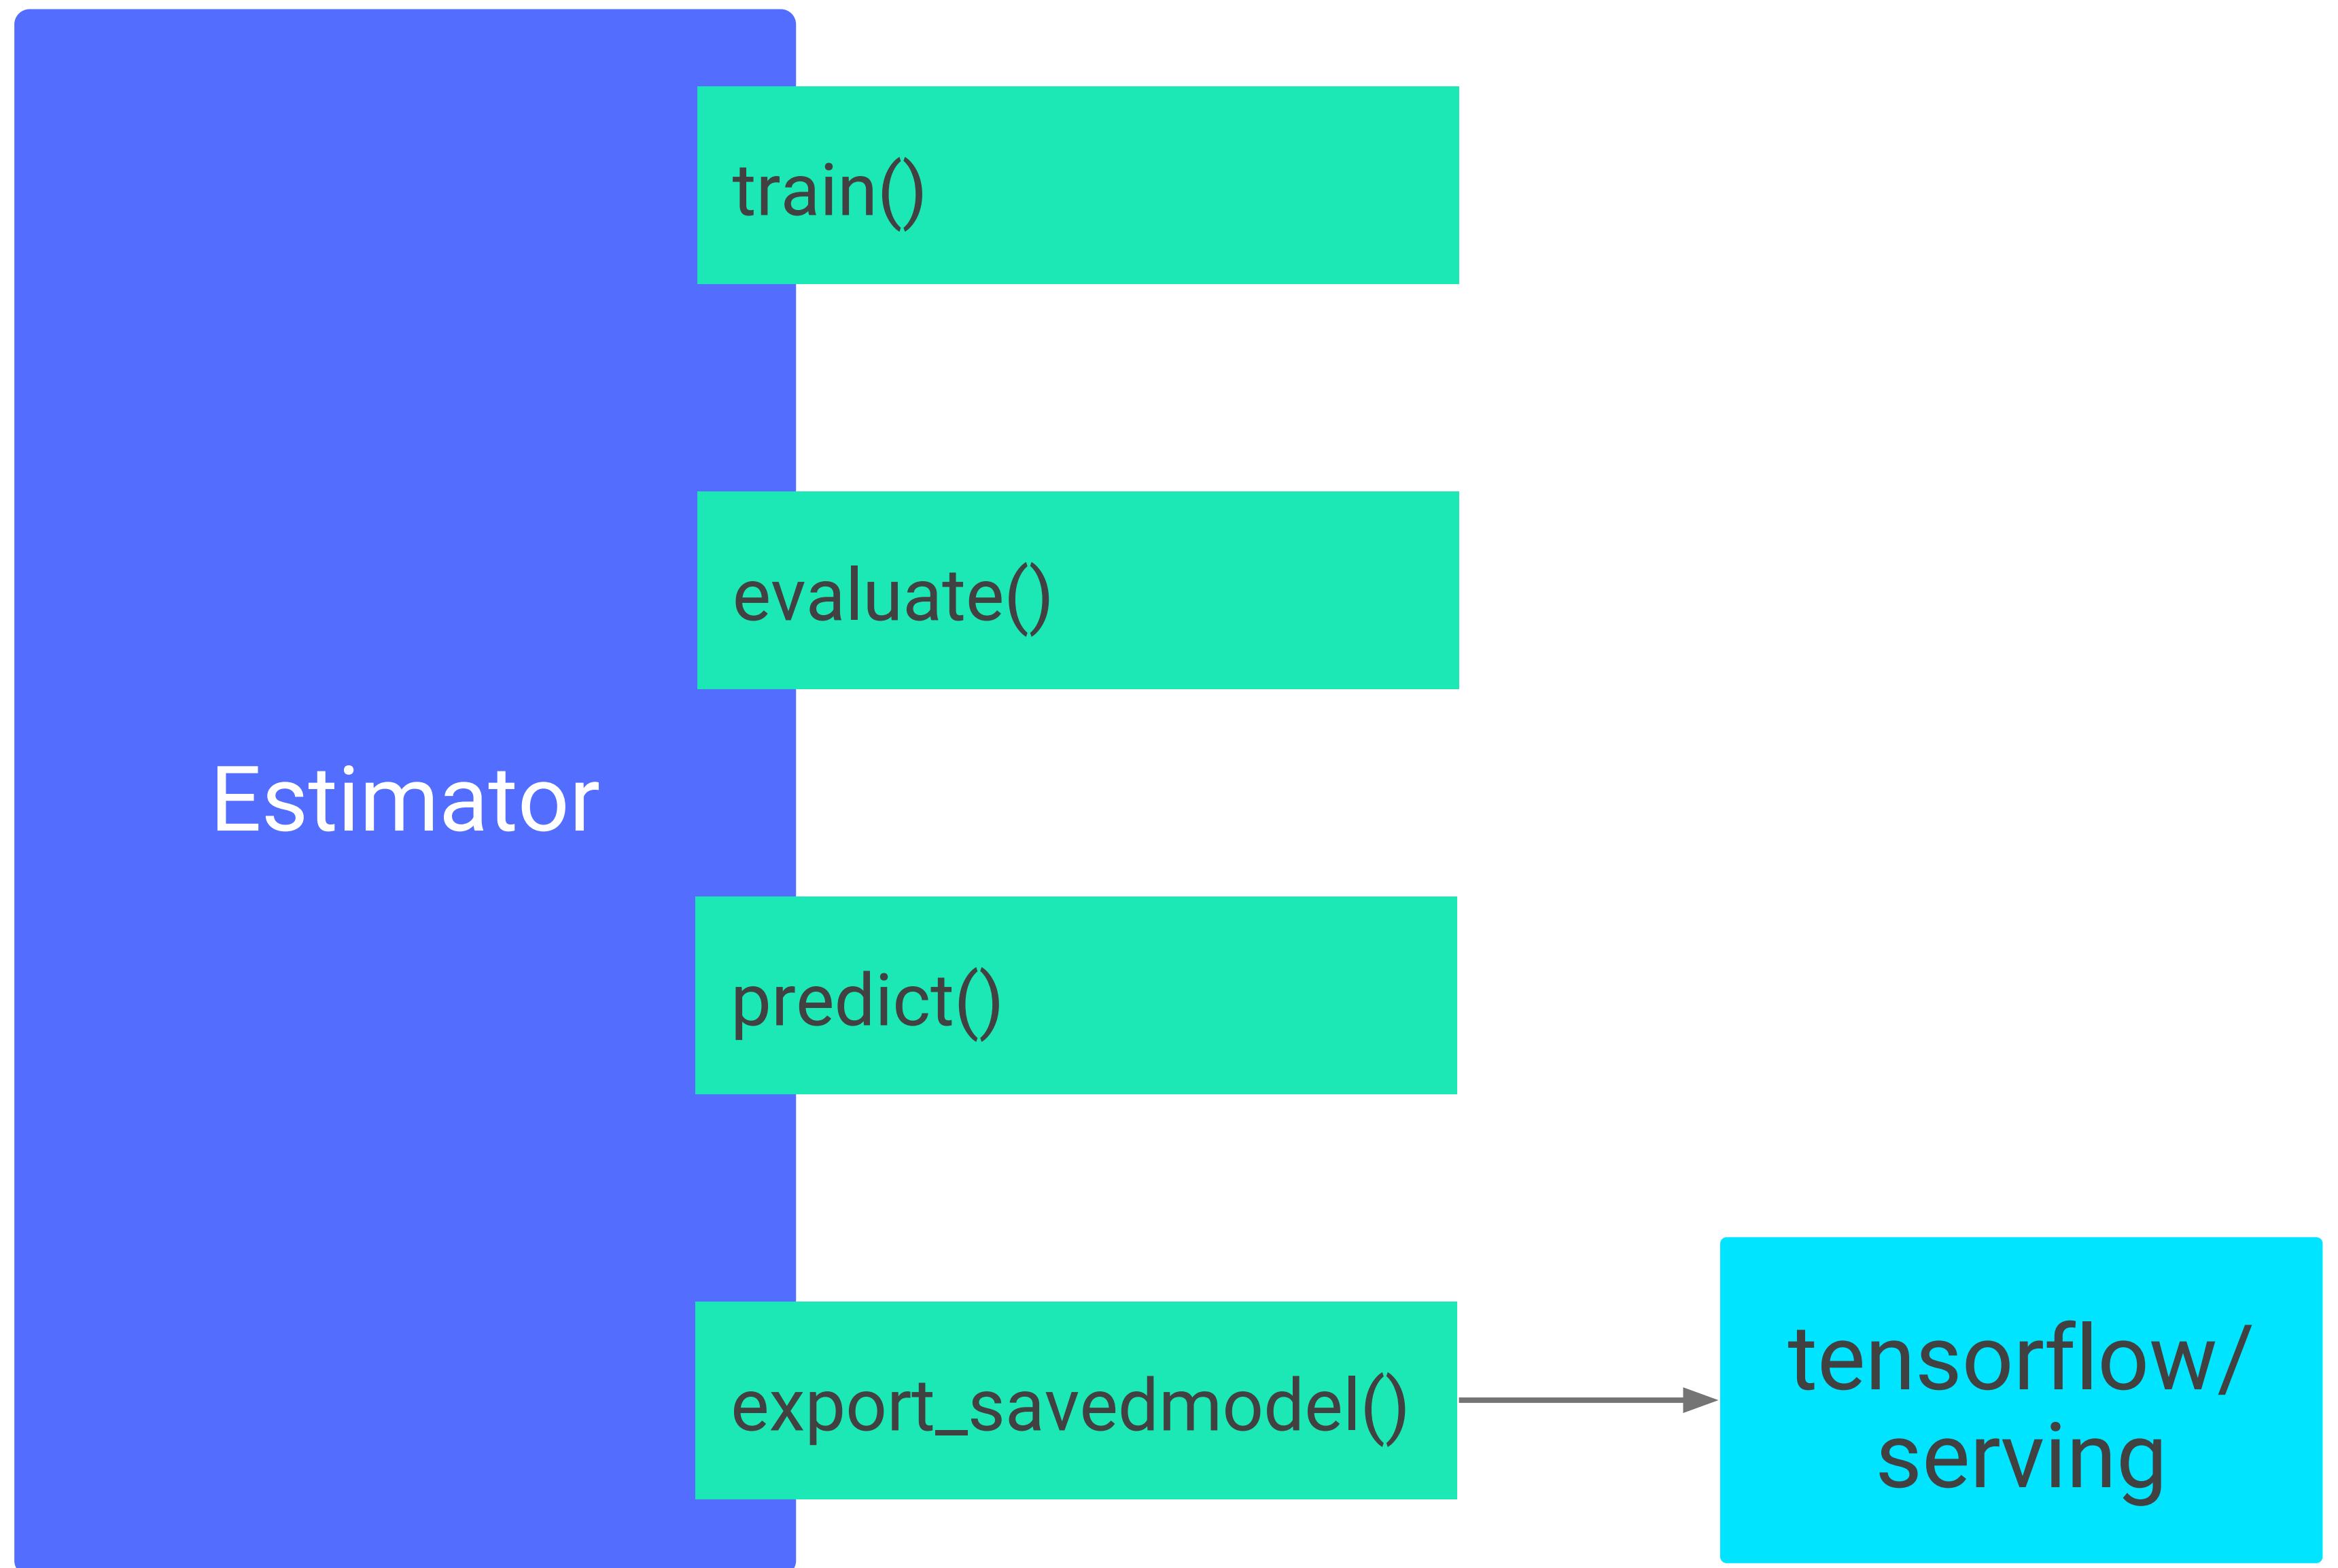
Subgraph: 1364 nodes

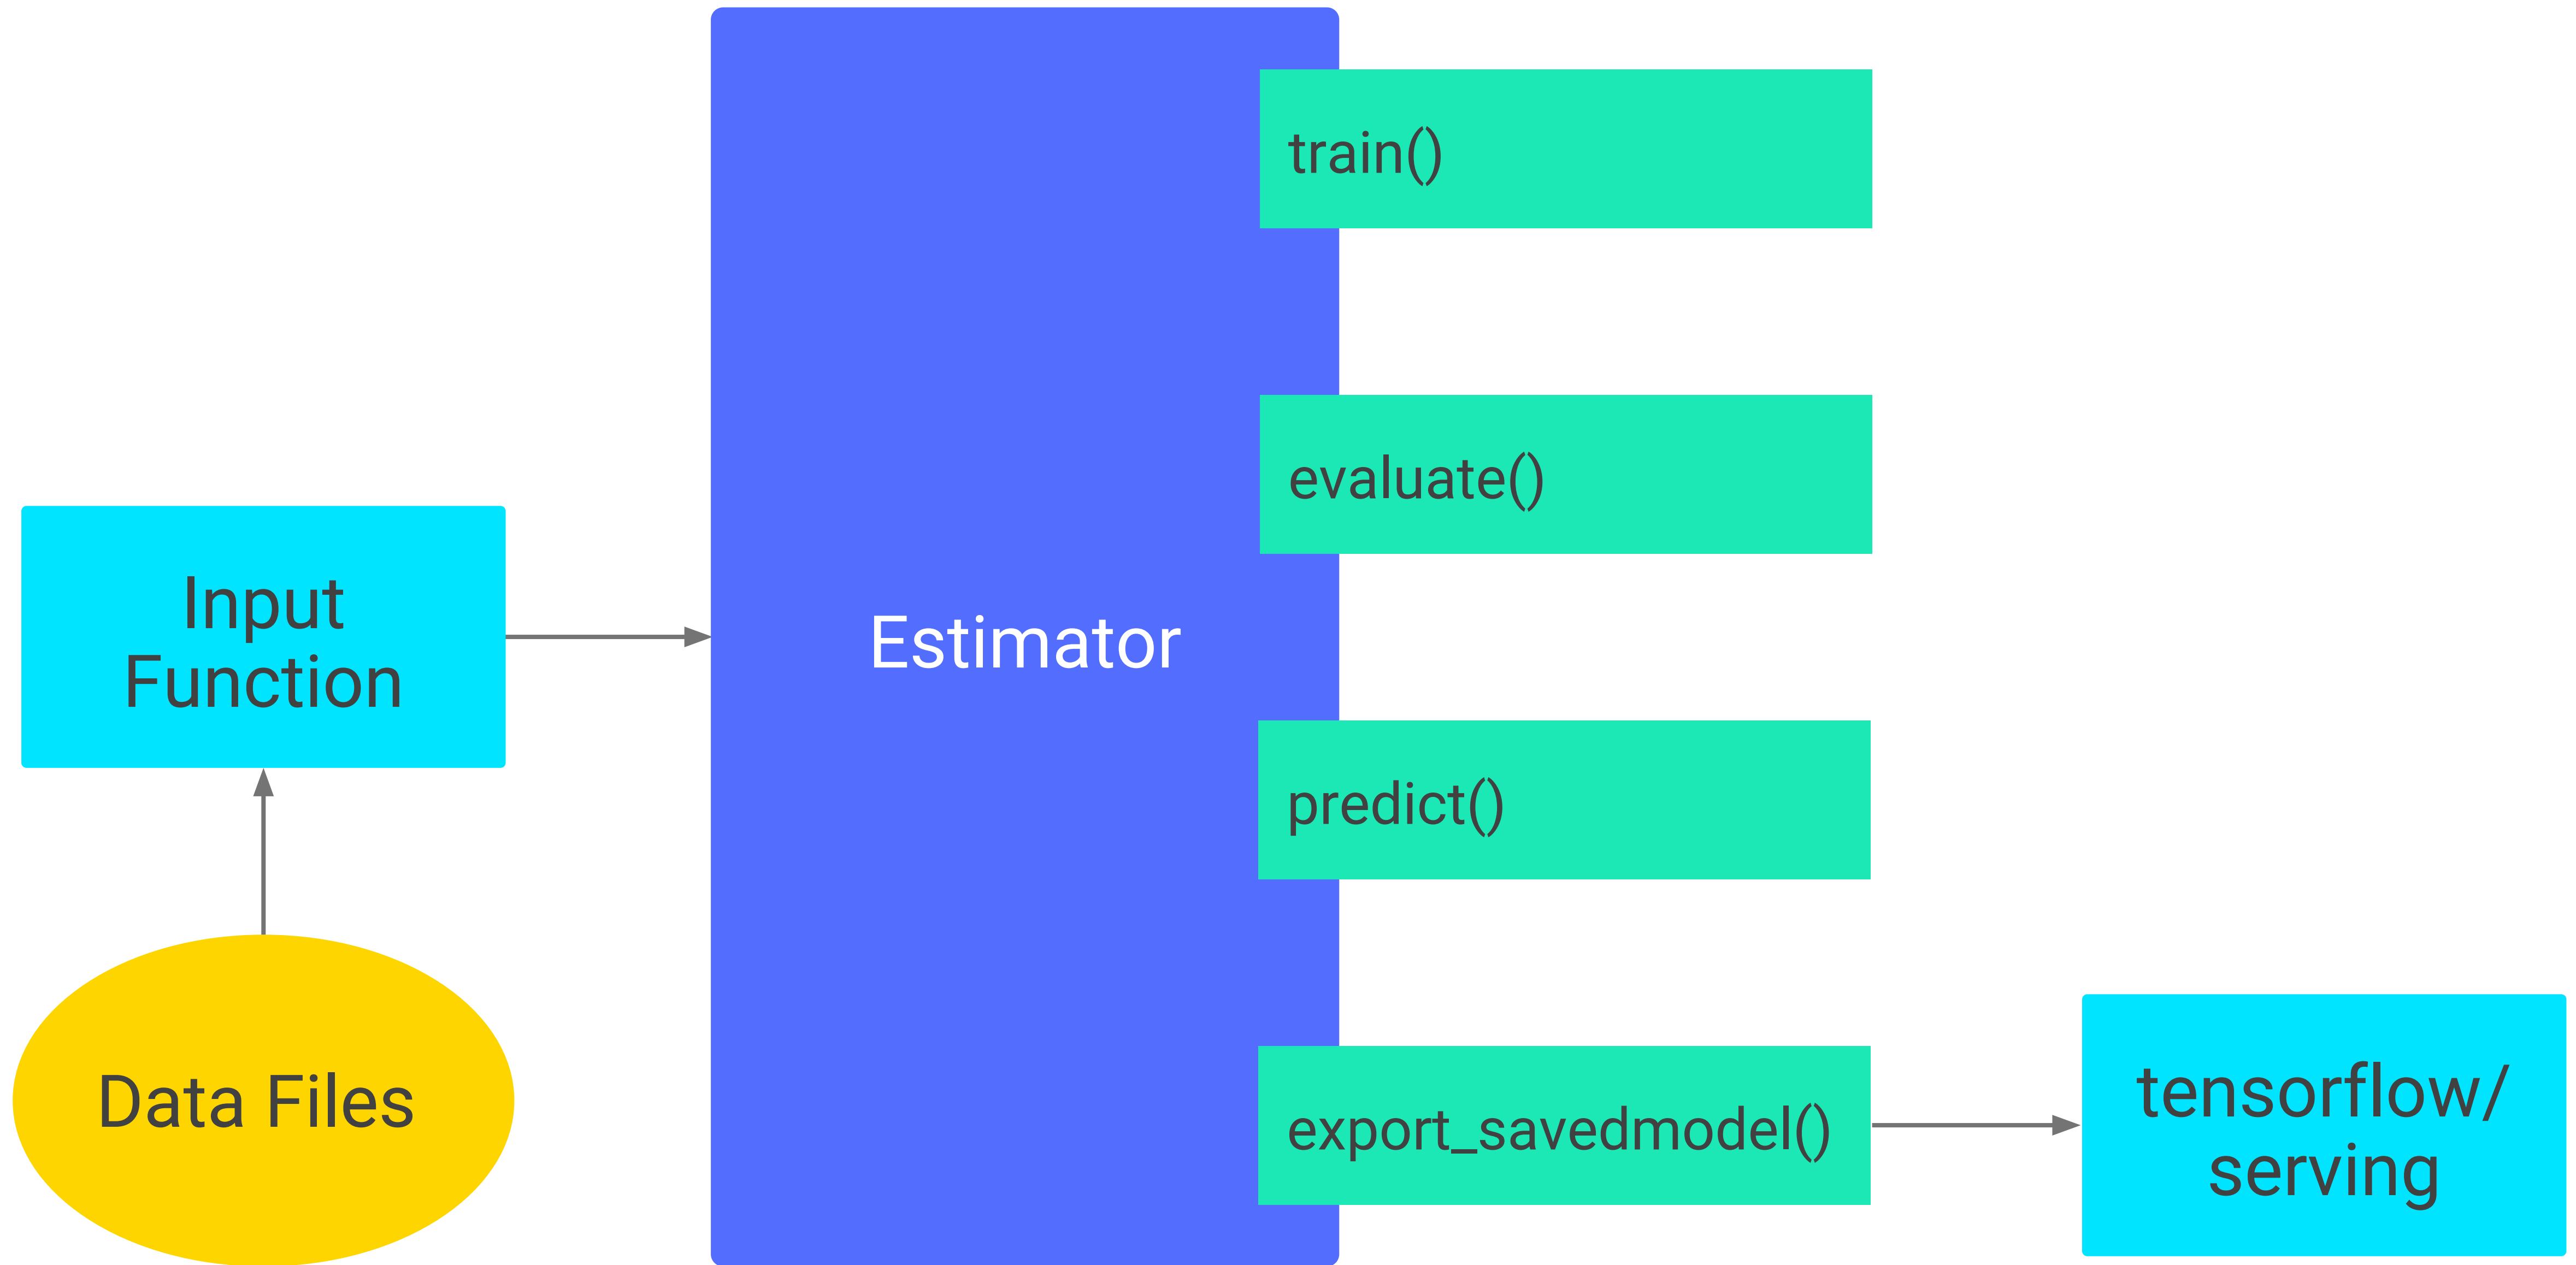


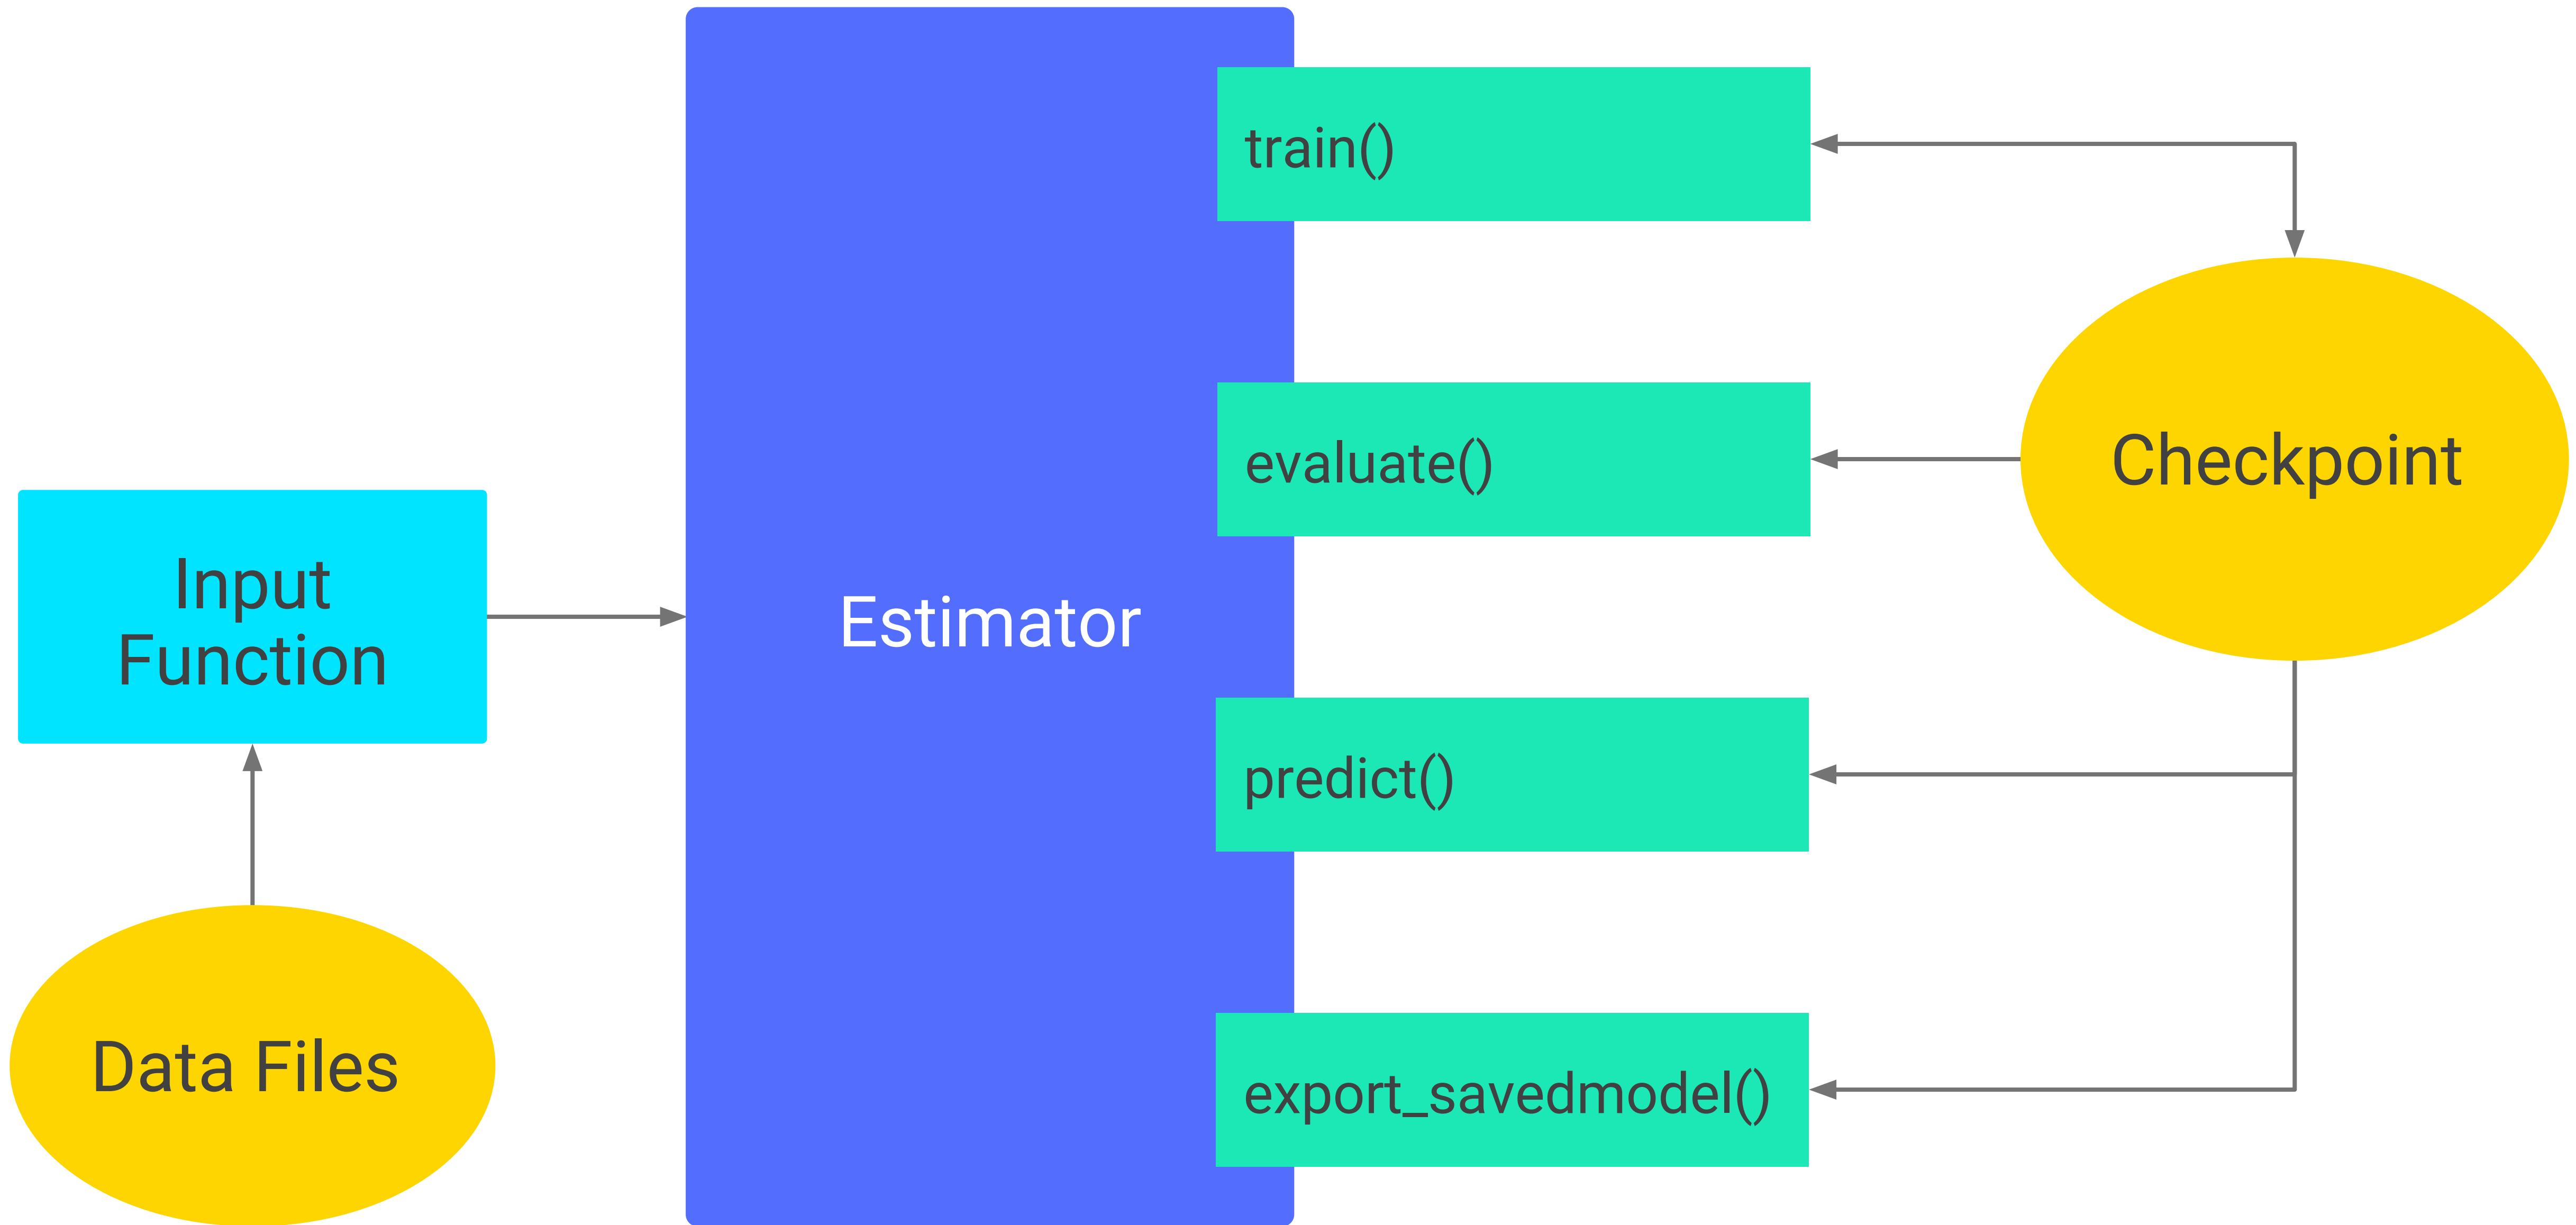


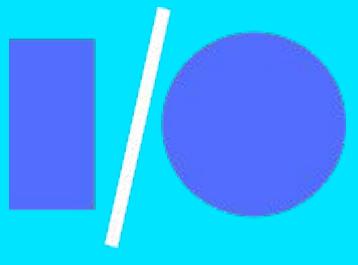




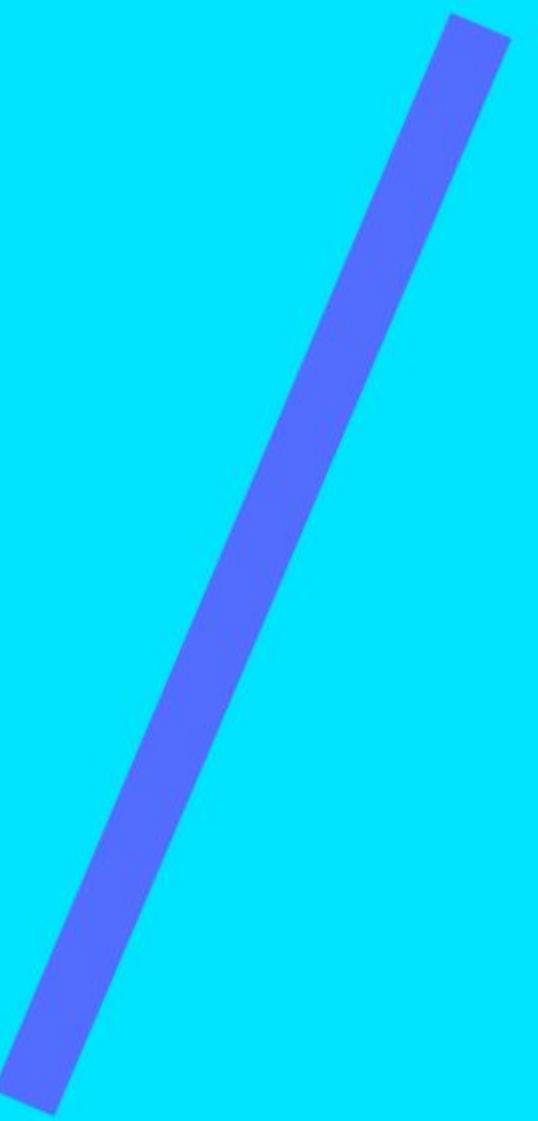






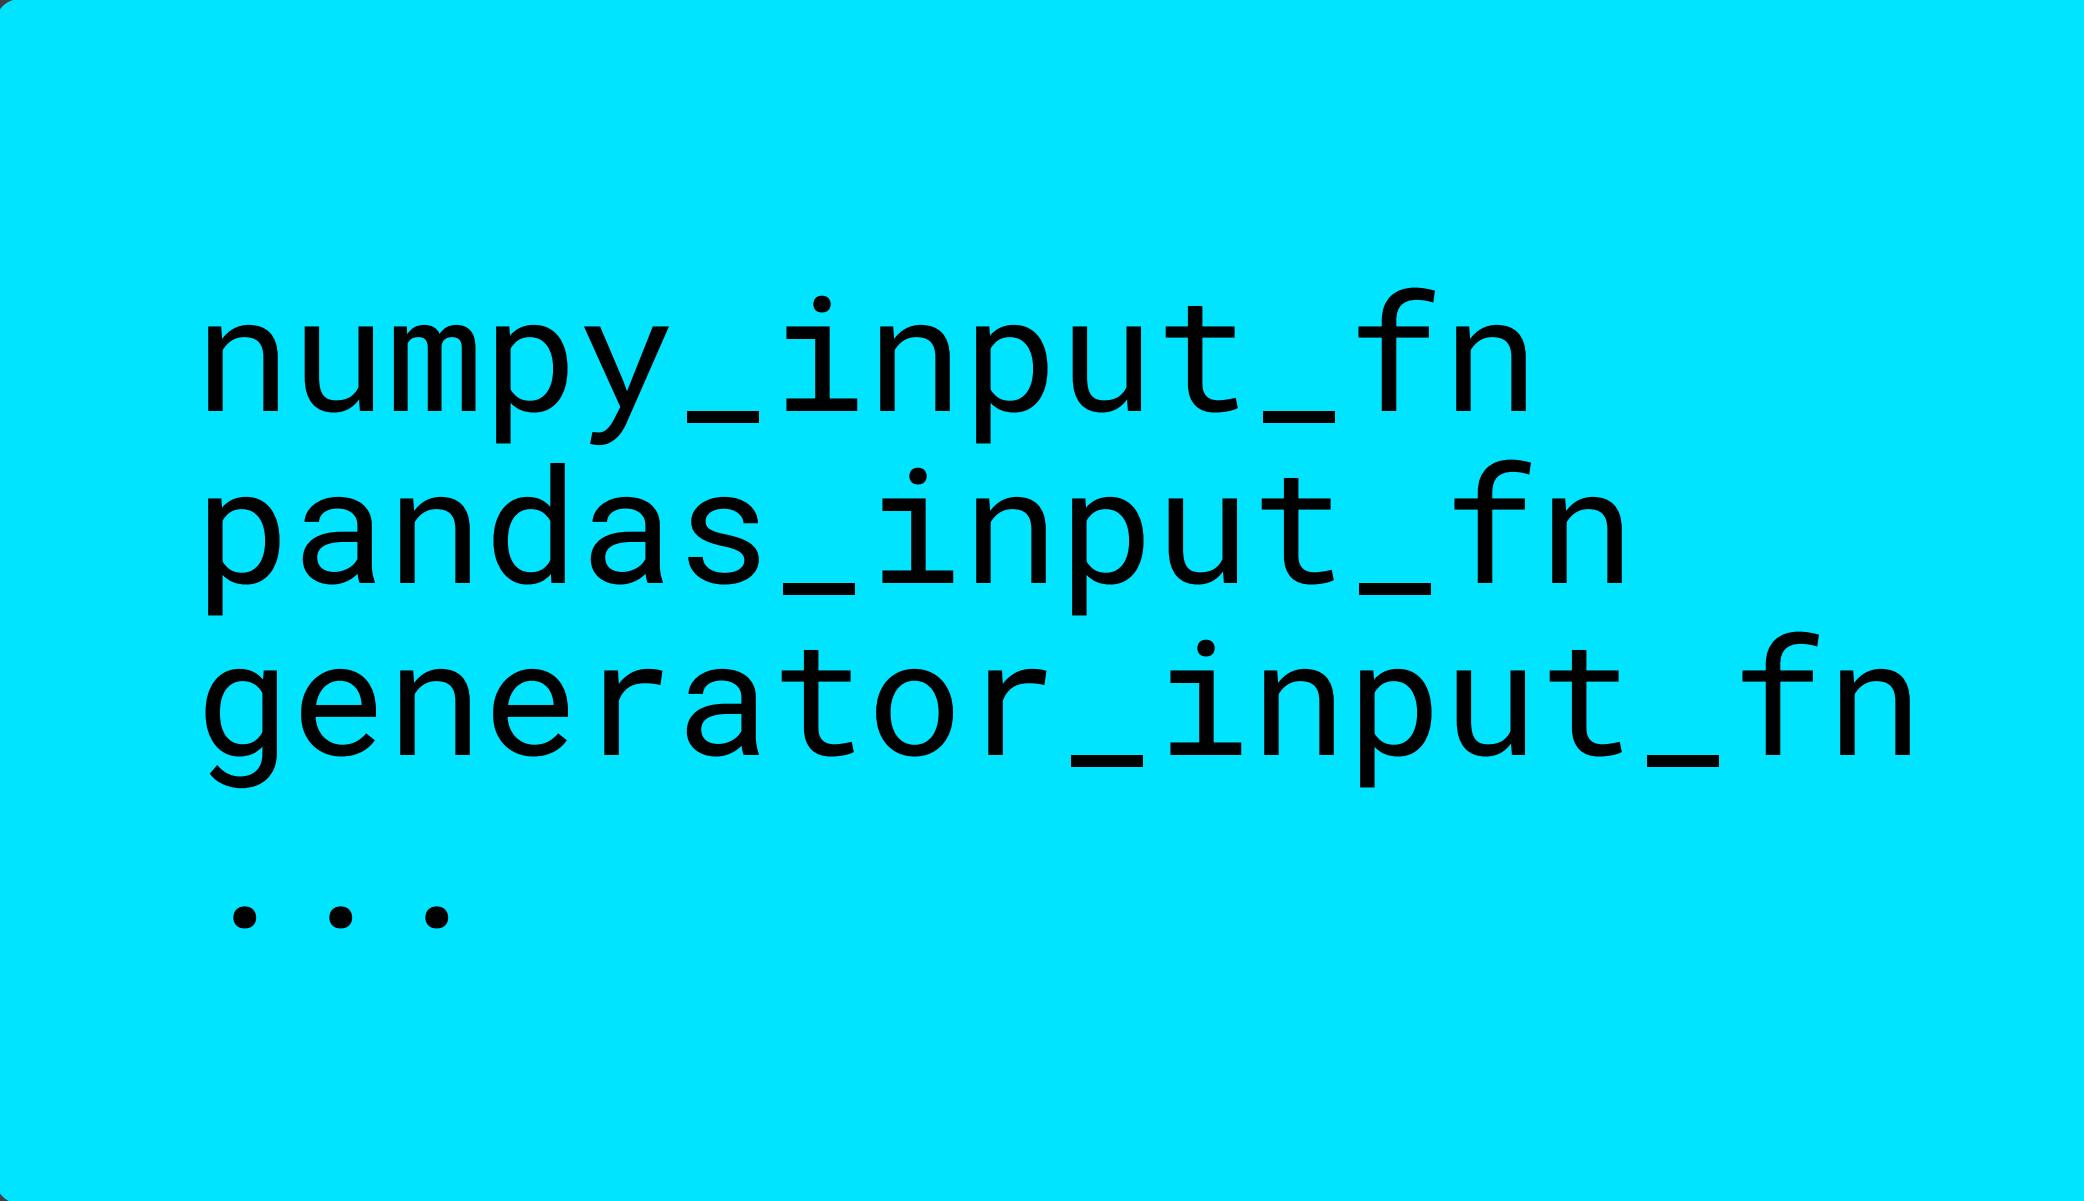


I want
... something else



```
...  
regressor = LinearRegressor(feature_columns=[make, horsepower,  
                                              cylinders, ...])
```

```
train_input_fn = pandas_input_fn(x=input_data, y=input_label,  
                                 batch_size=64, shuffle=True,  
                                 num_epochs=None)
```

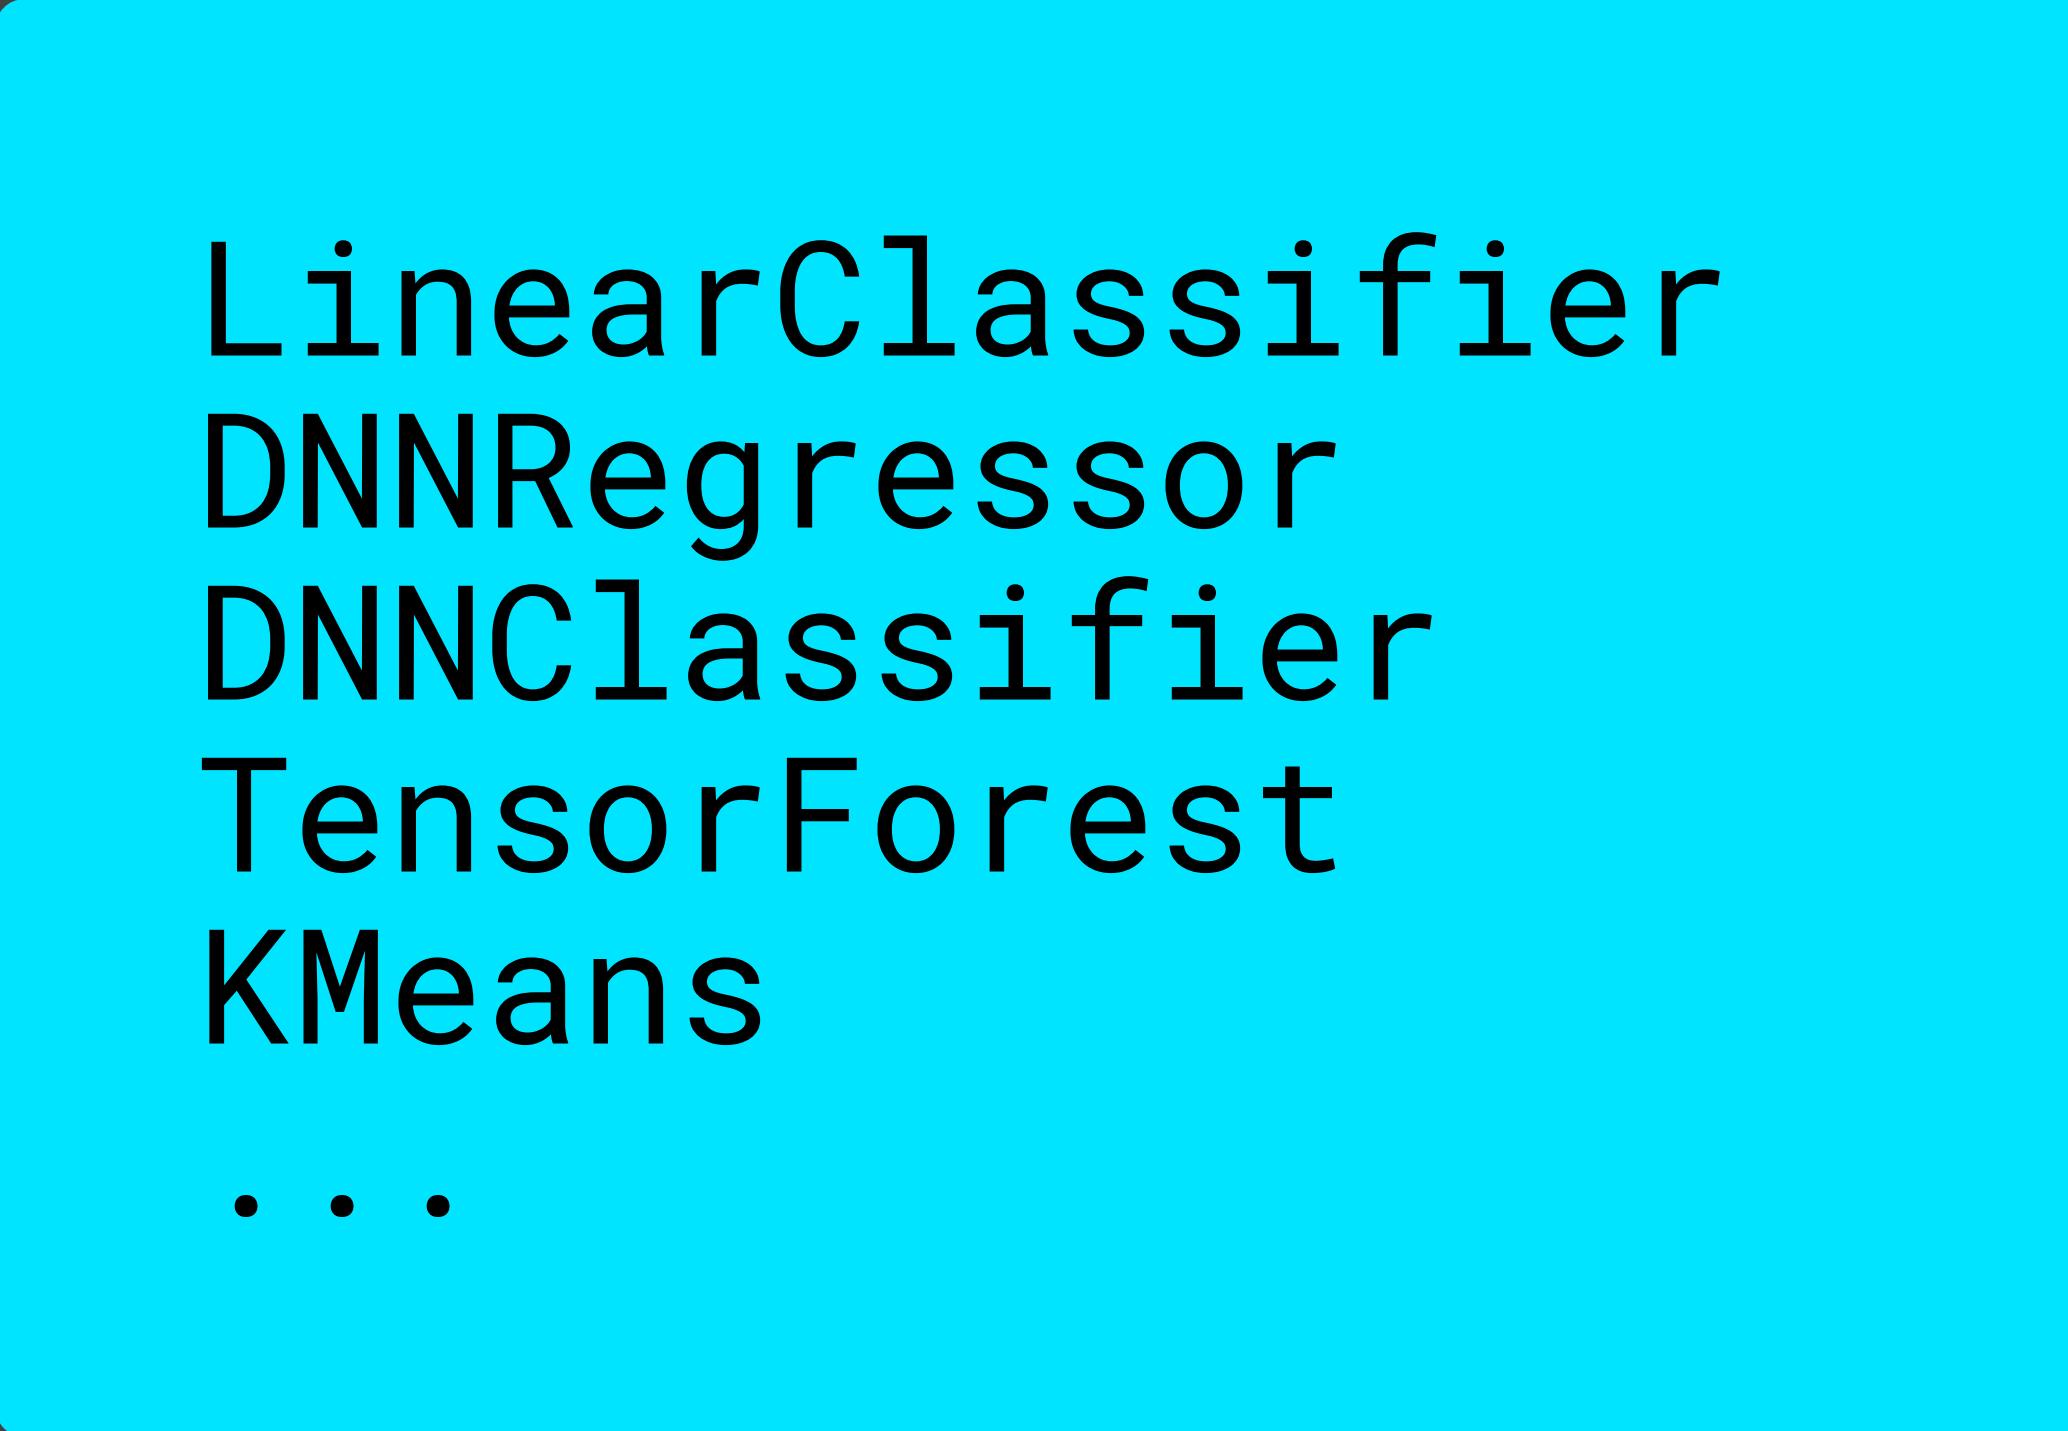


numpy_input_fn
pandas_input_fn
generator_input_fn
...

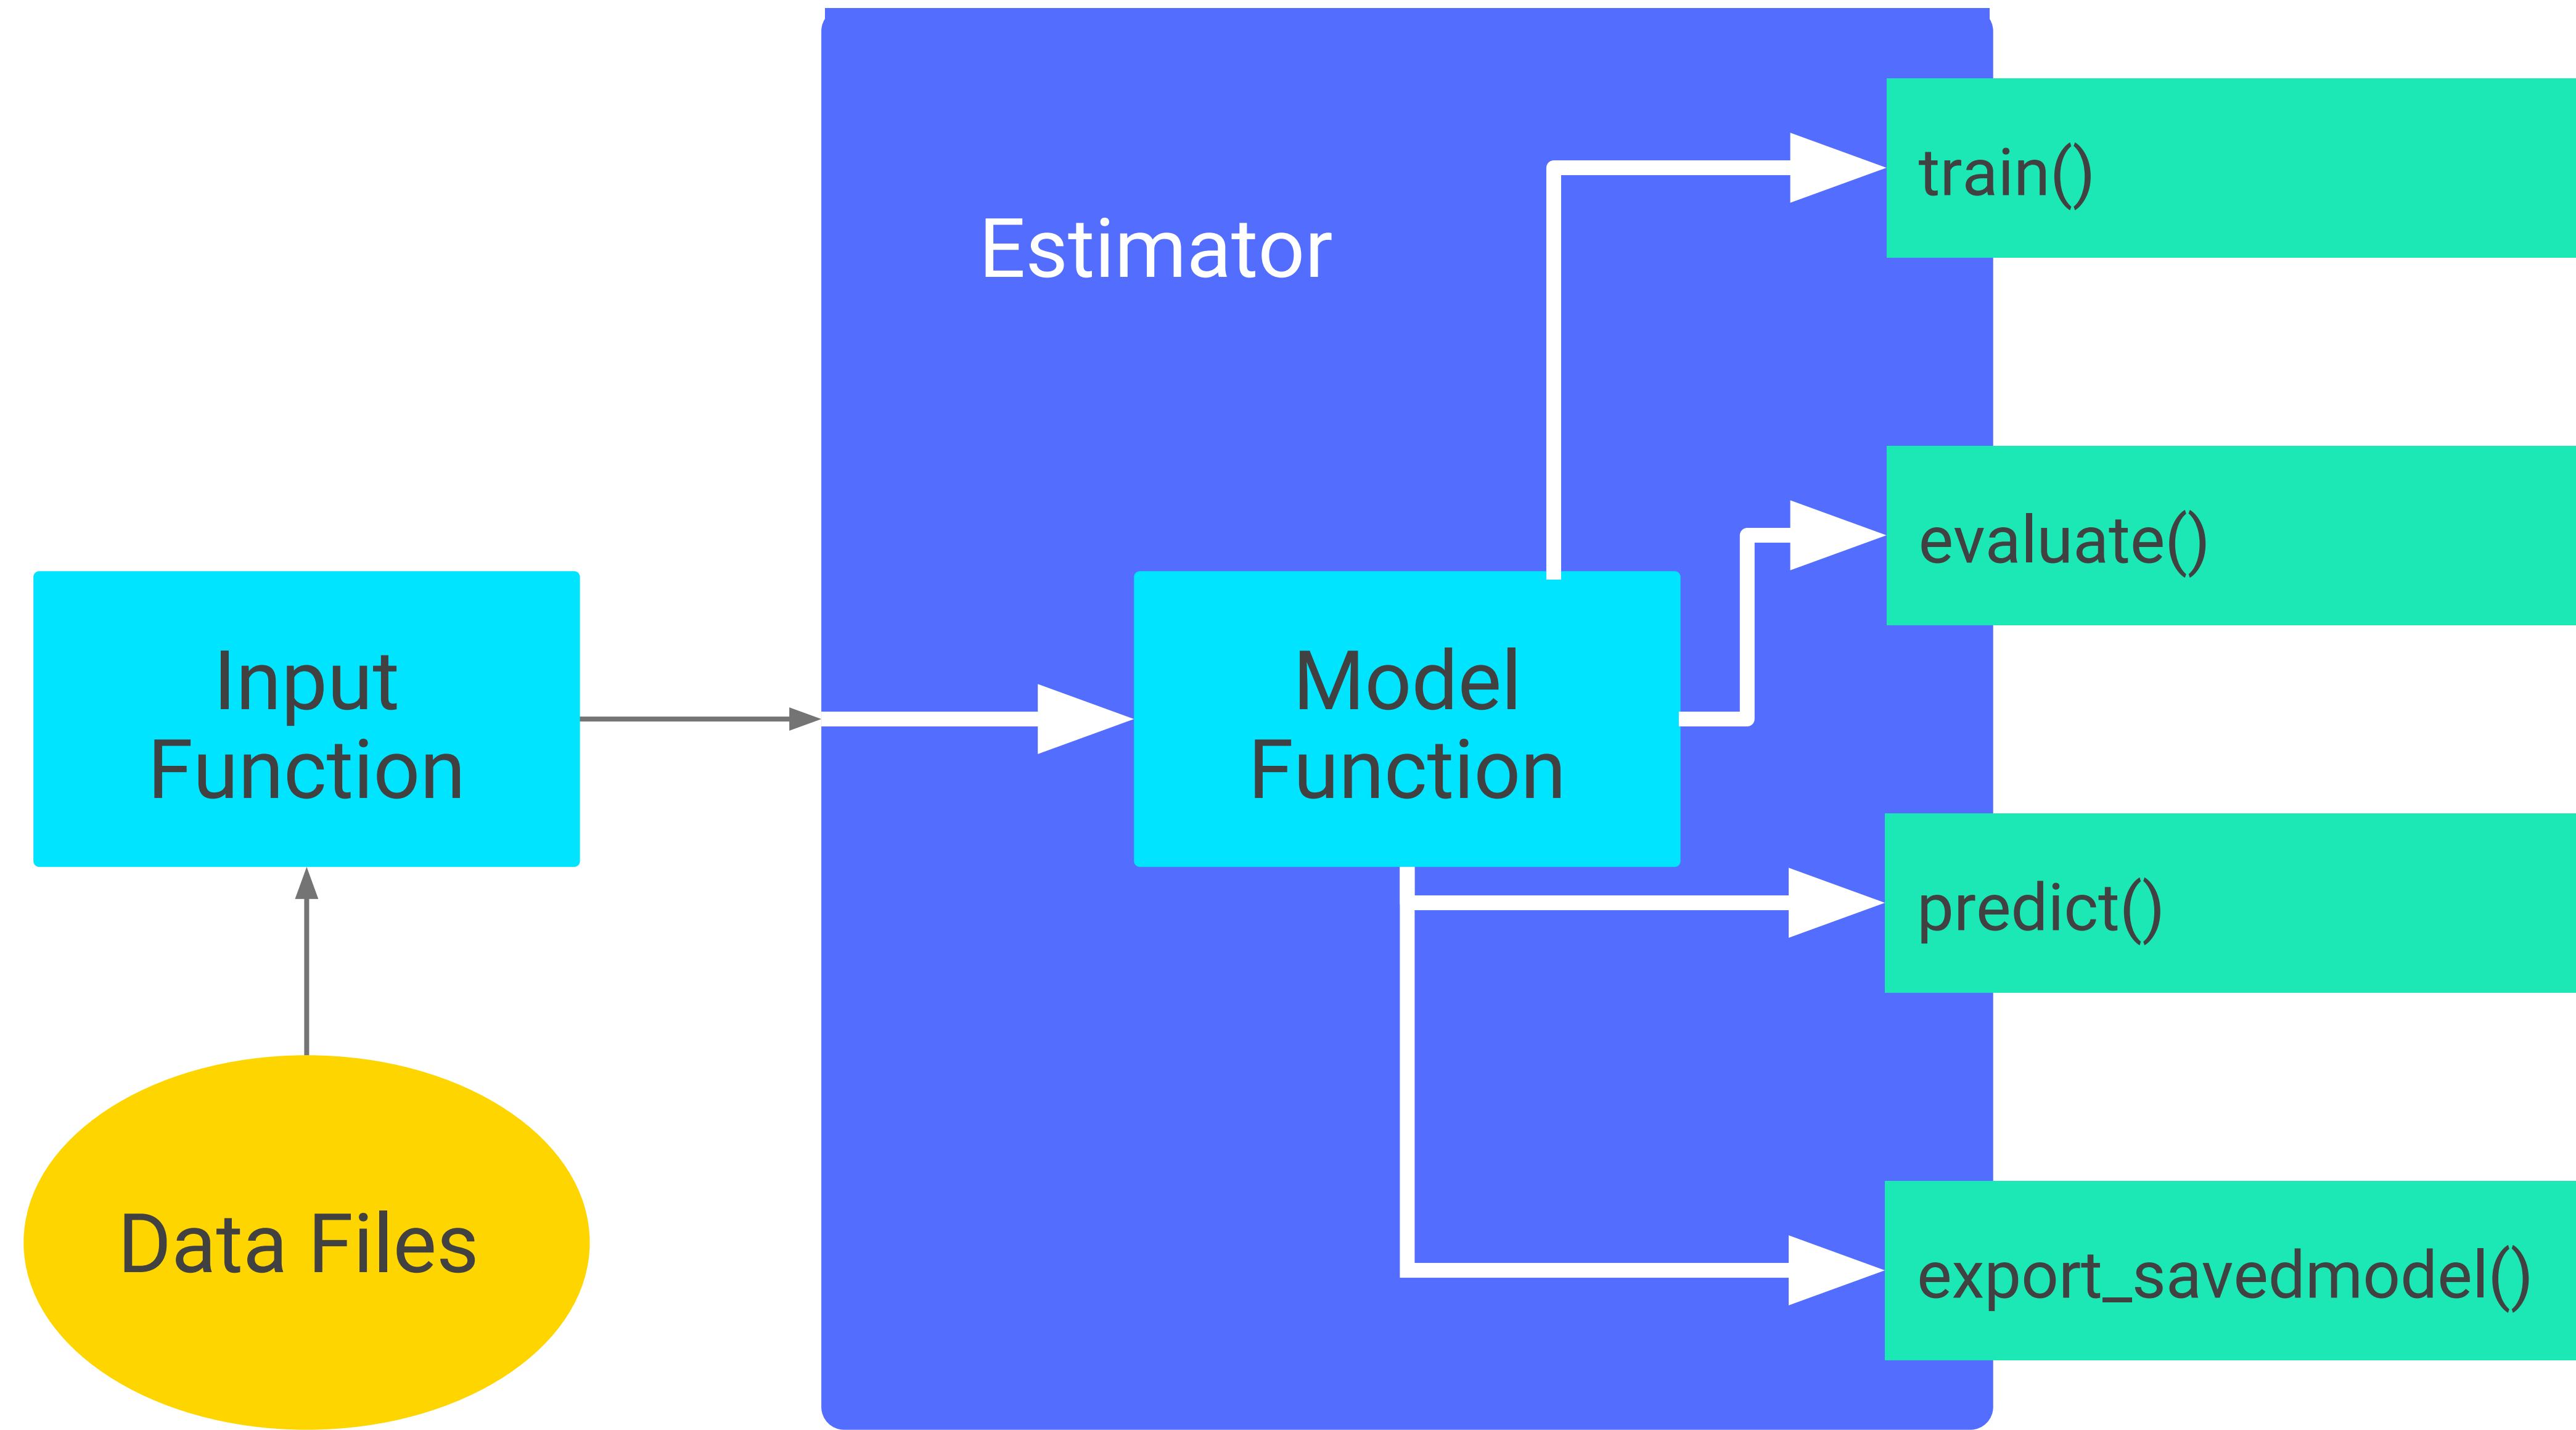
```
cylinders = categorical_column_with_vocabulary_list(  
    'num-of-cylinders', ['two', 'three', 'four', ...])
```

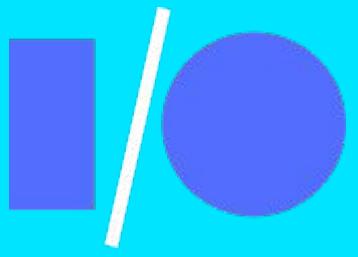
```
...
```

```
regressor = LinearRegressor(feature_columns=[make, horsepower,  
                                              cylinders, ...])
```

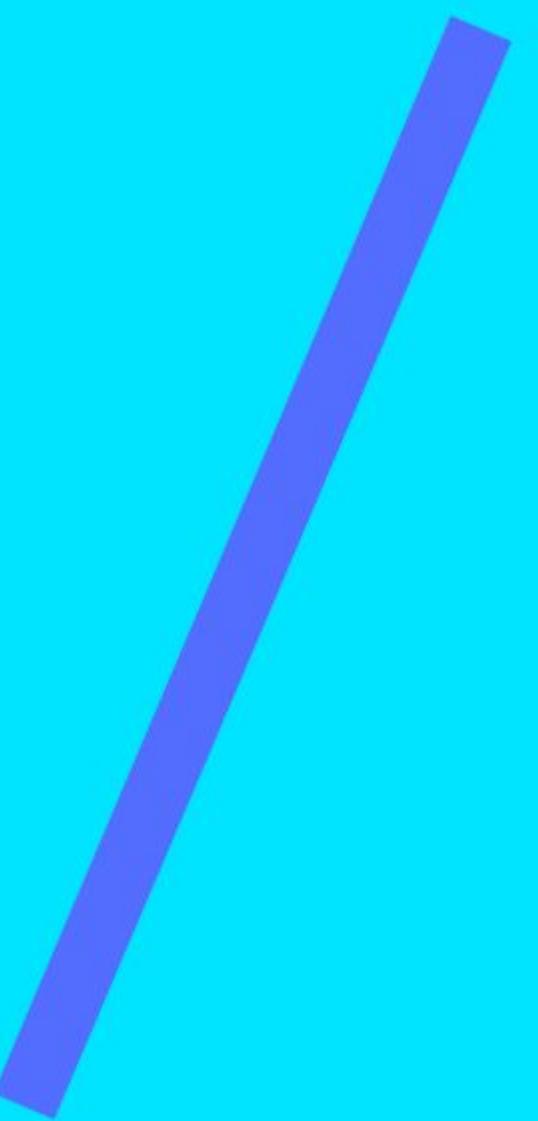


LinearClassifier
DNNRegressor
DNNClassifier
TensorForest
KMeans
...





Distributed TensorFlow



```
def experiment_fn(run_config, hparams):  
    # Make Estimator  
    regressor = DNNRegressor(..., config=run_config,  
                            hidden_units=hparams['units'])  
  
    # Collect information for training  
    return Experiment(estimator=regressor,  
                      train_input_fn=pandas_input_fn(...),  
                      eval_input_fn=pandas_input_fn(...))  
  
if __name__ == '__main__':  
    learn_runner.run(experiment_fn,  
                     run_config=RunConfig(model_dir="/tmp/output_dir"))
```

```
def experiment_fn(run_config, hparams):  
    # Make Estimator  
    regressor = DNNRegressor(..., config=run_config,  
                            hidden_units=hparams['units'])  
  
    # Collect information for training  
    return Experiment(estimator=regressor,  
                      train_input_fn=pandas_input_fn(...),  
                      eval_input_fn=pandas_input_fn(...))  
  
if __name__ == '__main__':  
    learn_runner.run(experiment_fn,  
                     run_config=RunConfig(model_dir="/tmp/output_dir"))
```

```
def experiment_fn(run_config, hparams):  
    # Make Estimator  
    regressor = DNNRegressor(..., config=run_config,  
                            hidden_units=hparams['units'])  
  
    # Collect information for training  
    return Experiment(estimator=regressor,  
                      train_input_fn=pandas_input_fn(...),  
                      eval_input_fn=pandas_input_fn(...))  
  
if __name__ == '__main__':  
    learn_runner.run(experiment_fn,  
                     run_config=RunConfig(model_dir="/tmp/output_dir"))
```

```
def experiment_fn(run_config, hparams):  
    # Make Estimator  
    regressor = DNNRegressor(..., config=run_config,  
                            hidden_units=hparams['units'])  
  
    # Collect information for training  
    return Experiment(estimator=regressor,  
                      train_input_fn=pandas_input_fn(...),  
                      eval_input_fn=pandas_input_fn(...))  
  
if __name__ == '__main__':  
    learn_runner.run(experiment_fn,  
                     run_config=RunConfig(model_dir="/tmp/output_dir"))
```

```
cluster_spec = {
    'ps': ['host1:2222', 'host2:2222', ...],
    'worker': ['host3:2222', 'host4:2222', ...],
}

os.environ['TF_CONFIG'] = json.dumps({
    'cluster': cluster_spec,
    'task': {'type': 'worker', 'index': 1}
})
```

Setting up a cluster

github.com/tensorflow/ecosystem





Complete ML models, quickly



Complete ML models, quickly



Distributed training



Complete ML models, quickly



Distributed training



Runs on CPU/GPU/TPU/mobile

Full Code available here

goo.gl/00gXiL

TensorFlow Tutorials

<https://www.tensorflow.org/tutorials>

Estimator Tutorial

<https://www.tensorflow.org/extend/estimators>

From Research to Production with TensorFlow Serving

Tomorrow, 1:30p, Amphitheatre

The Keras API

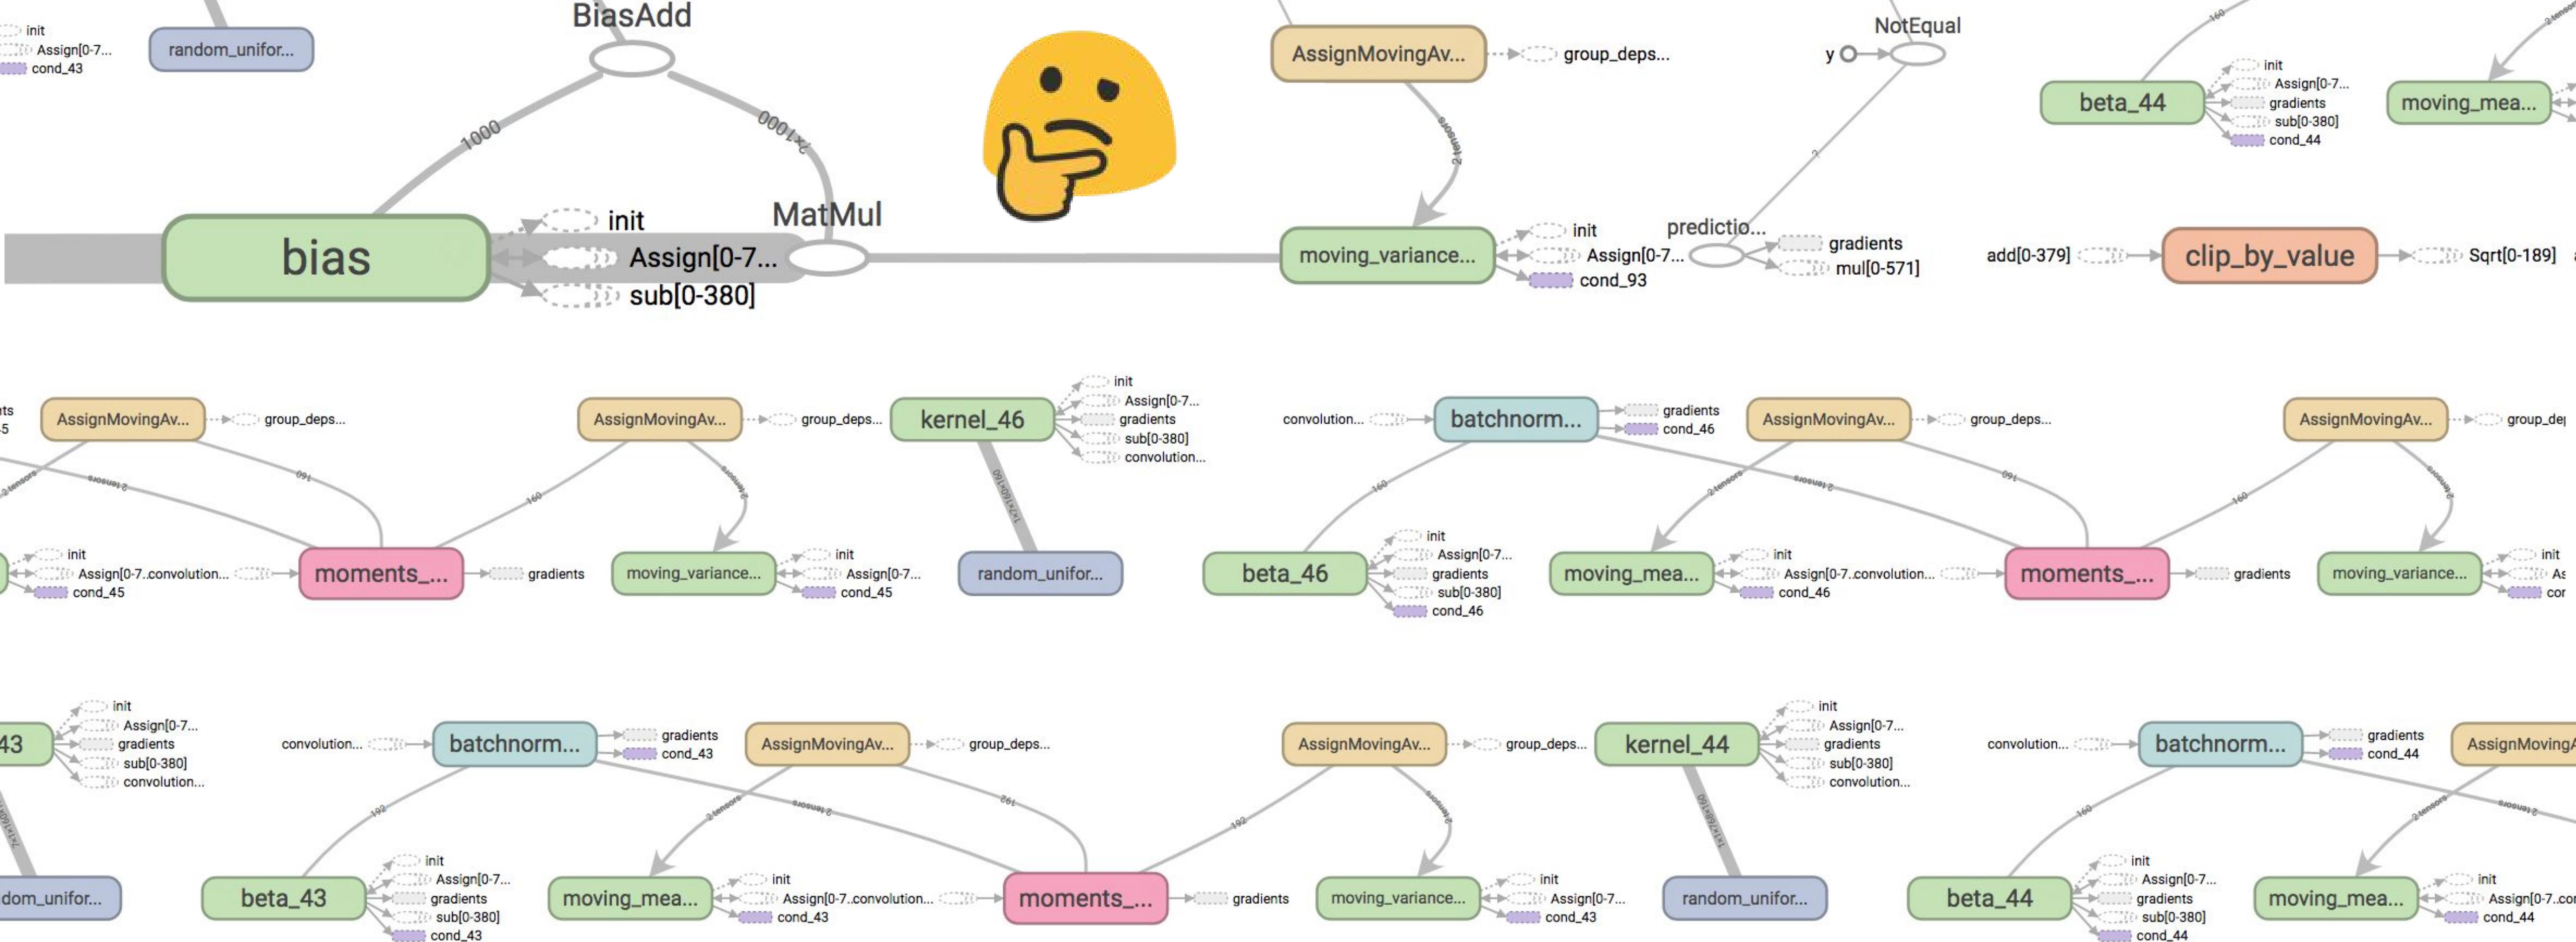
A simple, intuitive interface for TensorFlow



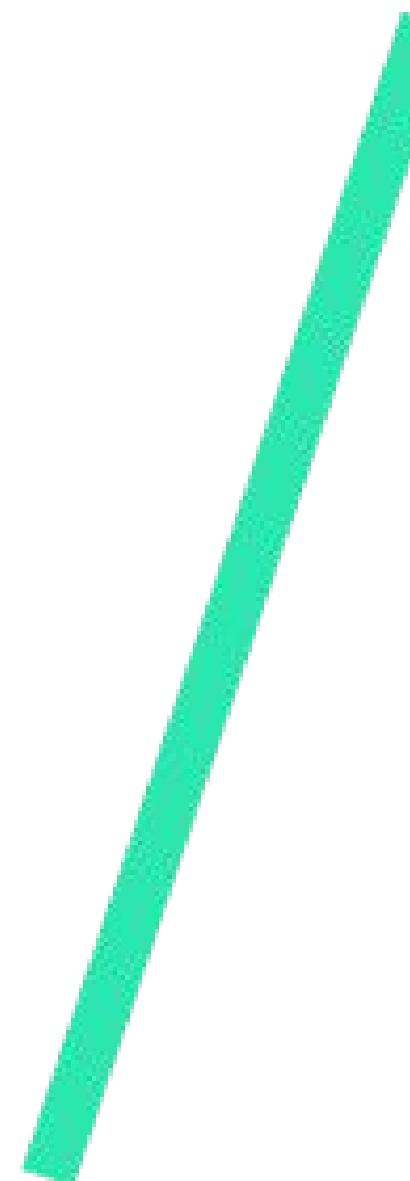
Baohua Liao
liaobaohua@gmail.com



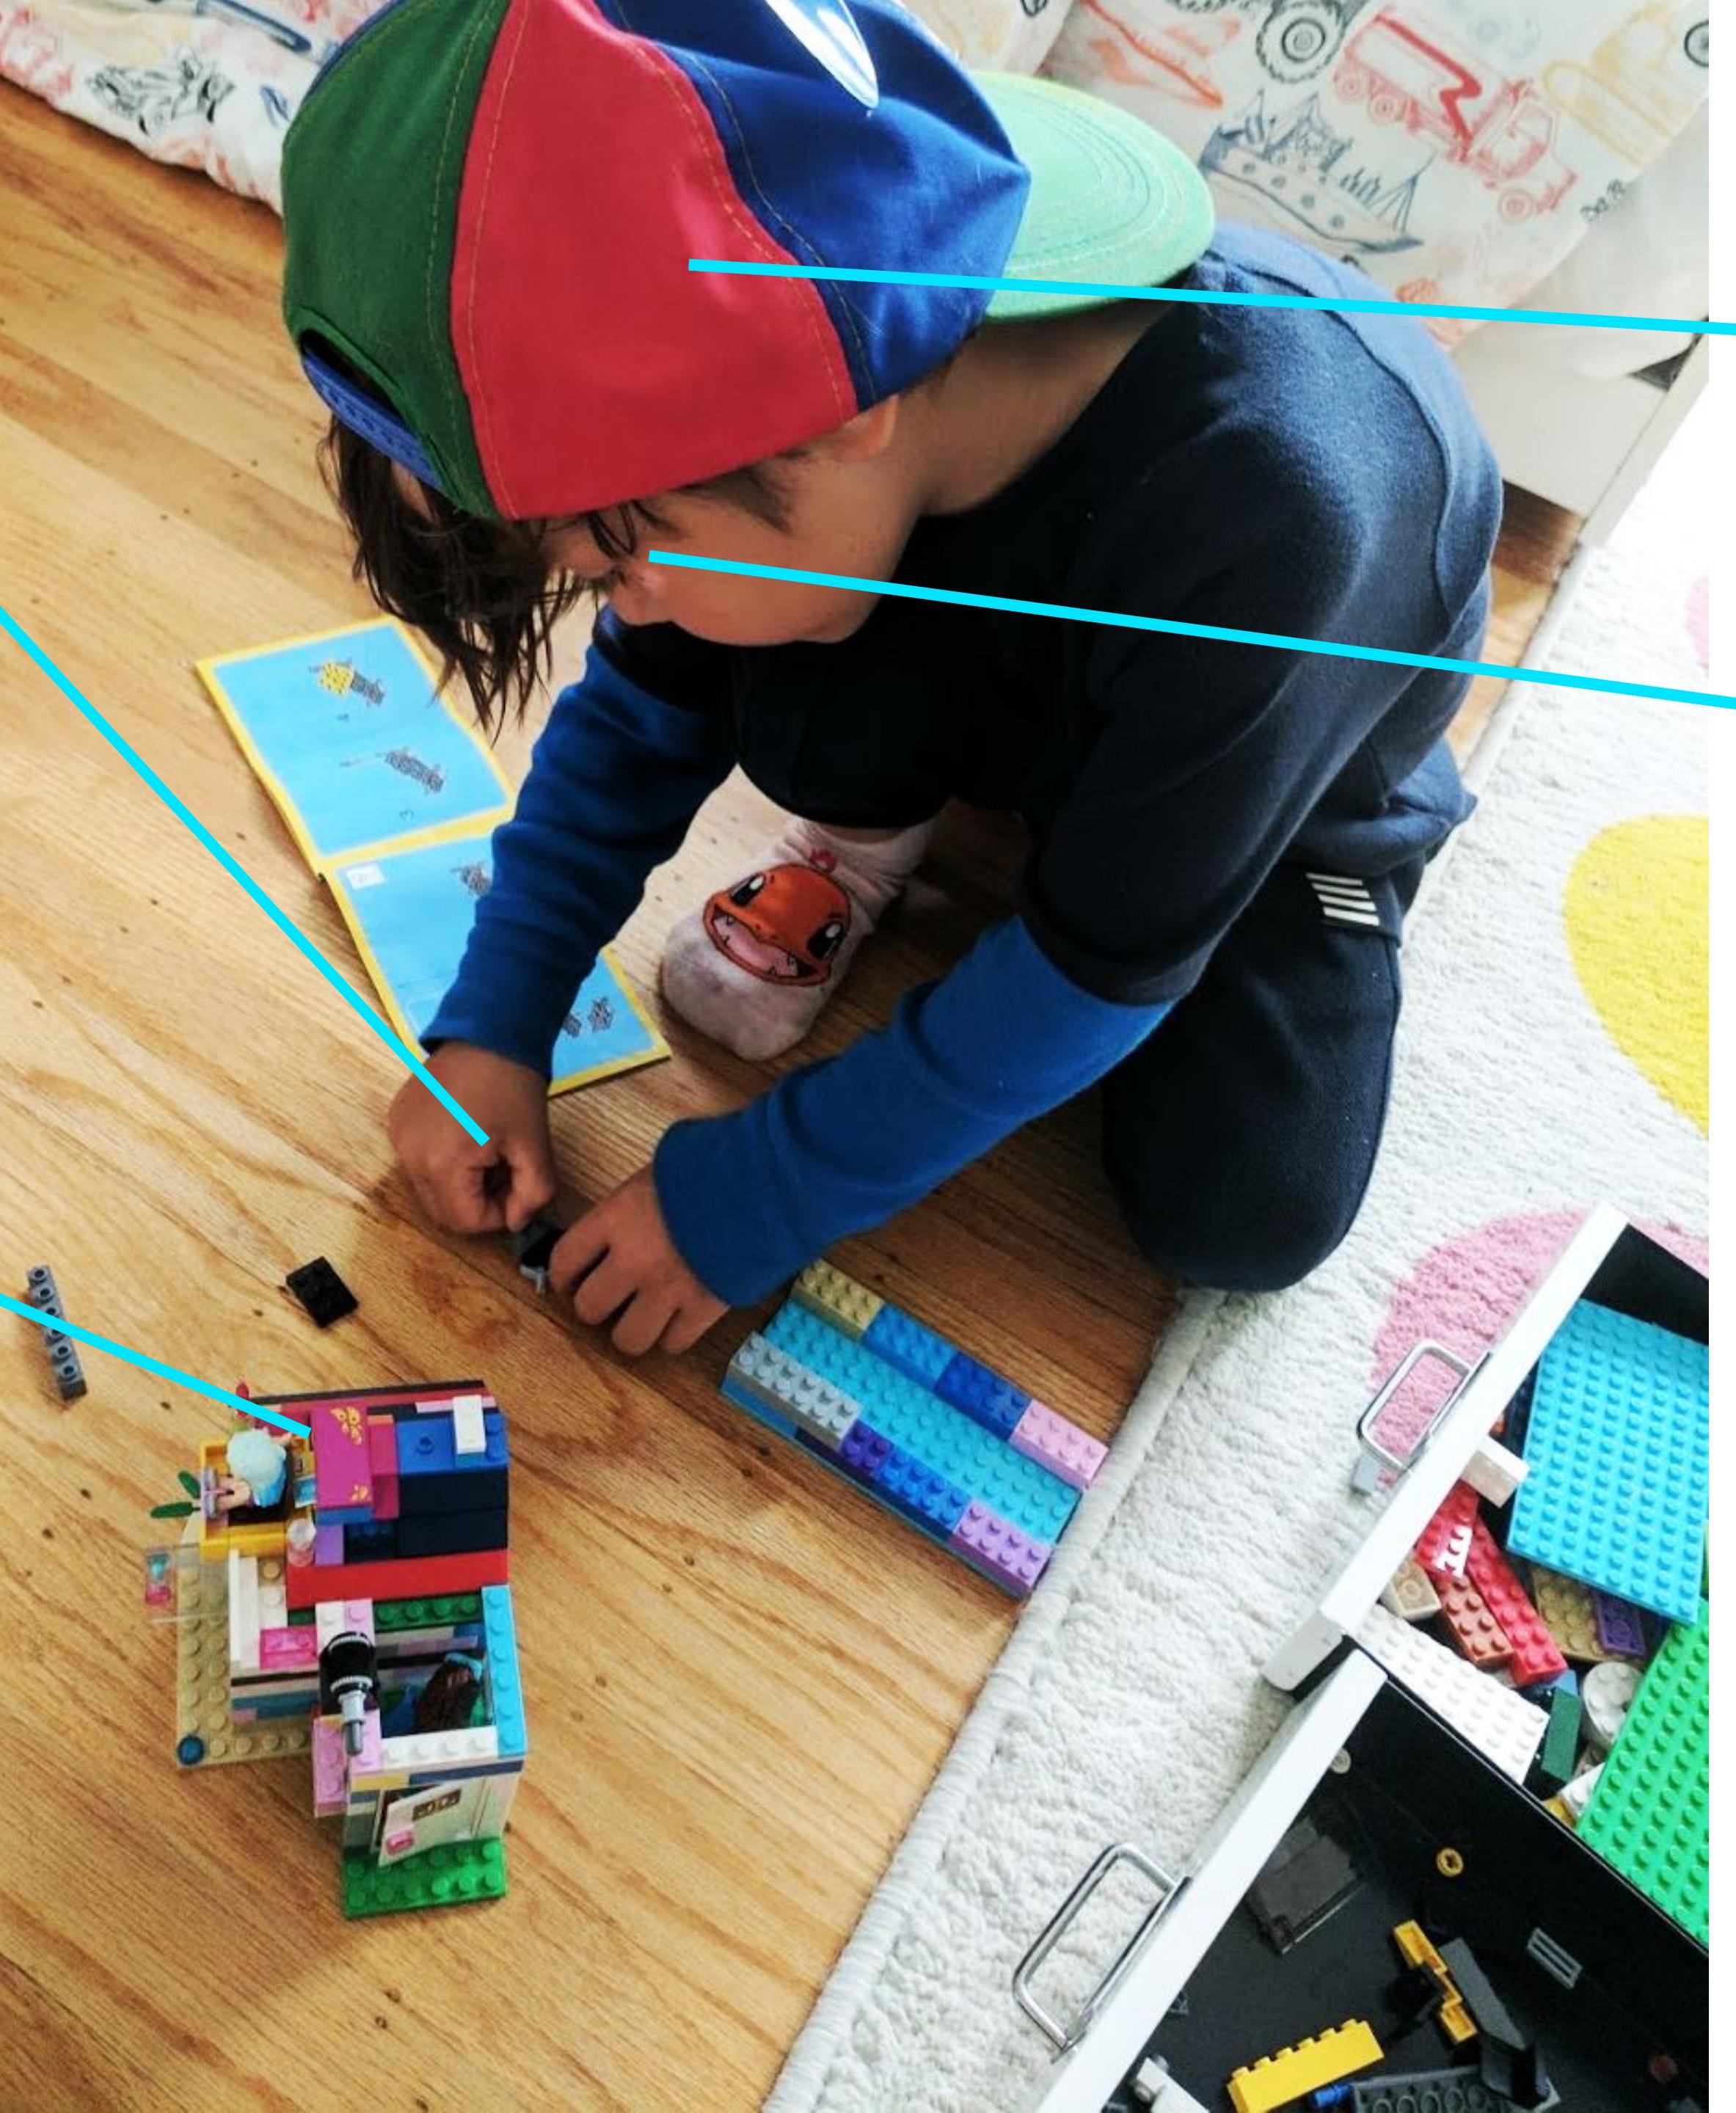
Keras is a high-level API for building
TensorFlow models, to use with
Estimators and Experiments.



Building TensorFlow models from scratch



The core ideas behind deep learning are simple, and so should be their implementation.



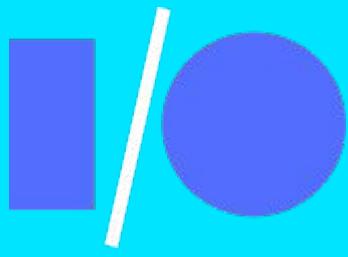
Intuitive affordances

Expressive:
build anything

Fast experimentation,
continuous learning

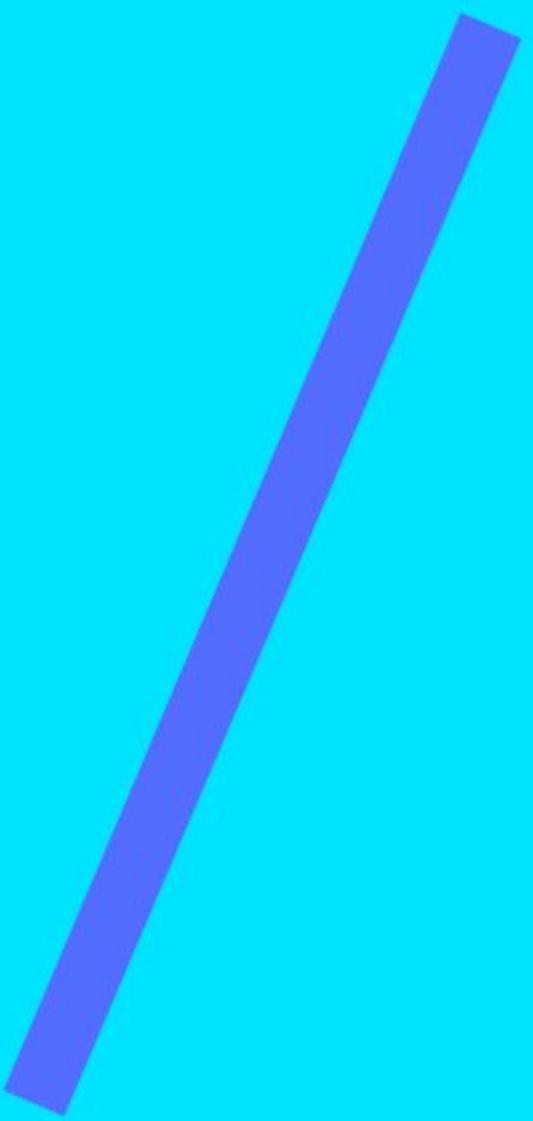
Instant, clear feedback

Fun, accessible!



Keras

Deep learning for everyone





A **high-level** API specification





A **high-level** API specification



Built with a focus on **User Experience**





/ A **high-level** API specification

/ Built with a focus on **User Experience**

/ “Deep learning accessible to everyone”

Keras API spec

TF-Keras

Theano-Keras

...

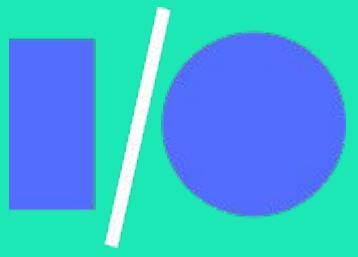
...

TensorFlow
workflow

The Keras API as `tf.keras`

For TensorFlow users: a simplified workflow

For Keras users: access to more powerful features



Example workflow

A toy video-QA problem

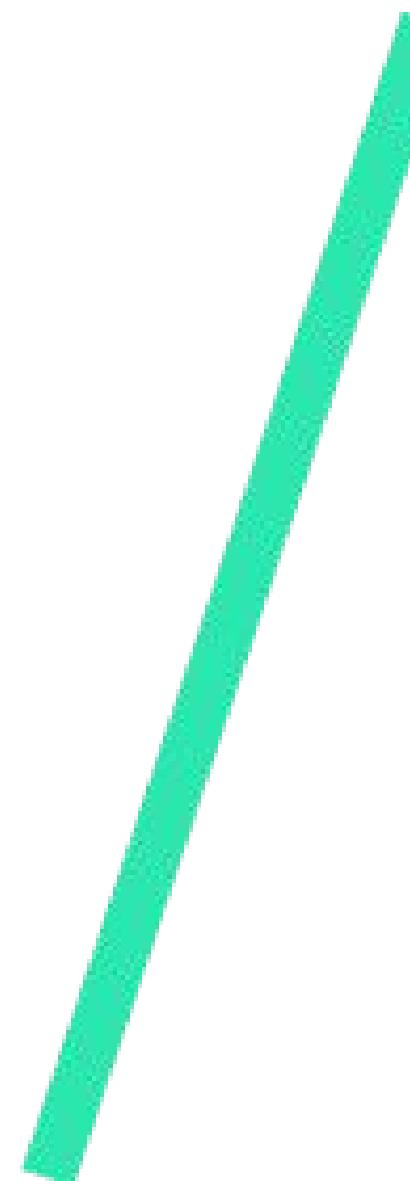


Toy video-QA problem

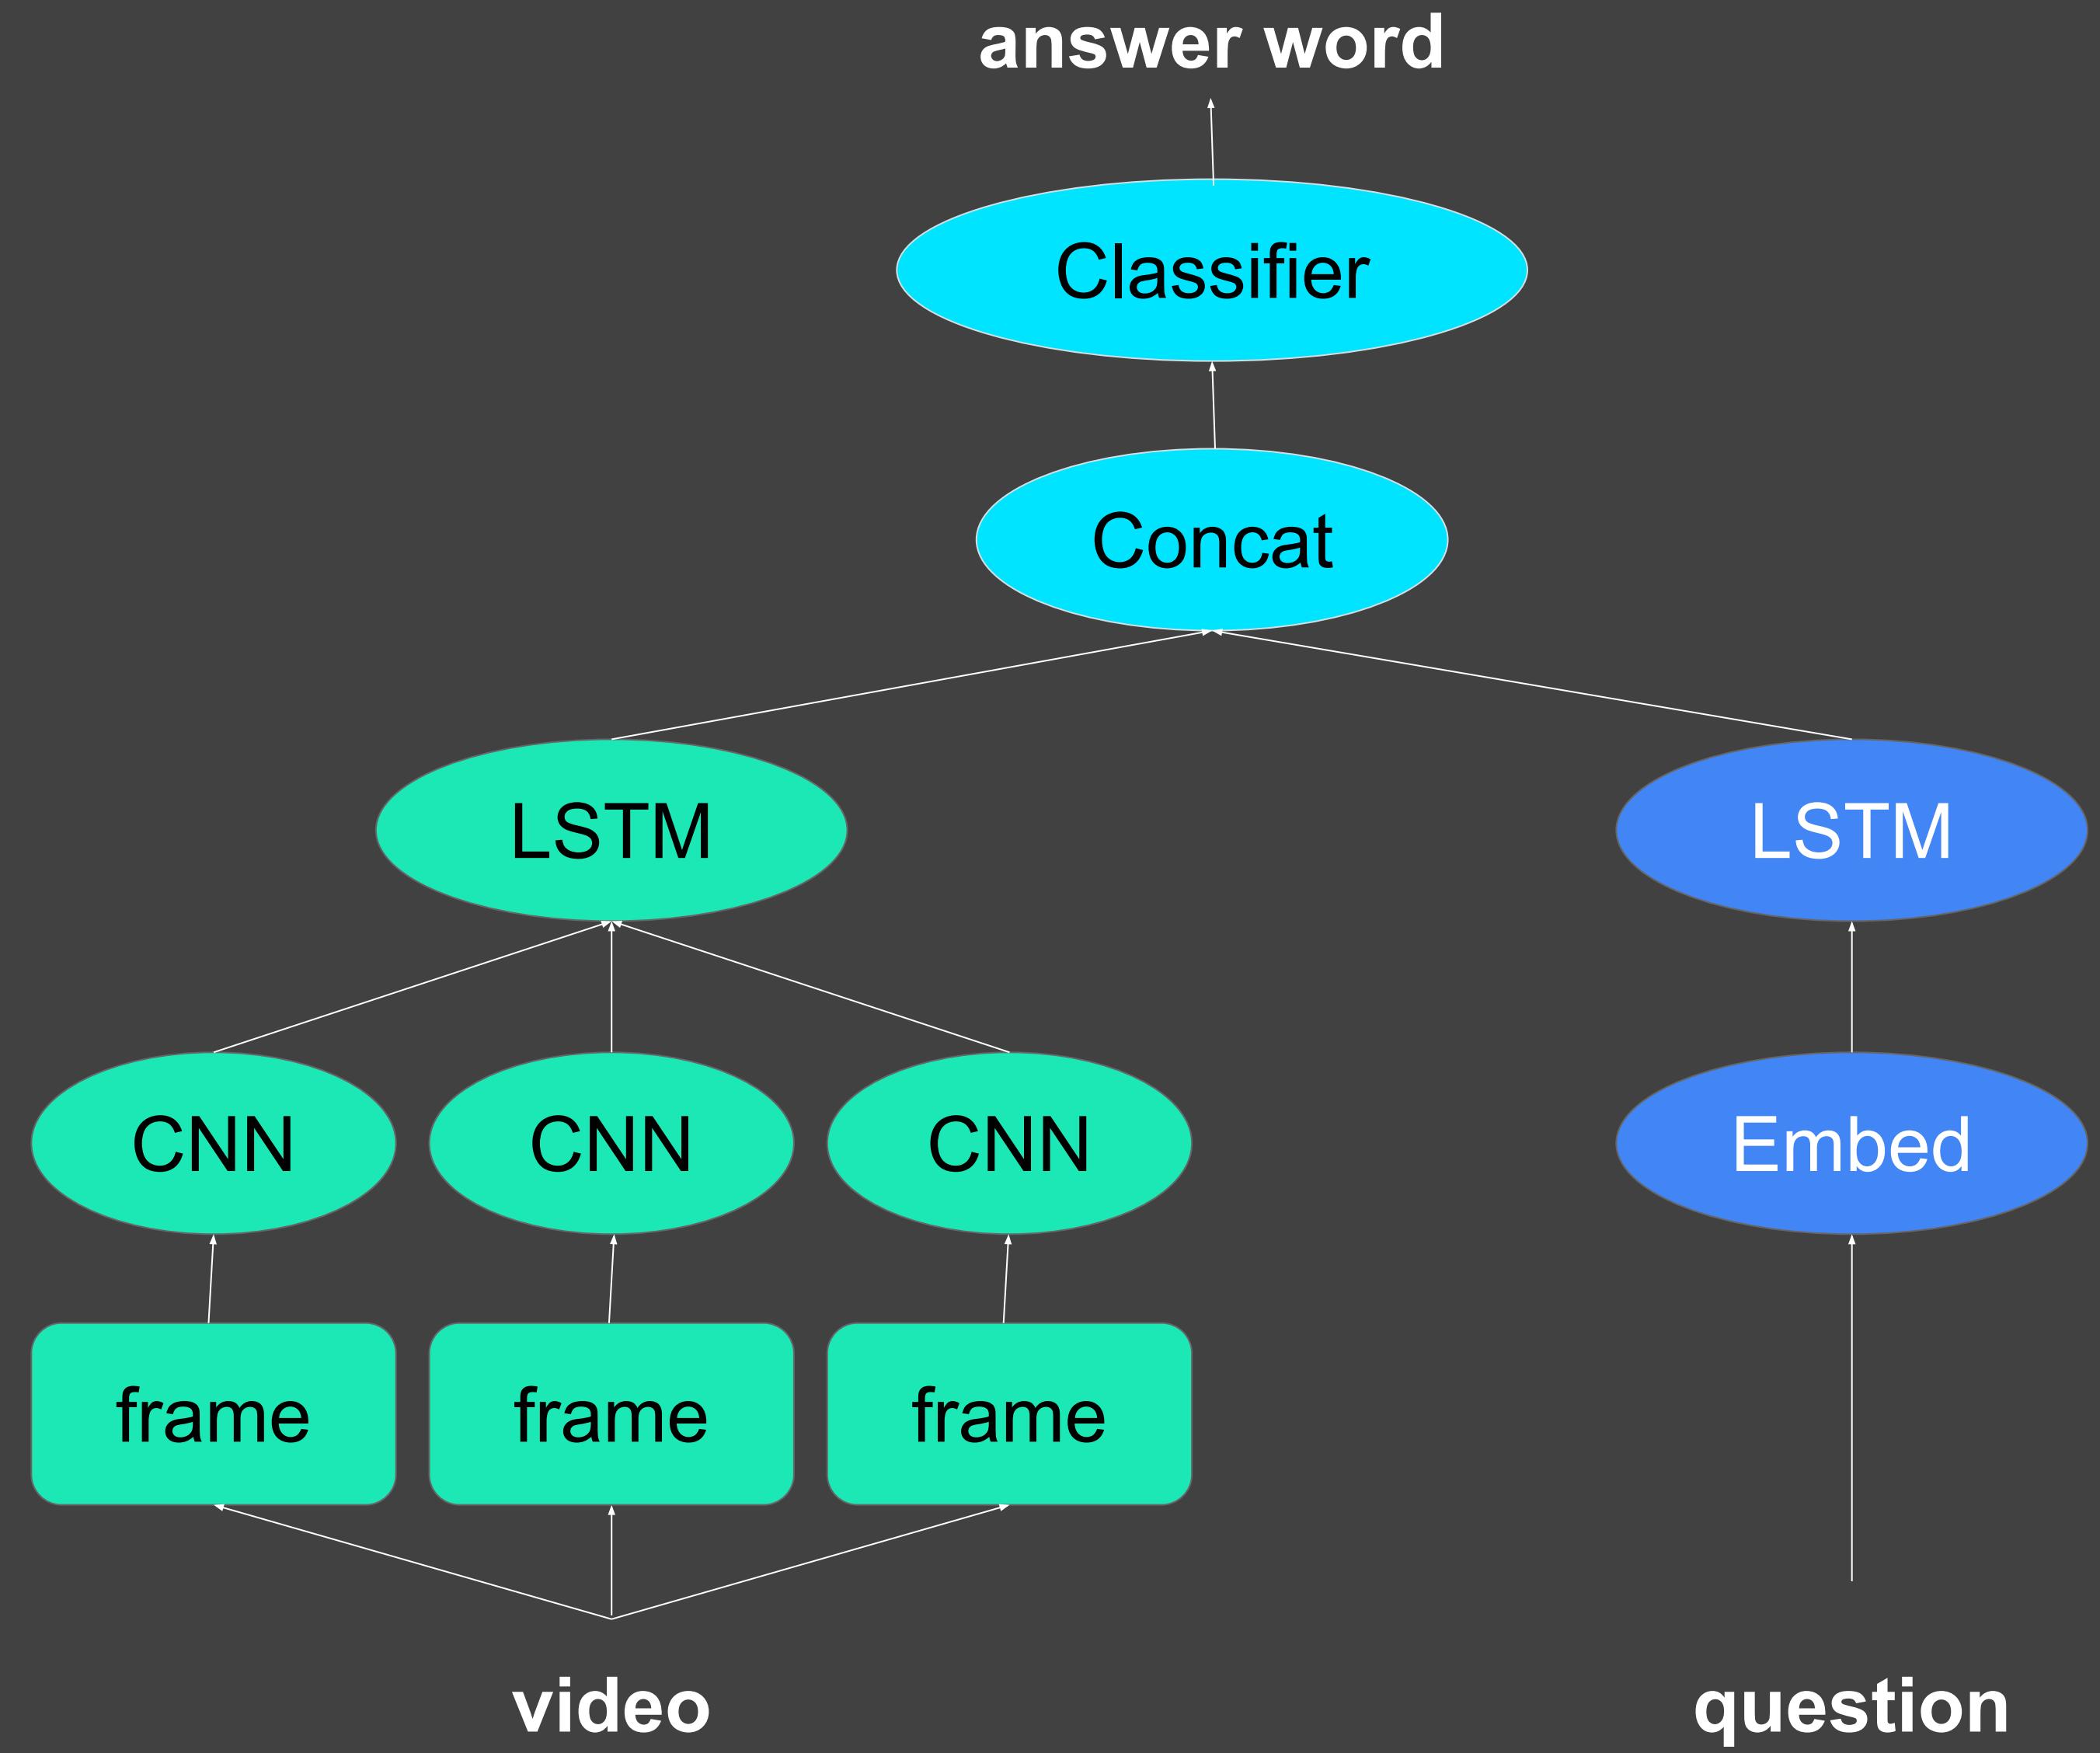


- > “**What is the man doing?**”
- > **packing**

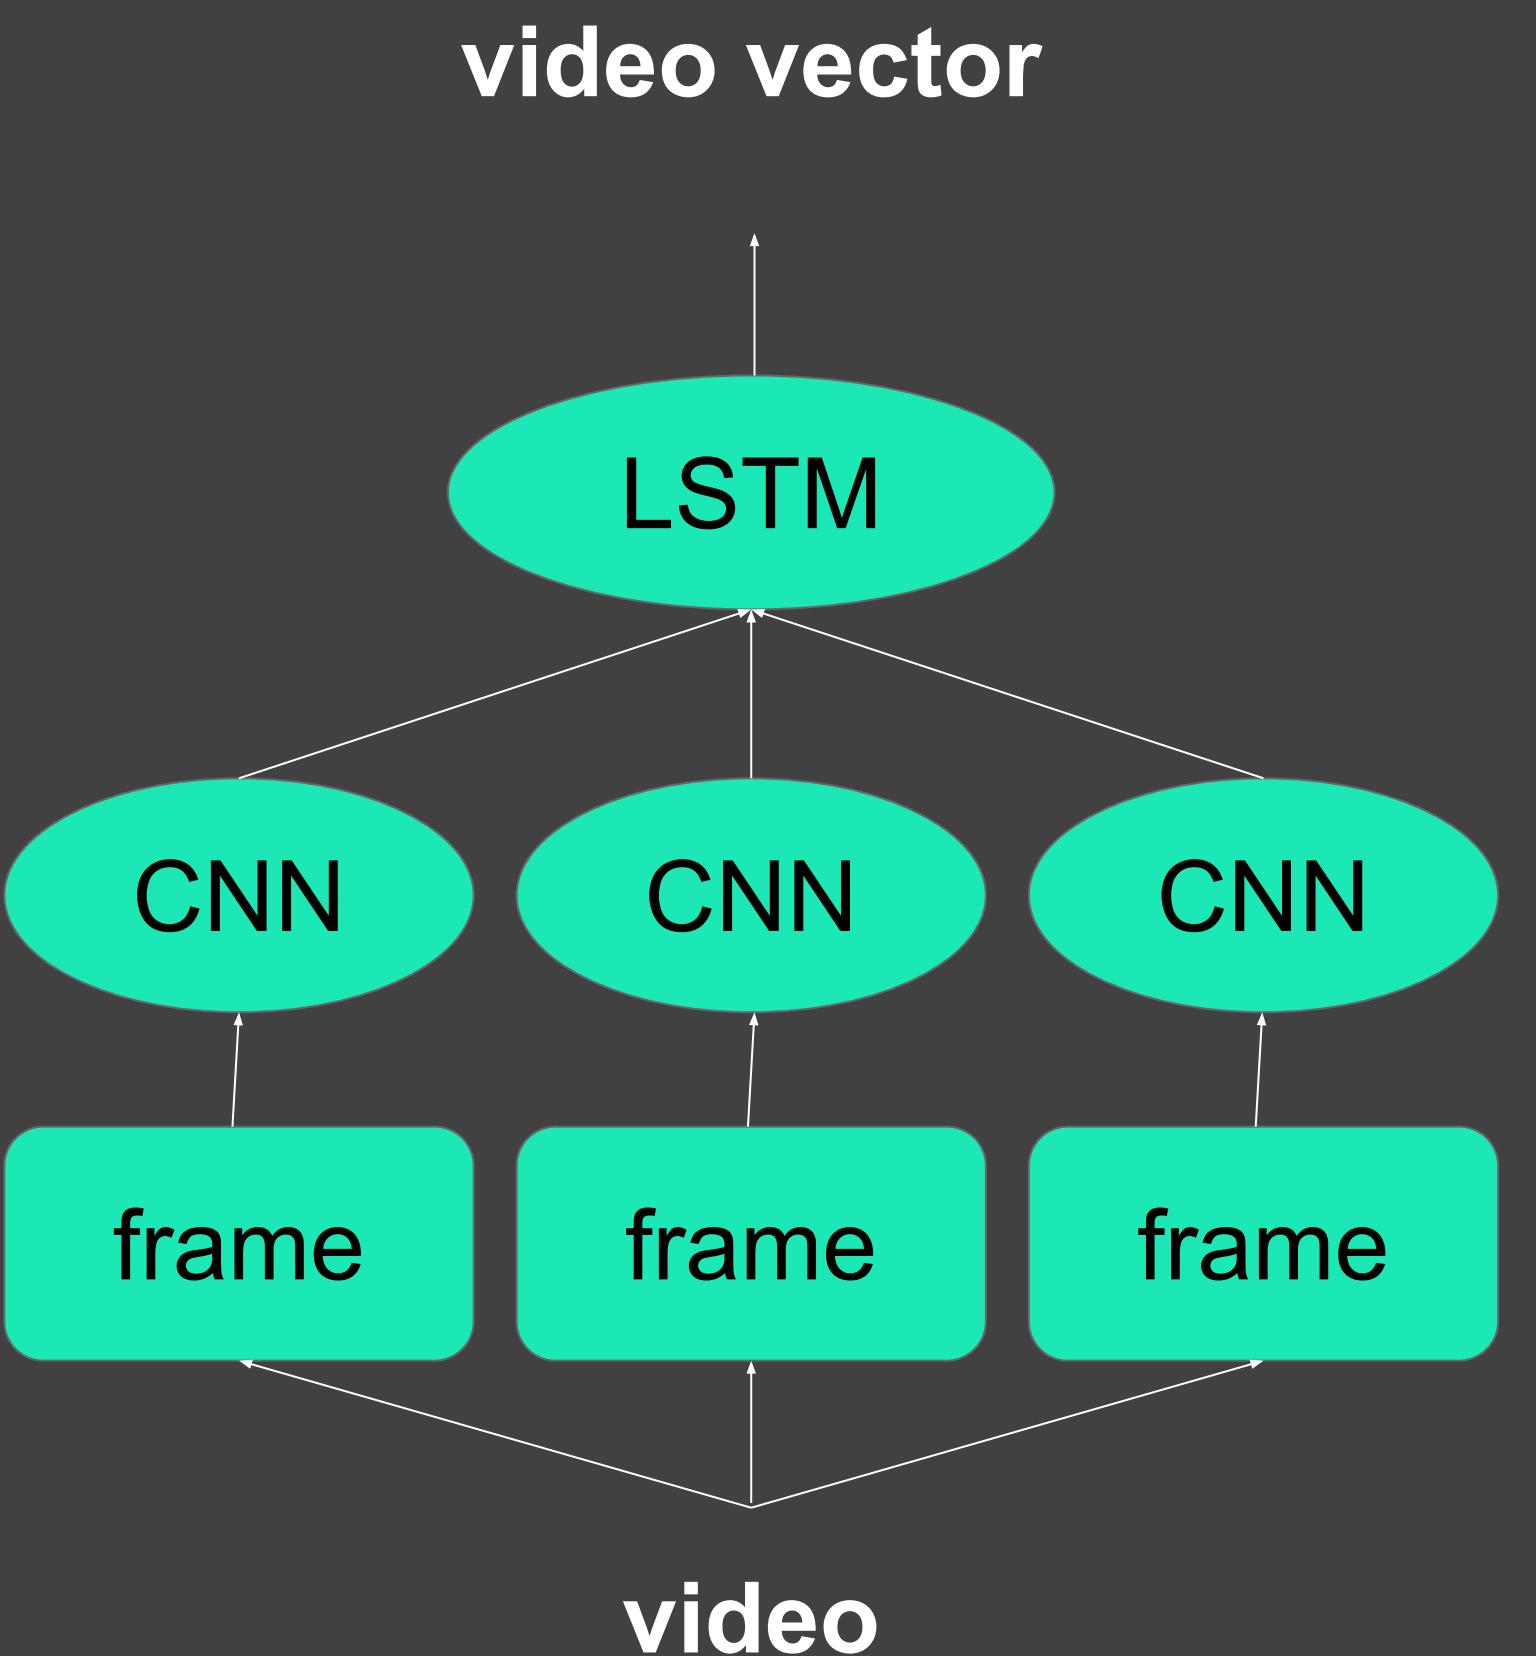
- > “**What color is his shirt?**”
- > **blue**

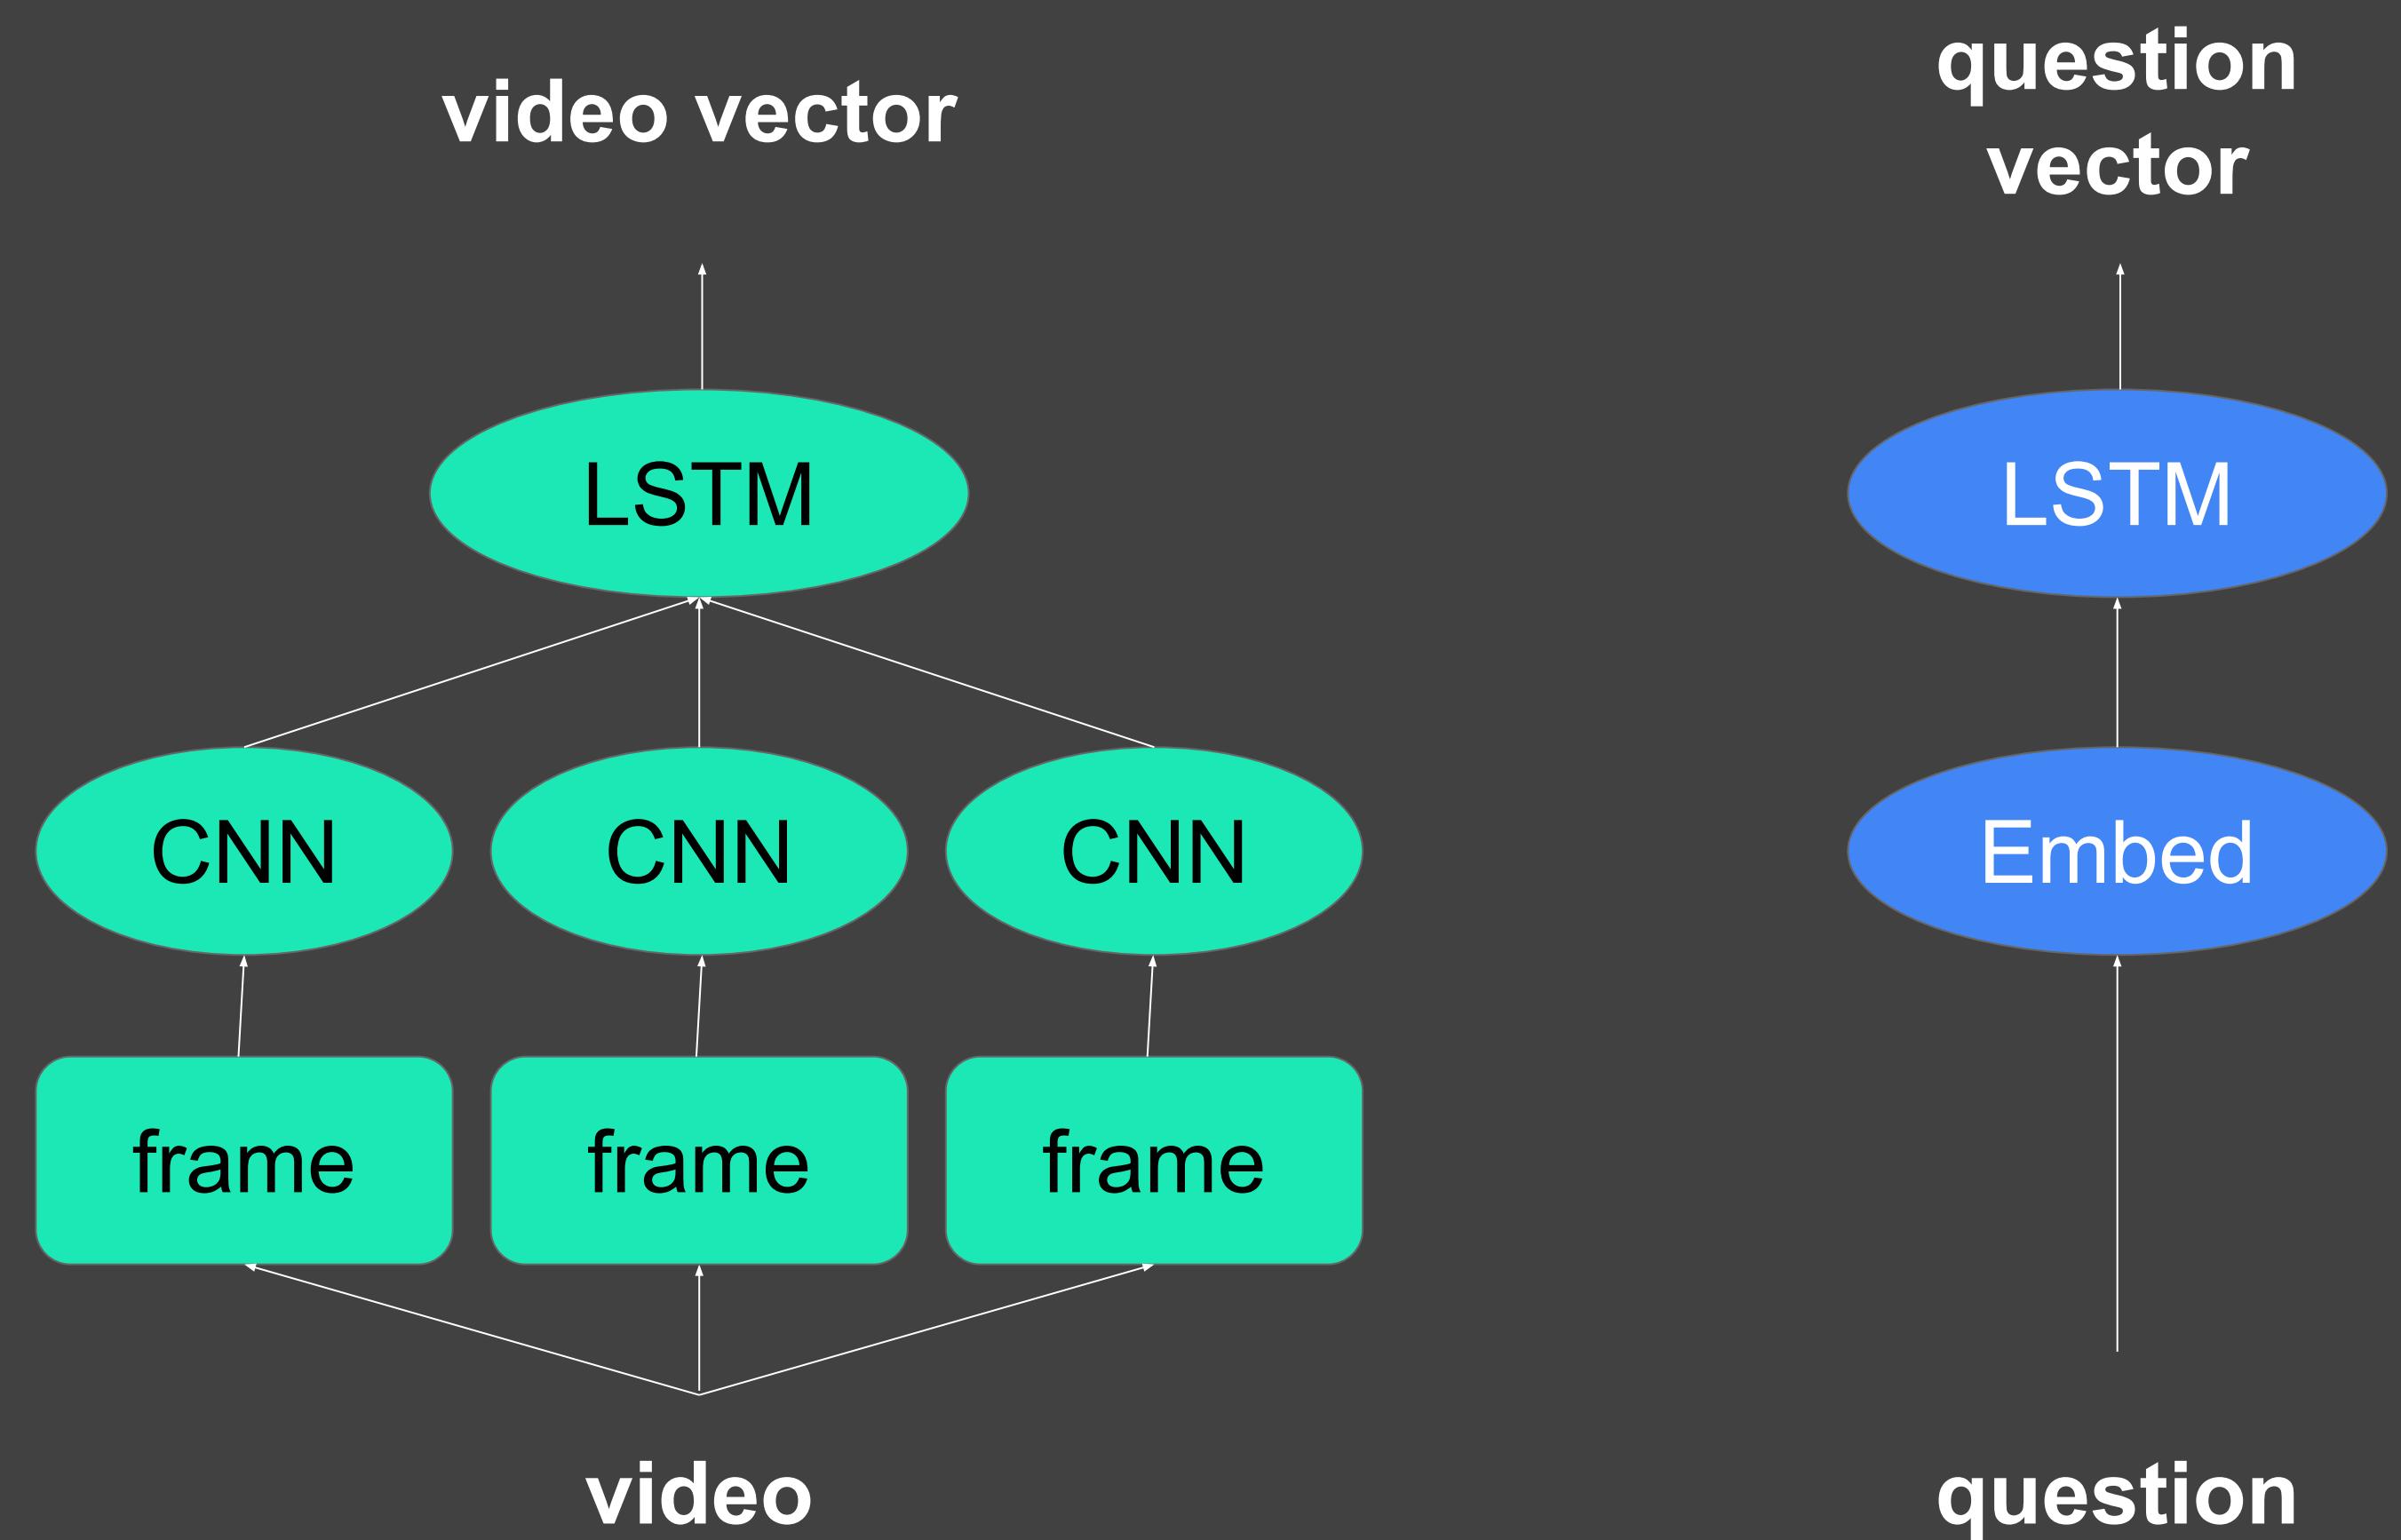


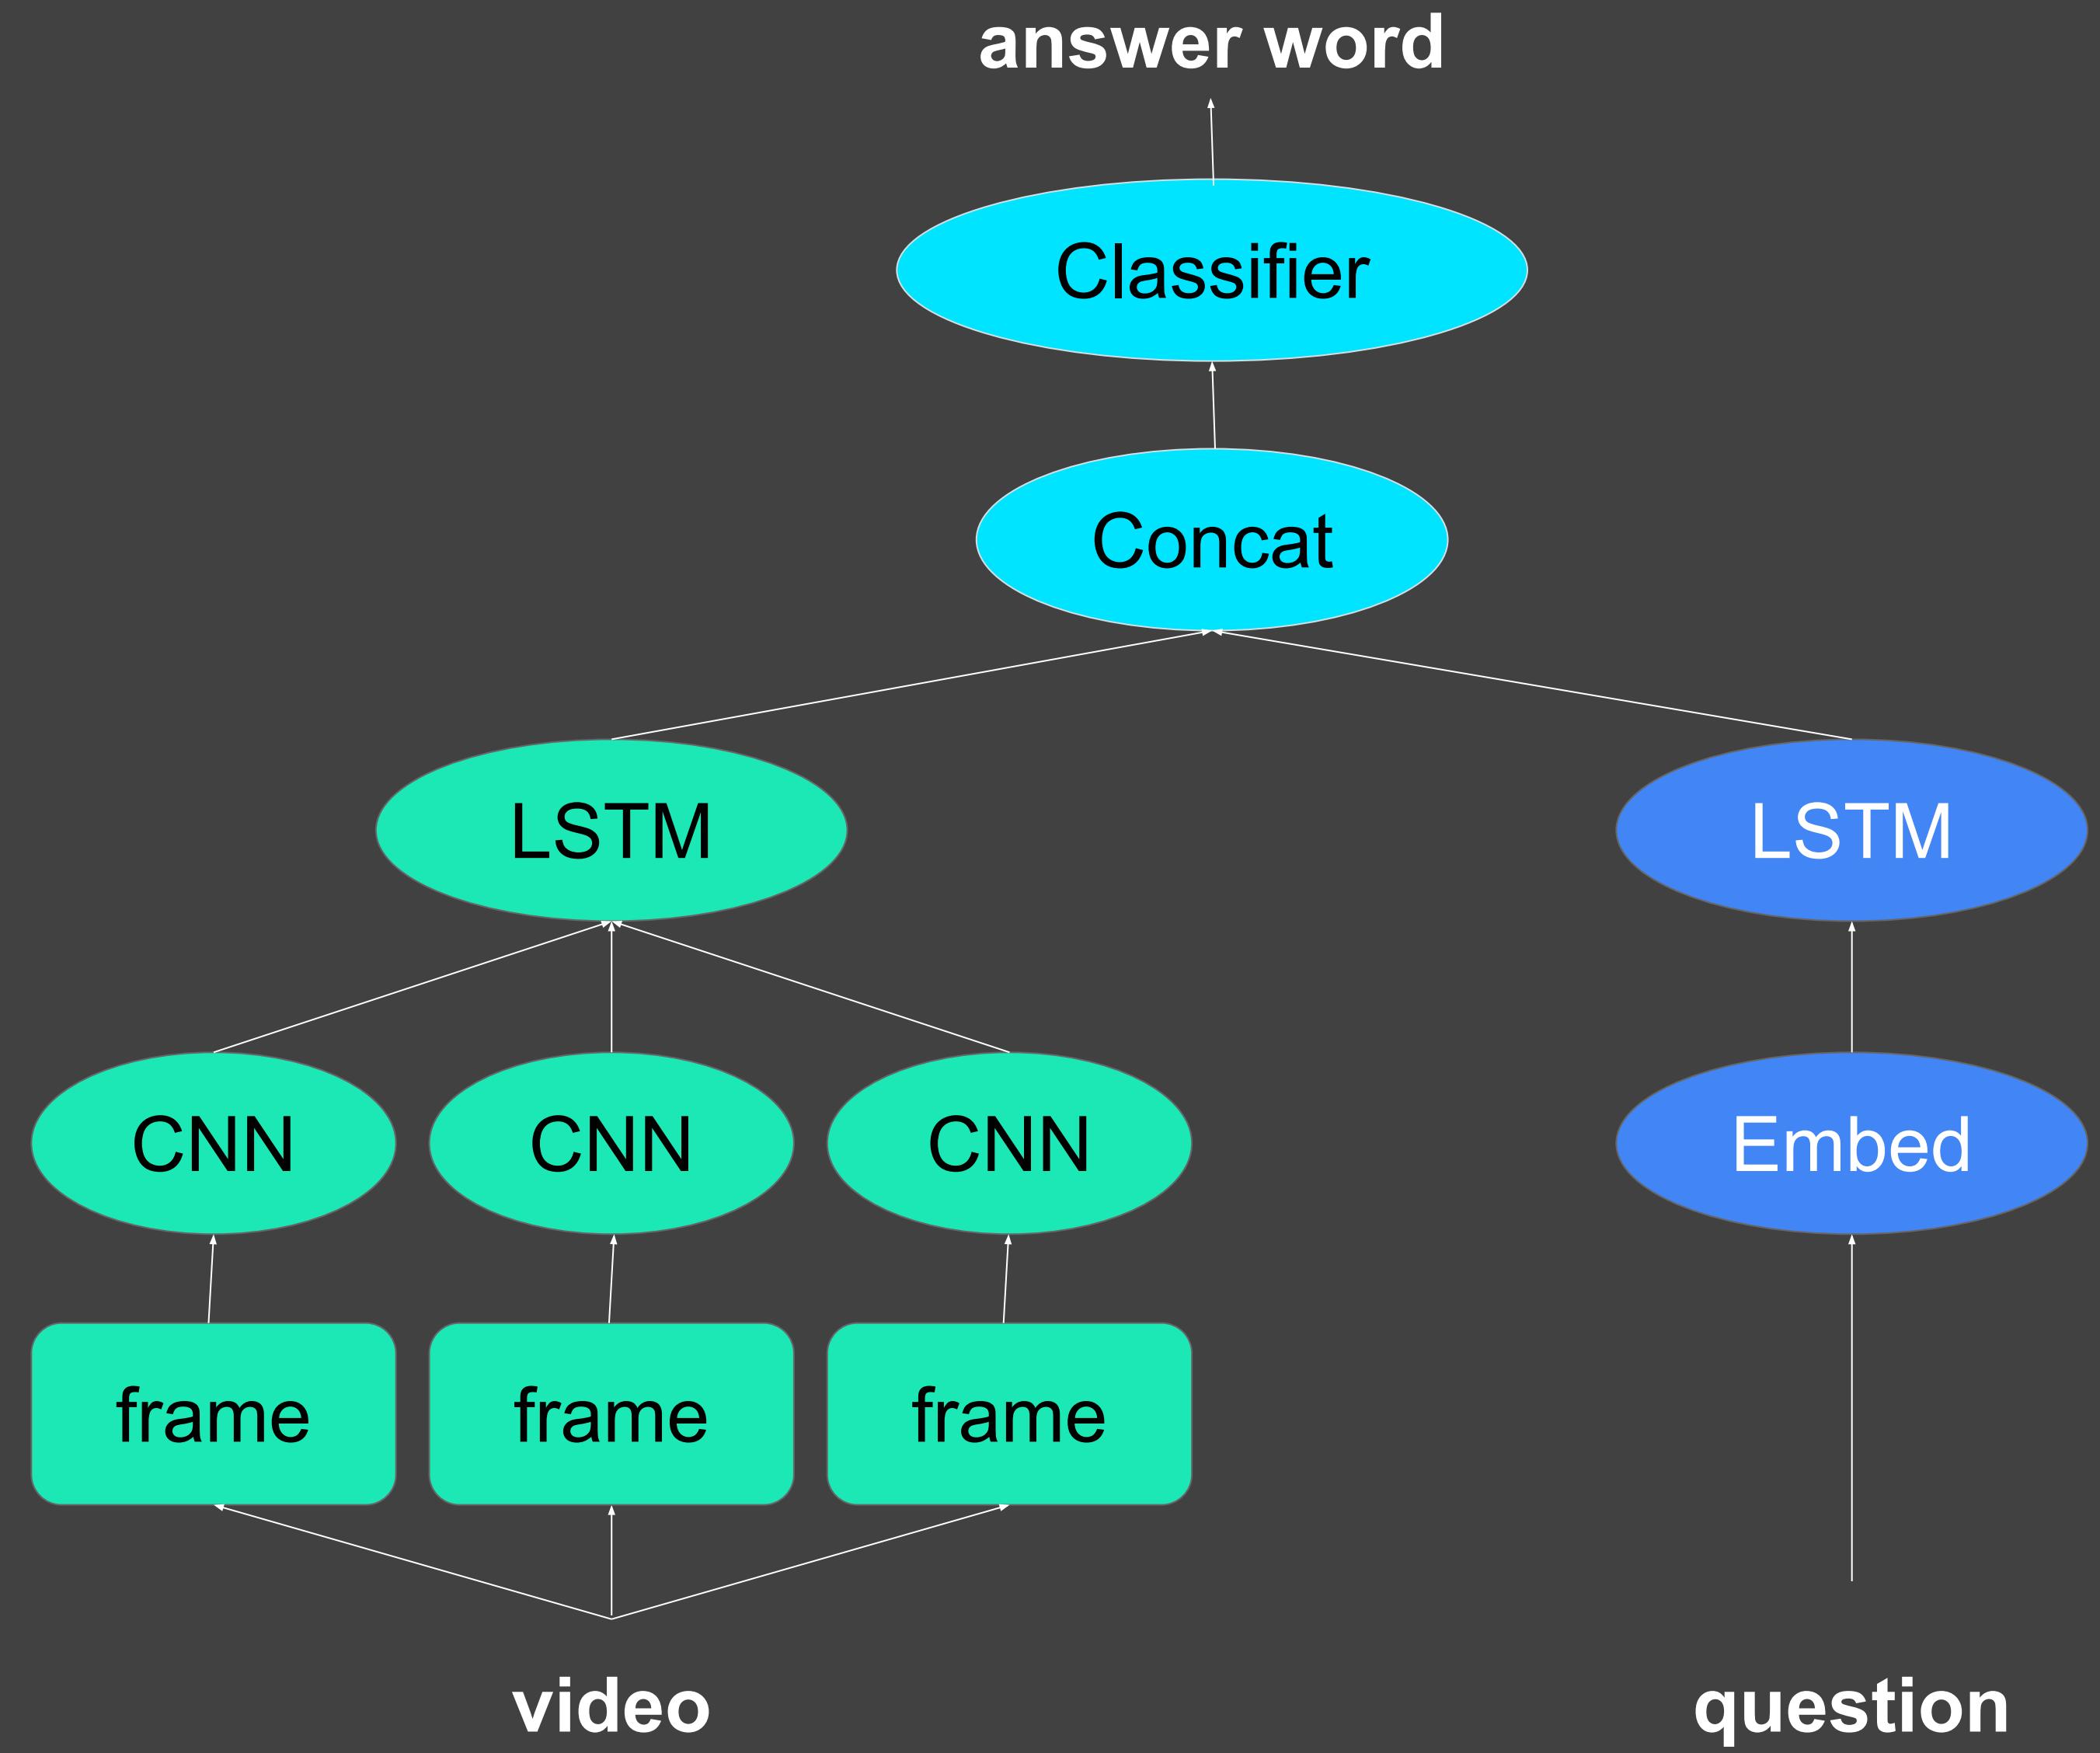
A previously very hard problem,
made accessible to anyone with
basic Python scripting abilities

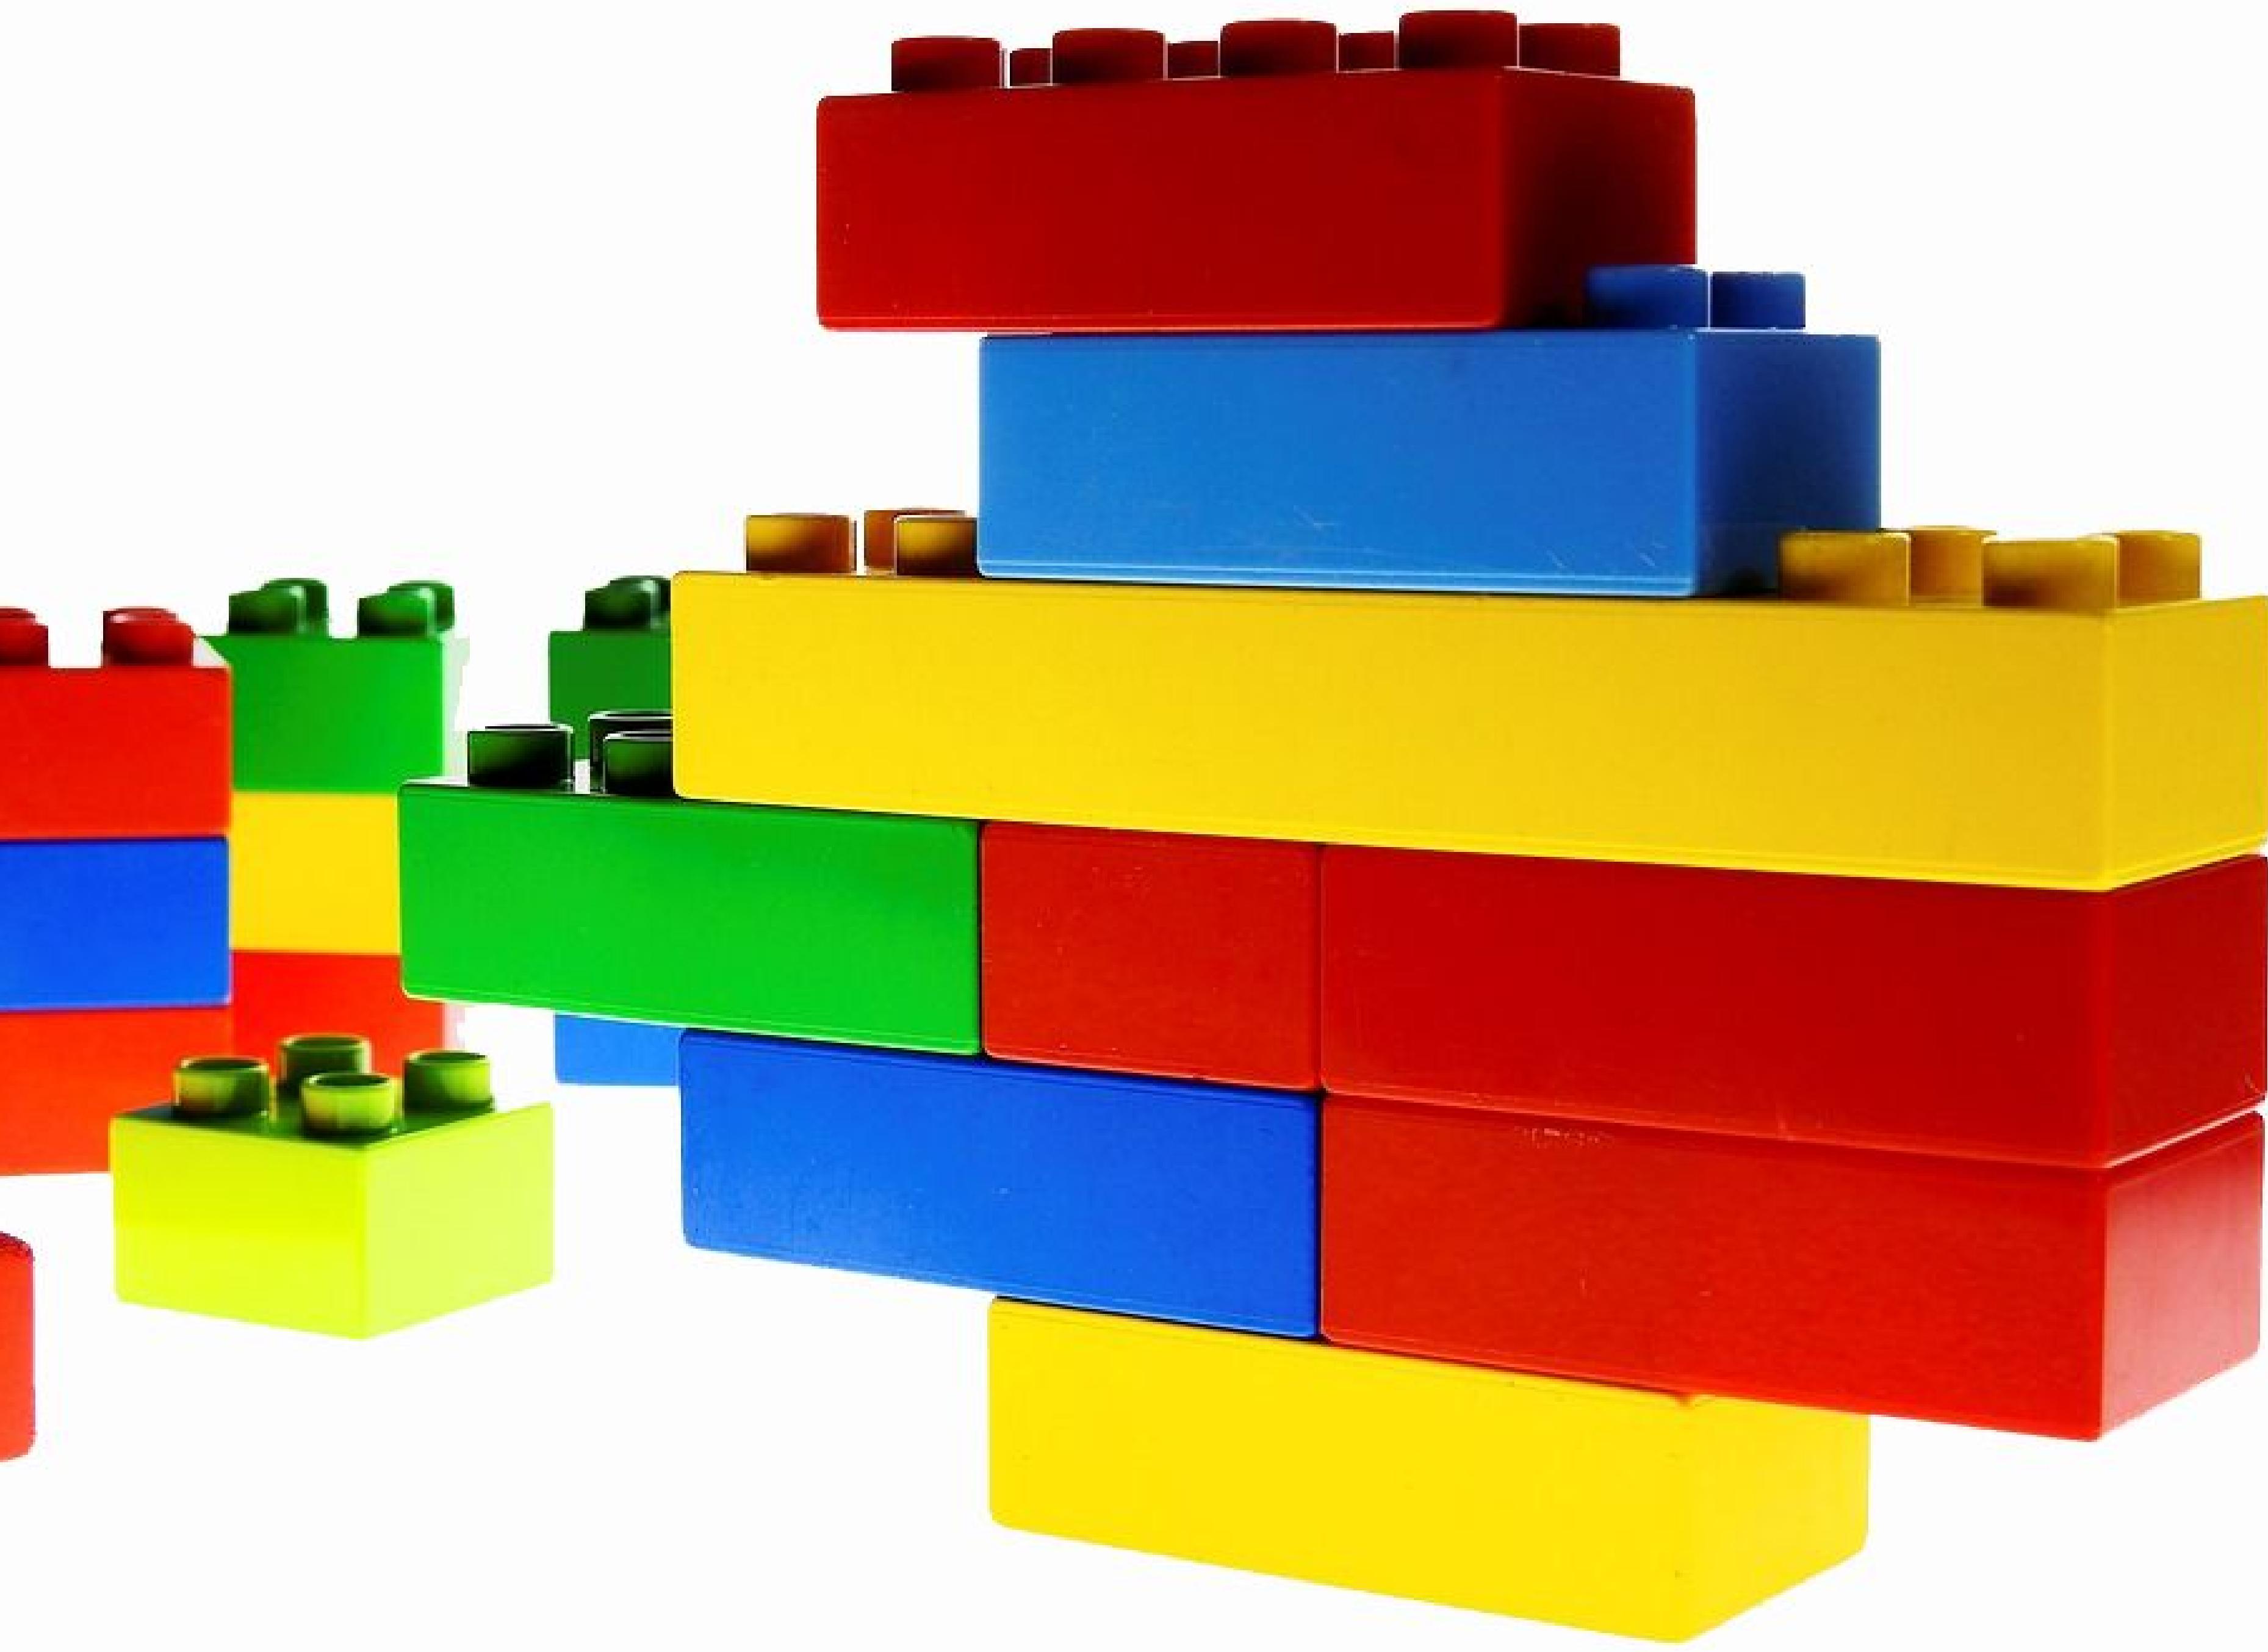


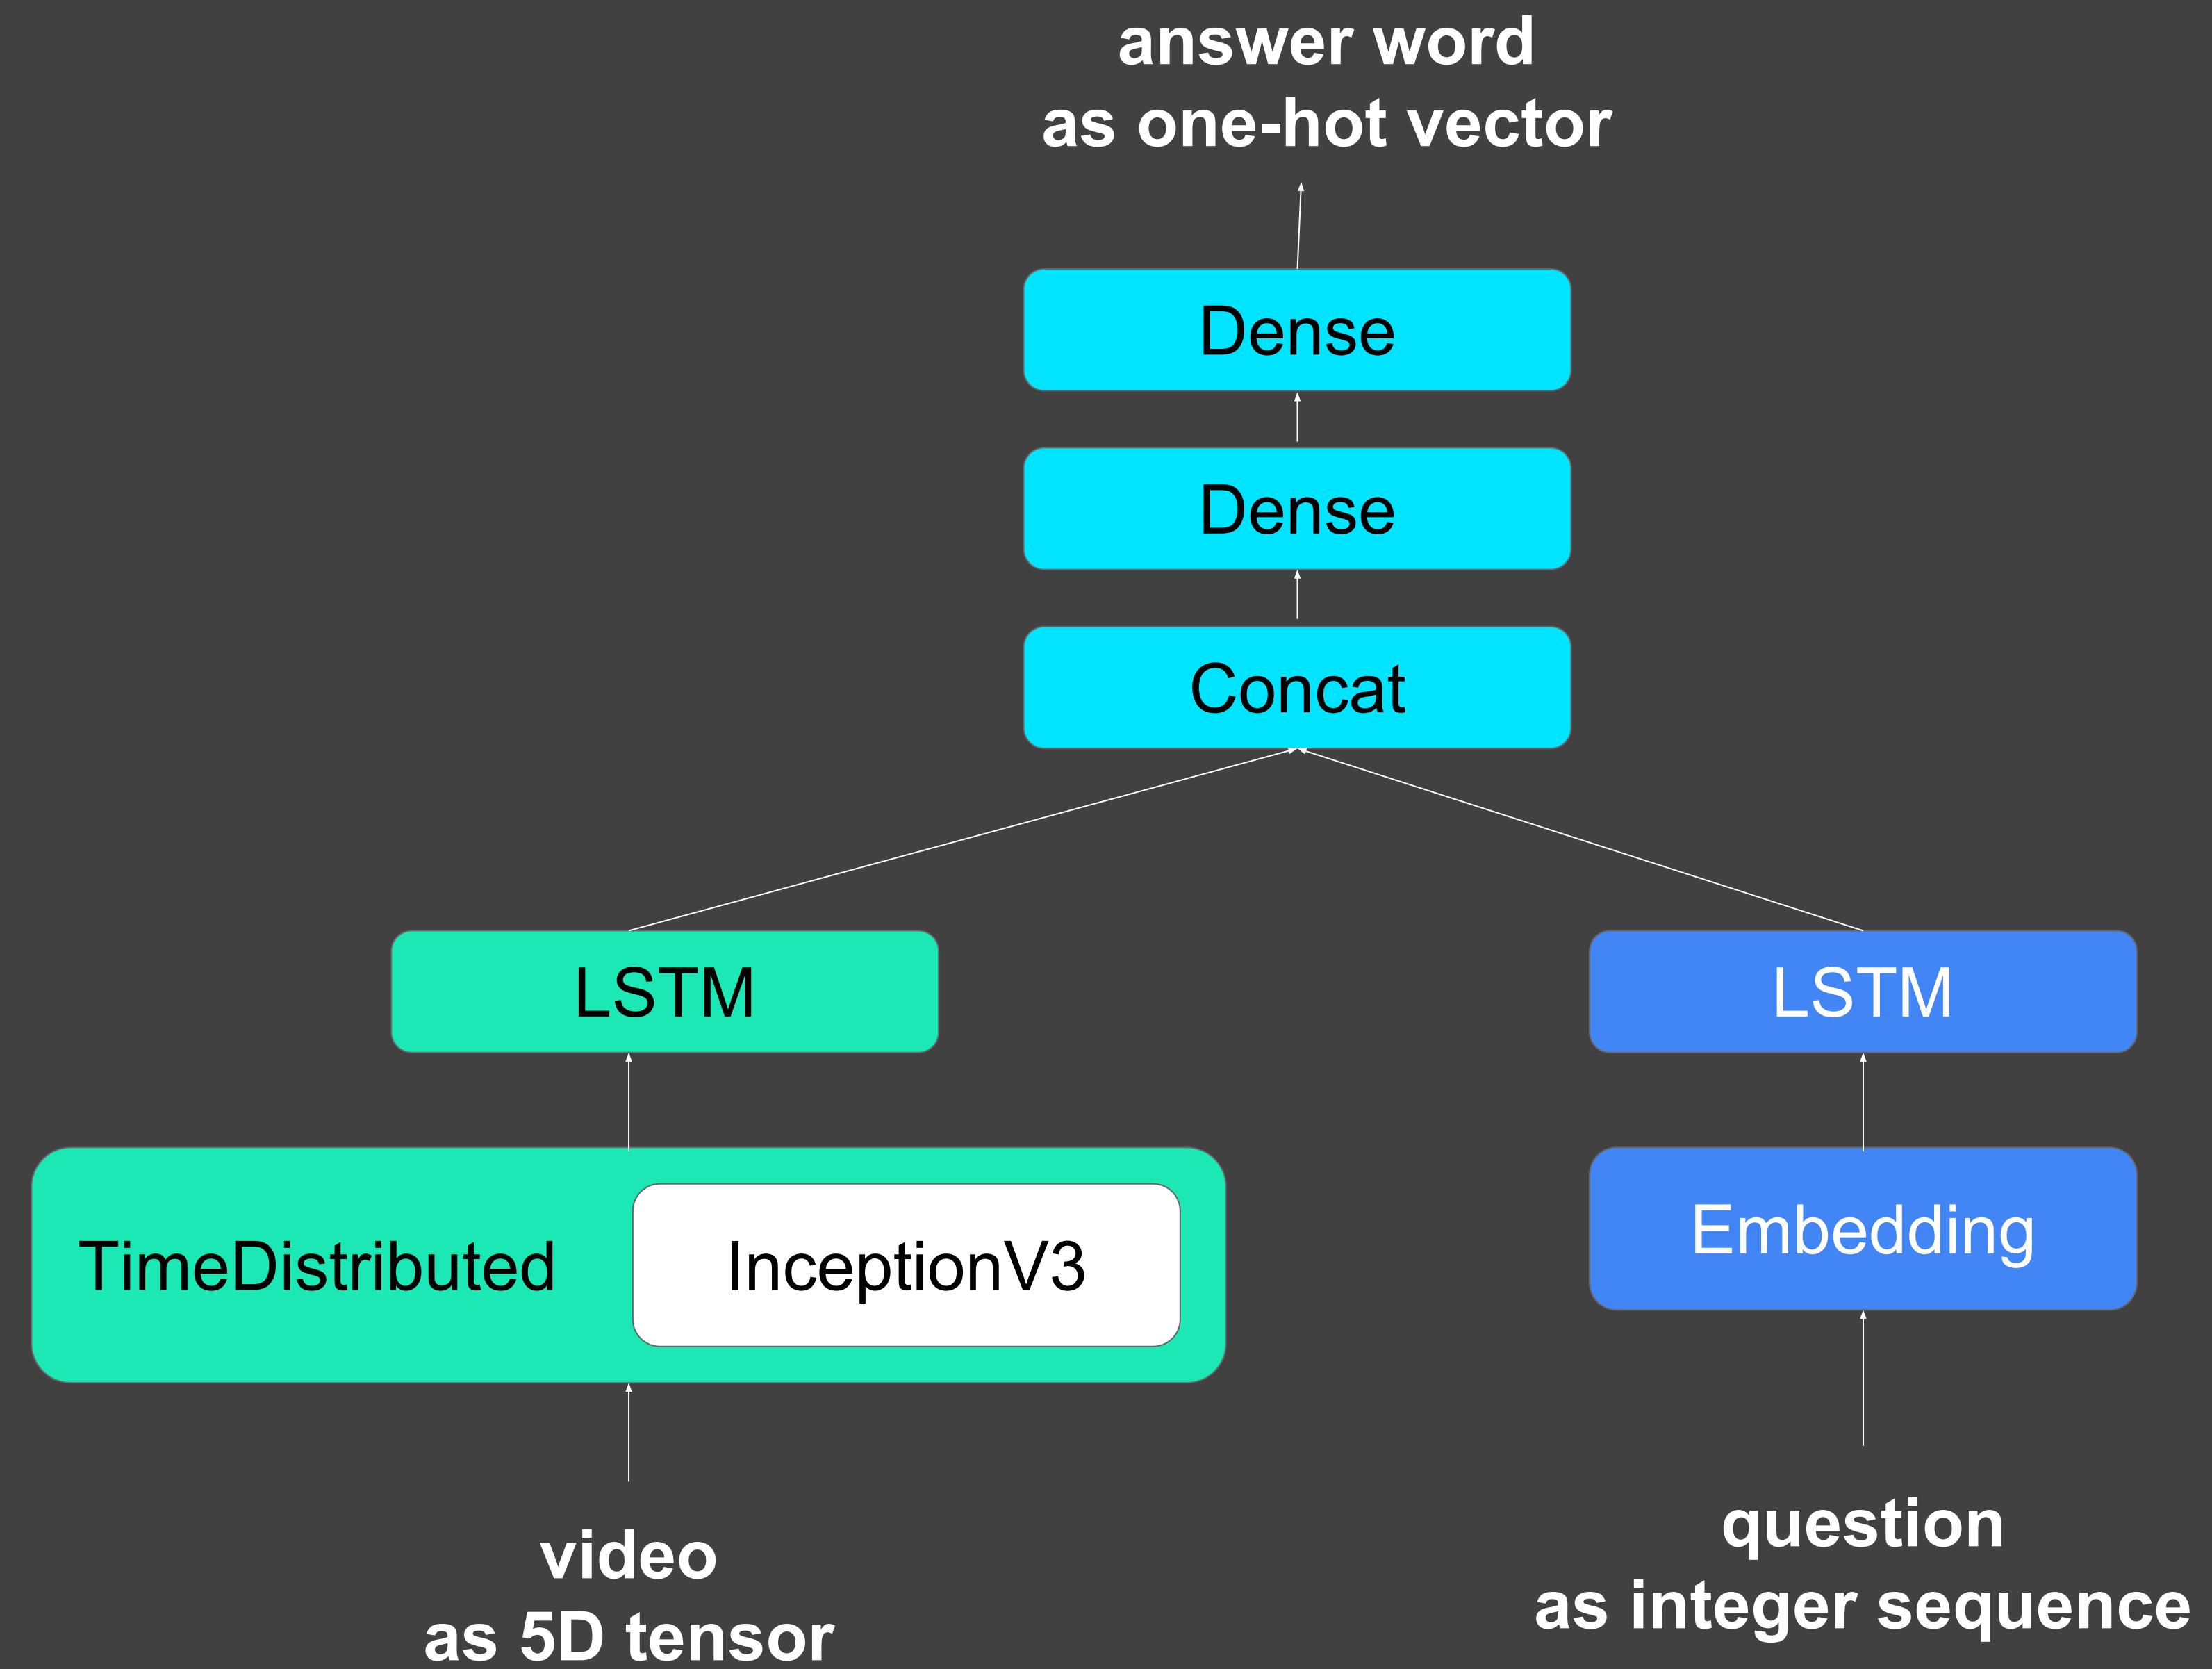
from frames to a vector











Turning frames into a vector, with pre-trained representations

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import InceptionV3

video = keras.Input(shape=(None, 150, 150, 3), name='video')
cnn = InceptionV3(weights='imagenet',
                    include_top=False,
                    pooling='avg')
cnn.trainable = False
frame_features = layers.TimeDistributed(cnn)(video)
video_vector = layers.LSTM(256)(frame_features)
```

Turning frames into a vector, with pre-trained representations

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import InceptionV3

video = keras.Input(shape=(None, 150, 150, 3), name='video')
cnn = InceptionV3(weights='imagenet',
                    include_top=False,
                    pooling='avg')

cnn.trainable = False
frame_features = layers.TimeDistributed(cnn)(video)
video_vector = layers.LSTM(256)(frame_features)
```

Turning frames into a vector, with pre-trained representations

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import InceptionV3

video = keras.Input(shape=(None, 150, 150, 3), name='video')
cnn = InceptionV3(weights='imagenet',
                    include_top=False,
                    pooling='avg')
cnn.trainable = False
frame_features = layers.TimeDistributed(cnn)(video)
video_vector = layers.LSTM(256)(frame_features)
```

Turning frames into a vector, with pre-trained representations

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import InceptionV3

video = keras.Input(shape=(None, 150, 150, 3), name='video')
cnn = InceptionV3(weights='imagenet',
                    include_top=False,
                    pooling='avg')
cnn.trainable = False
frame_features = layers.TimeDistributed(cnn)(video)
video_vector = layers.LSTM(256)(frame_features)
```

Turning frames into a vector, with pre-trained representations

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import InceptionV3

video = keras.Input(shape=(None, 150, 150, 3), name='video')
cnn = InceptionV3(weights='imagenet',
                    include_top=False,
                    pooling='avg')
cnn.trainable = False
frame_features = layers.TimeDistributed(cnn)(video)
video_vector = layers.LSTM(256)(frame_features)
```

Turning a sequence of words into a vector

```
question = keras.Input(shape=(None,), dtype='int32', name='question')
embedded_words = layers.Embedding(input_voc_size, 256)(question)
question_vector = layers.LSTM(128)(embedded_words)
```

Predicting an answer word

```
x = layers.concatenate([video_vector, question_vector])
x = layers.Dense(128, activation=tf.nn.relu)(x)
predictions = layers.Dense(output_voc_size, name='predictions')(x)
```

Setting up the training configuration

```
model = keras.models.Model([video, question], predictions)
model.compile(optimizer=tf.AdamOptimizer(),
              loss=tf.softmax_crossentropy_with_logits)
```

Leveraging Experiment for distributed training

```
def experiment_fn(config, params):
    model = ...
    estimator = model.get_estimator(config=...)
    return Experiment(estimator,
                      train_input_fn=pandas_input_fn(...),
                      eval_input_fn=pandas_input_fn(...))

if __name__ == '__main__':
    tf.contrib.train.run_experiment(experiment_fn)
```

Takeaways

Keras: a high-level API with focus on UX

Now part of TensorFlow as **tf.keras**

To use with **Estimators** and **Experiments** for distributed training

A big step towards making deep learning accessible to everyone

Thank you!



Baohua Liao

liaobaohua@gmail.com